


Lecture 7

BST (Insertion & Deletion), Red-Black Trees


Modifying a BST: Insertion

Modifying a BST: Insertion




Modifying a BST: Insertion

Example:



Modifying a BST: Insertion


Example: Insert a node with 24 as key in the following BST.



Modifying a BST: Insertion

Example: Insert a node with 24 as key in the following BST.


Idea:



Modifying a BST: Insertion

Example: Insert a node with 24 as key in the following BST.


Idea: Find the correct leaf where it can be inserted.



Modifying a BST: Insertion

Example: Insert a node with **24** as key in the following BST.

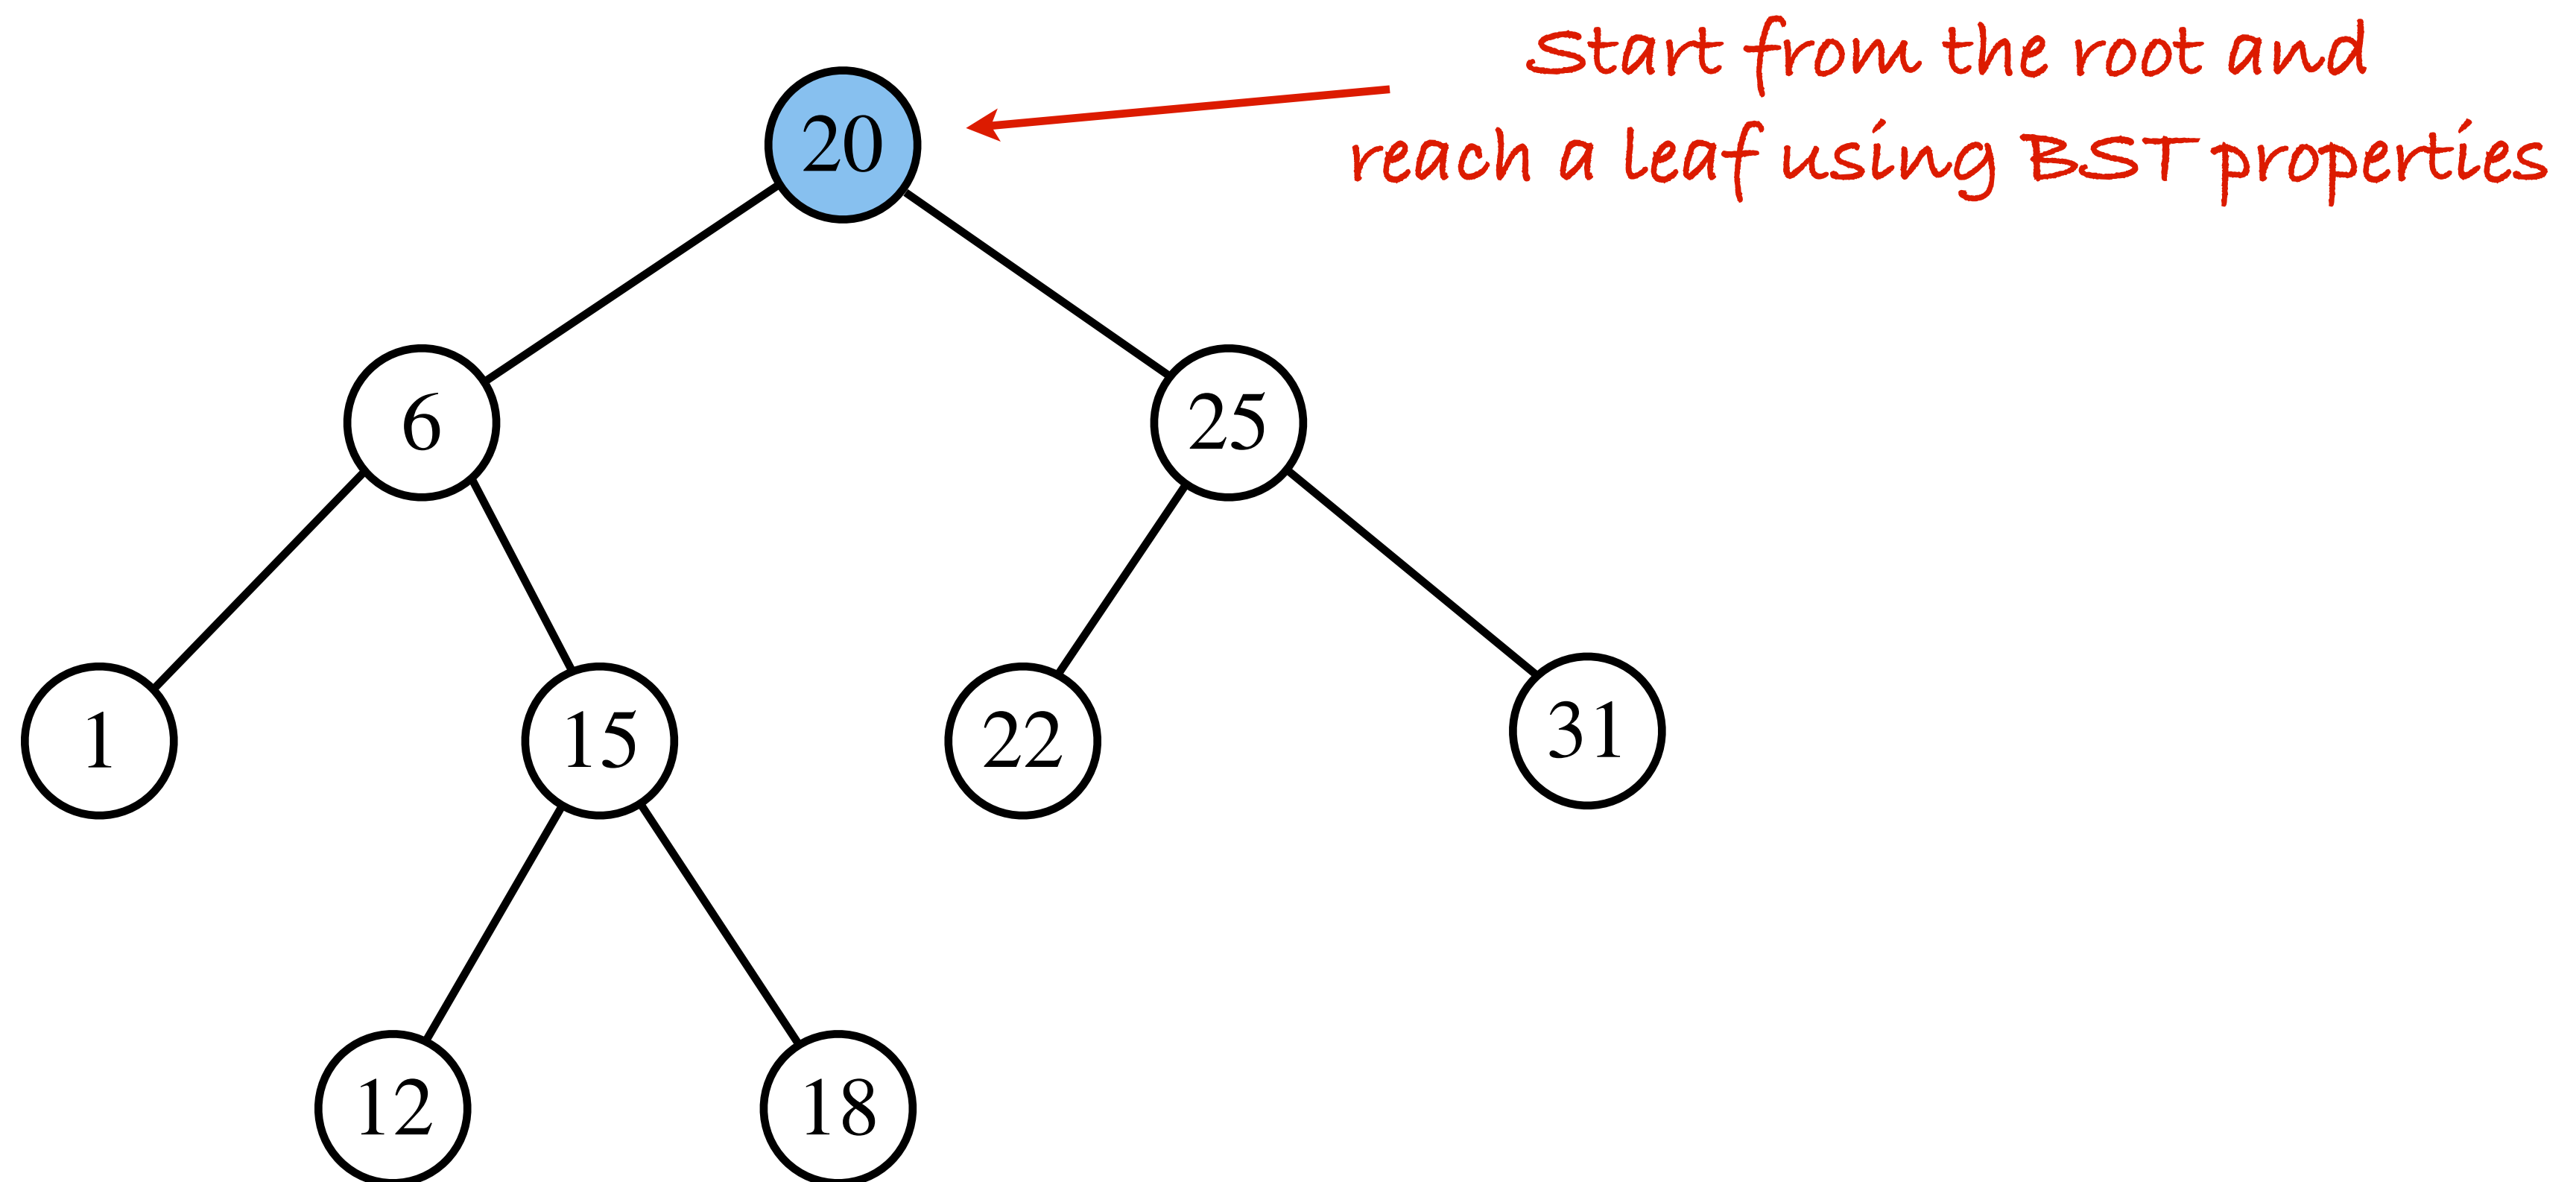
Idea: Find the correct leaf where it can be inserted.



Modifying a BST: Insertion

Example: Insert a node with 24 as key in the following BST.


Idea: Find the correct leaf where it can be inserted.



Modifying a BST: Insertion

Example: Insert a node with 24 as key in the following BST.


Idea: Find the correct leaf where it can be inserted.



Modifying a BST: Insertion

Example: Insert a node with 24 as key in the following BST.


Idea: Find the correct leaf where it can be inserted.



Modifying a BST: Insertion

Example: Insert a node with **24** as key in the following BST.


Idea: Find the correct leaf where it can be inserted.



Modifying a BST: Insertion

Example: Insert a node with **24** as key in the following BST.


Idea: Find the correct leaf where it can be inserted.



Modifying a BST: Insertion

Example: Insert a node with 24 as key in the following BST.

Idea: Find the correct leaf where it can be inserted.



*Searching for the correct leaf
and insertion takes $O(h)$ time.*

Modifying a BST: Deletion

Modifying a BST: Deletion

Deletion can be **more tricky** than Insertion.

Modifying a BST: Deletion

Deletion can be **more tricky** than Insertion.

Let z be the node we want to delete.

Modifying a BST: Deletion

Deletion can be **more tricky** than Insertion.

Let z be the node we want to delete. Then, the following cases are possible:

Modifying a BST: Deletion

Deletion can be **more tricky** than Insertion.

Let z be the node we want to delete. Then, the following cases are possible:

- **Case 1:** z has no children.

Modifying a BST: Deletion

Deletion can be **more tricky** than Insertion.

Let z be the node we want to delete. Then, the following cases are possible:

- **Case 1:** z has no children.
- **Case 2:** z has only single child.

Modifying a BST: Deletion

Deletion can be **more tricky** than Insertion.

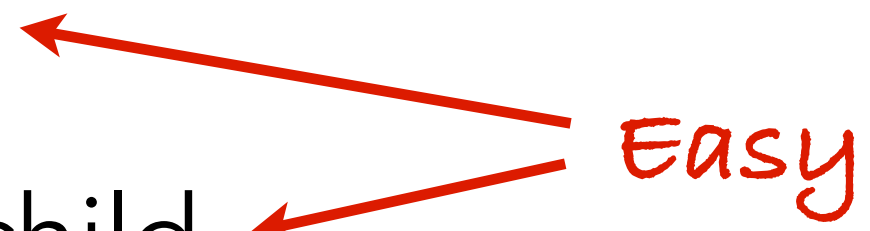
Let z be the node we want to delete. Then, the following cases are possible:

- **Case 1:** z has no children.
- **Case 2:** z has only single child.
- **Case 3:** z has two children.

Modifying a BST: Deletion

Deletion can be **more tricky** than Insertion.

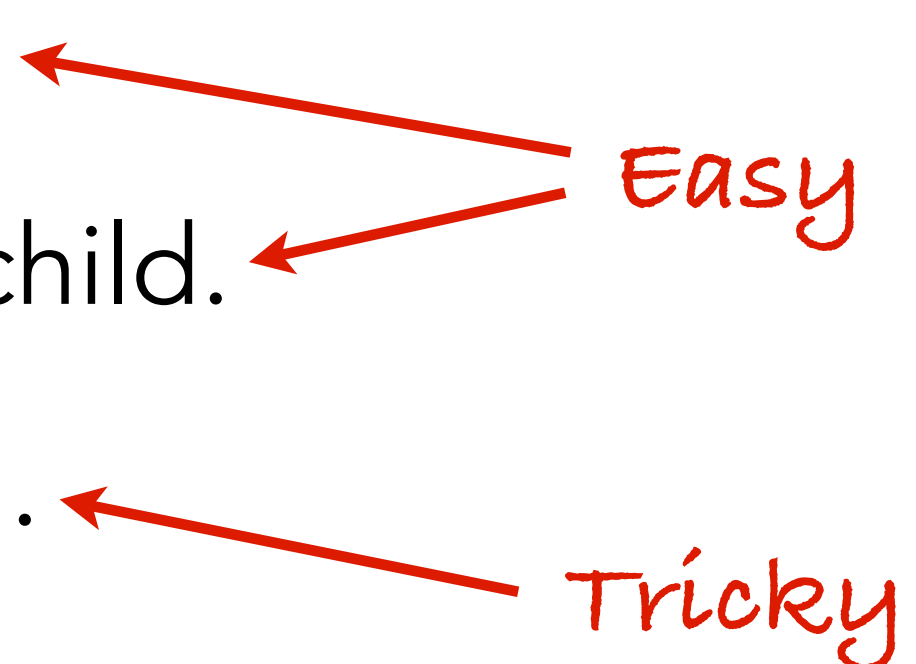
Let z be the node we want to delete. Then, the following cases are possible:

- **Case 1:** z has no children.
 - **Case 2:** z has only single child.
 - **Case 3:** z has two children.
- 
- The word "Easy" is written in red cursive script. Two red arrows originate from it: one points to "Case 1: z has no children." and the other points to "Case 2: z has only single child."

Modifying a BST: Deletion

Deletion can be **more tricky** than Insertion.

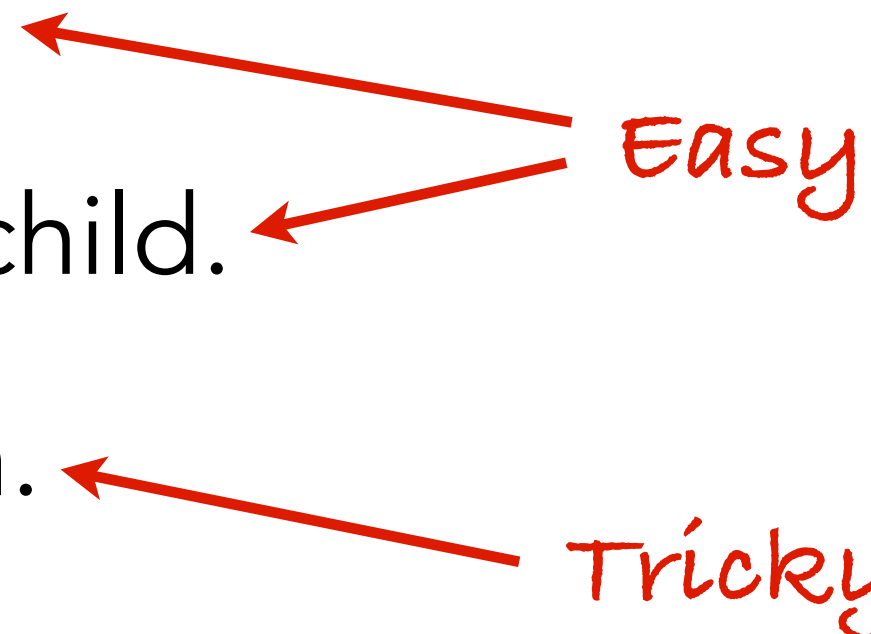
Let z be the node we want to delete. Then, the following cases are possible:

- **Case 1:** z has no children.
 - **Case 2:** z has only single child.
 - **Case 3:** z has two children.
- 
- Easy*
- Tricky*

Modifying a BST: Deletion

Deletion can be **more tricky** than Insertion.

Let z be the node we want to delete. Then, the following cases are possible:

- **Case 1:** z has no children.
 - **Case 2:** z has only single child.
 - **Case 3:** z has two children.
- 
- The diagram consists of two red arrows. One arrow originates from the word 'Easy' and points to 'Case 1'. The other arrow originates from the word 'Easy' and points to 'Case 2'. A second arrow originates from the word 'Tricky' and points to 'Case 3'.

Note: Node z is provided as input.

Modifying a BST: Deletion

Modifying a BST: Deletion


Case 1: z has no children.

Modifying a BST: Deletion

Case 1: z has no children. (WLOG assume z is a right child.)


Modifying a BST: Deletion

Case 1: z has no children. (WLOG assume z is a right child.)




Modifying a BST: Deletion

Case 1: z has no children. (WLOG assume z is a right child.)





Modifying a BST: Deletion

Case 1: z has no children. (WLOG assume z is a right child.)



Modifying a BST: Deletion

Case 1: z has no children. (WLOG assume z is a right child.)



Make q 's right child NIL


Modifying a BST: Deletion

Modifying a BST: Deletion

Case 2: z has one child. (WLOG assume z is a right child.)


Modifying a BST: Deletion

Case 2: z has one child. (WLOG assume z is a right child.)




Modifying a BST: Deletion

Case 2: z has one child. (WLOG assume z is a right child.)




Modifying a BST: Deletion

Case 2: z has one child. (WLOG assume z is a right child.)



Modifying a BST: Deletion

Case 2: z has one child. (WLOG assume z is a right child.)




Modifying a BST: Deletion

Case 2: z has one child. (WLOG assume z is a right child.)


Modifying a BST: Deletion

Case 2: z has one child. (WLOG assume z is a right child.)




Modifying a BST: Deletion

Case 2: z has one child. (WLOG assume z is a right child.)



Modifying a BST: Deletion

Case 2: z has one child. (WLOG assume z is a right child.)




Modifying a BST: Deletion

Modifying a BST: Deletion

Case 3: z has two children. (WLOG assume z is a right child.)


Modifying a BST: Deletion

Case 3: z has two children. (WLOG assume z is a right child.)




Modifying a BST: Deletion

Case 3: z has two children. (WLOG assume z is a right child.)



Modifying a BST: Deletion


Case 3: z has two children. (WLOG assume z is a right child.)



?

Modifying a BST: Deletion

Case 3: z has two children. (WLOG assume z is a right child.)



Two sub-cases:


- r has no left child.
- r has a left child.

Modifying a BST: Deletion

Case 3a: z has two children where its right child has no left child.


Modifying a BST: Deletion

Case 3a: z has two children where its right child has no left child.




Modifying a BST: Deletion

Case 3a: z has two children where its right child has no left child.




Modifying a BST: Deletion

Case 3a: z has two children where its right child has no left child.



Modifying a BST: Deletion

Case 3a: z has two children where its right child has no left child.




Modifying a BST: Deletion

Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.

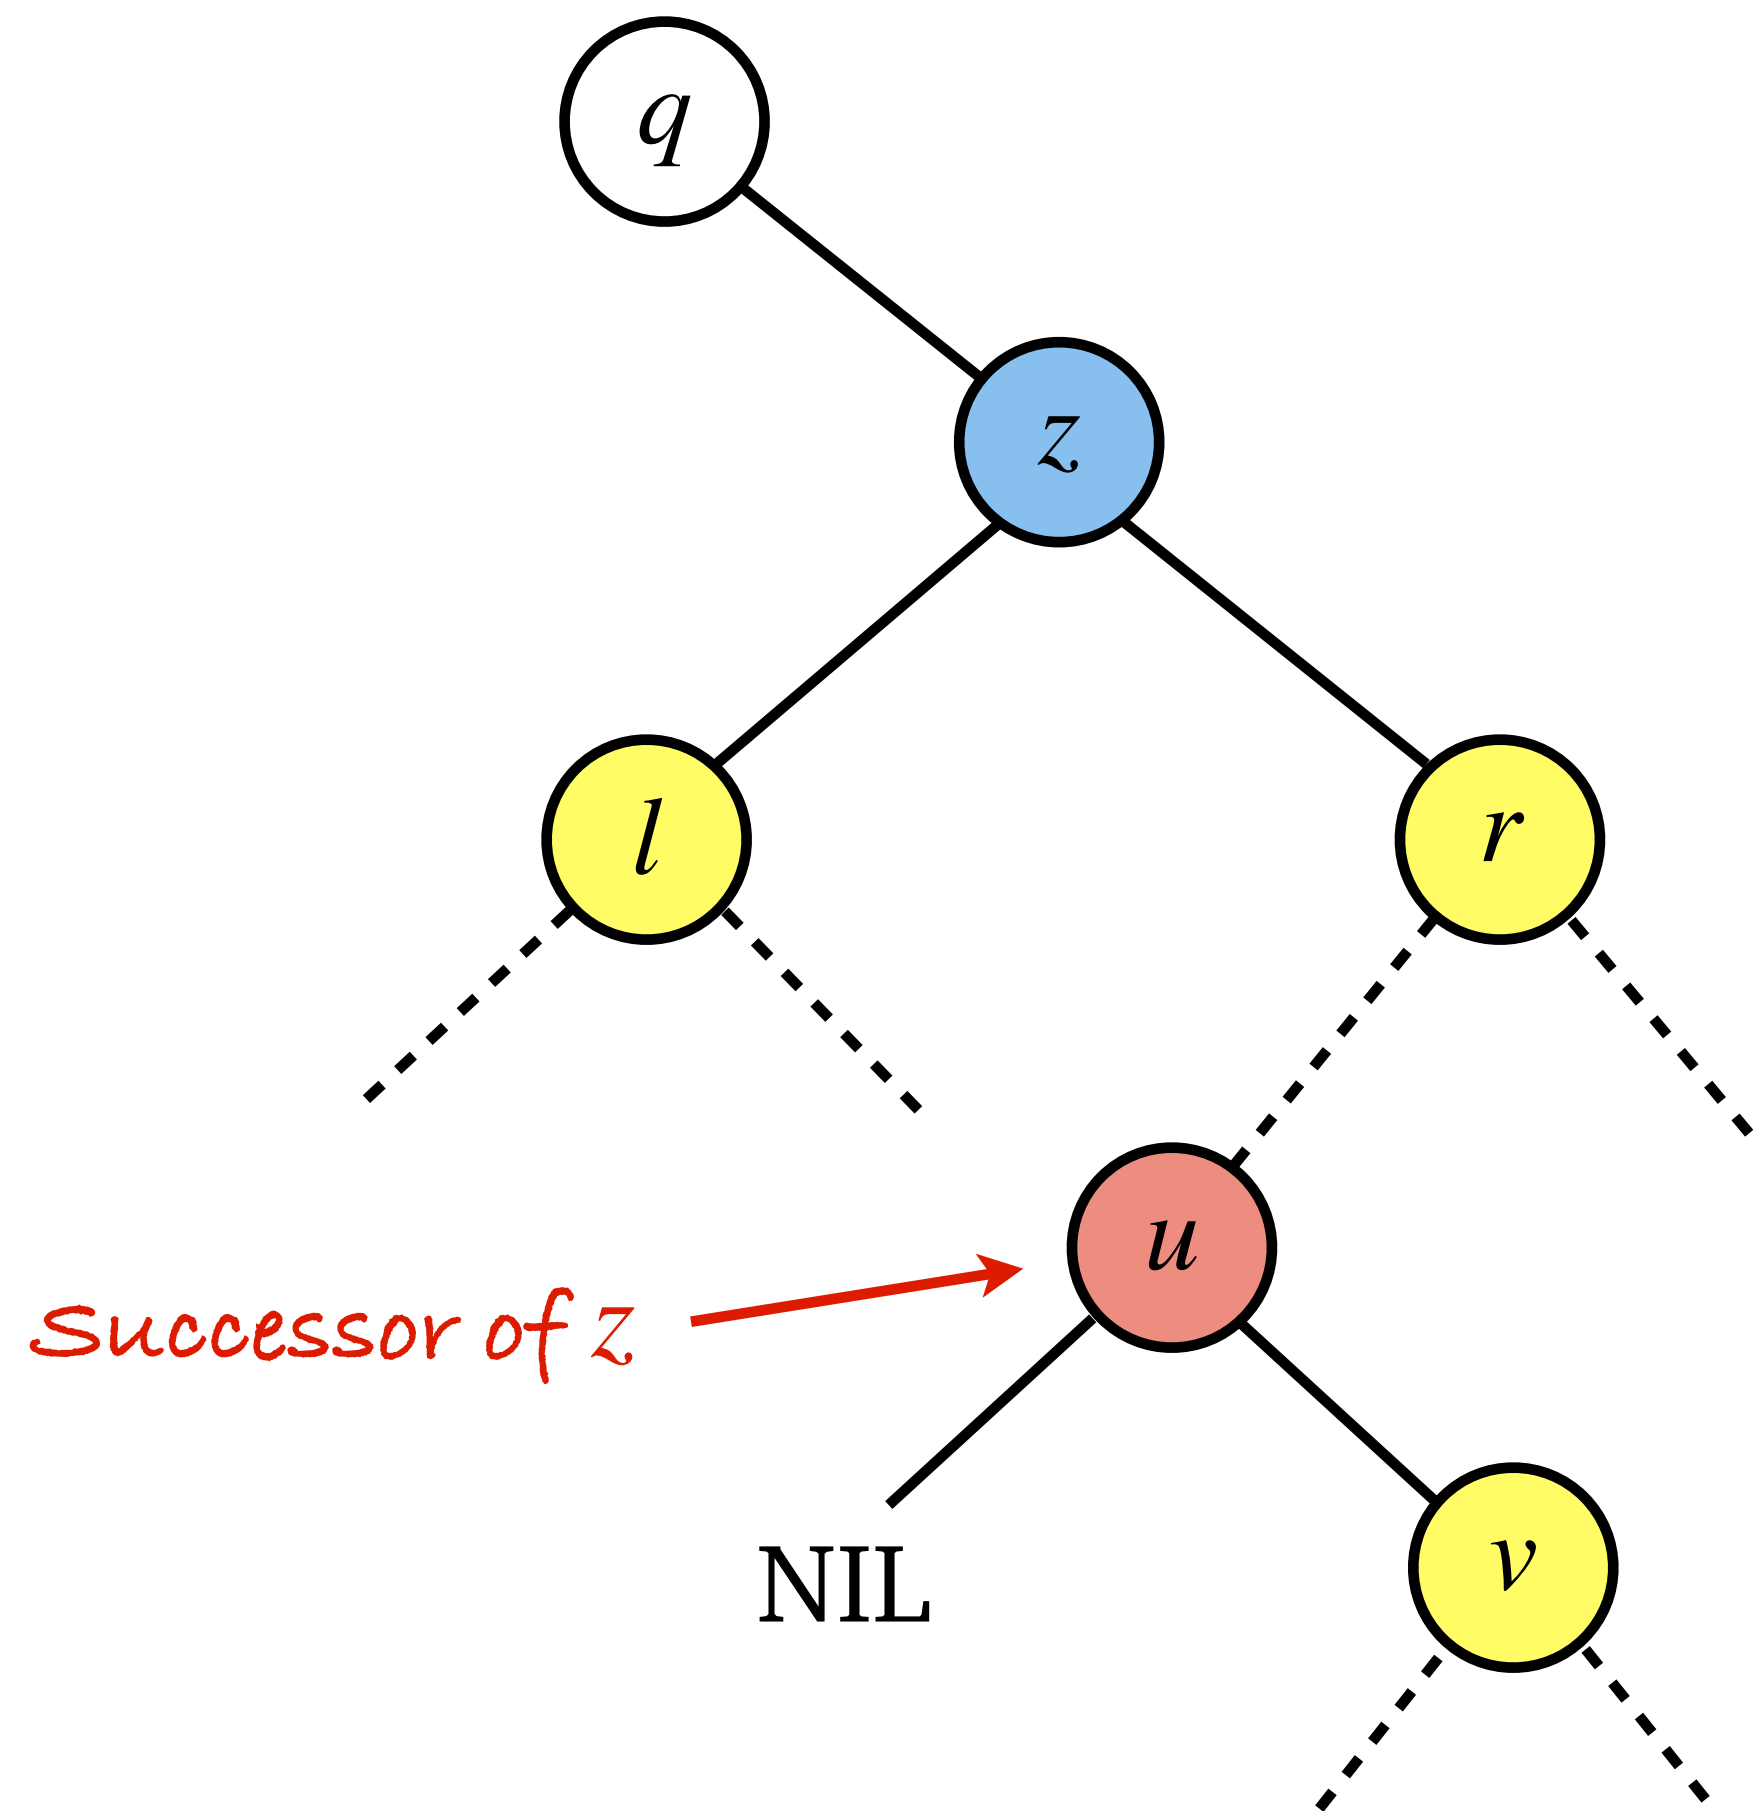
Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.




Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.



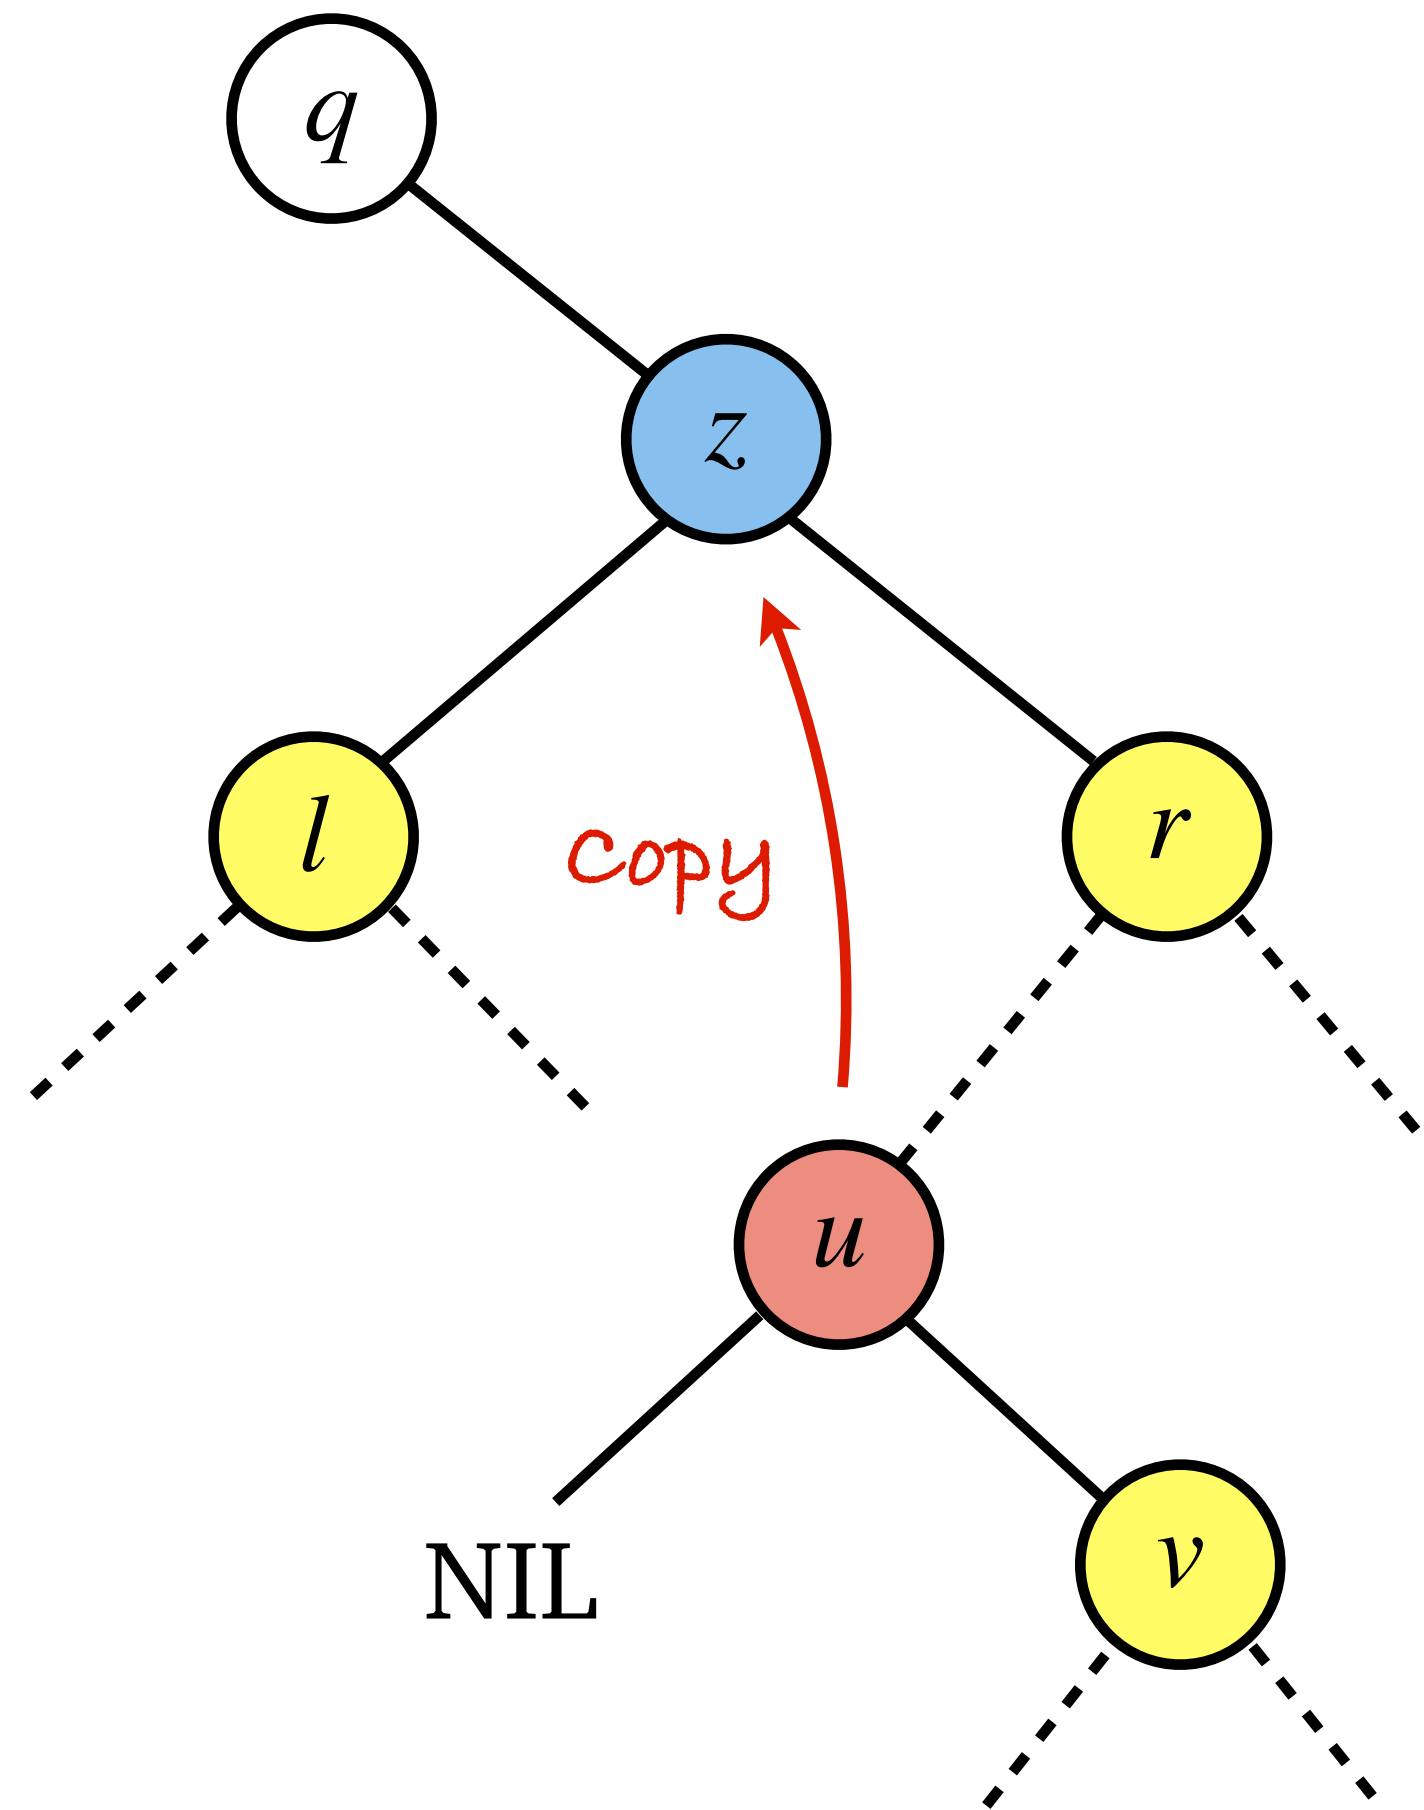
Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.




Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.




Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.




Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.




Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.




Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.




Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.



Modifying a BST: Deletion

Case 3b: z has two children where its right child has a left child.



Are BSTs Good Enough?

Are BSTs Good Enough?

BSTs can perform **Insert**, **Delete**, **Search**, etc., in $O(h)$ time.

Are BSTs Good Enough?


BSTs can perform **Insert**, **Delete**, **Search**, etc., in $O(h)$ time.

But,

Are BSTs Good Enough?

BSTs can perform **Insert**, **Delete**, **Search**, etc., in $O(h)$ time.


But,



Are BSTs Good Enough?

BSTs can perform **Insert**, **Delete**, **Search**, etc., in $O(h)$ time.

But,




Best case: $h = O(\log n)$


Are BSTs Good Enough?

BSTs can perform **Insert**, **Delete**, **Search**, etc., in $O(h)$ time.

But,




Best case: $h = O(\log n)$




Are BSTs Good Enough?

BSTs can perform **Insert**, **Delete**, **Search**, etc., in $O(h)$ time.

But,



Best case: $h = O(\log n)$




Worst case: $h = O(n)$


Are BSTs Good Enough?

BSTs can perform **Insert**, **Delete**, **Search**, etc., in $O(h)$ time.

But,




Best case: $h = O(\log n)$




Worst case: $h = O(n)$

Red-black (RB) trees provide a way keep height $O(\log n)$.


RB-Trees: How do they look like?



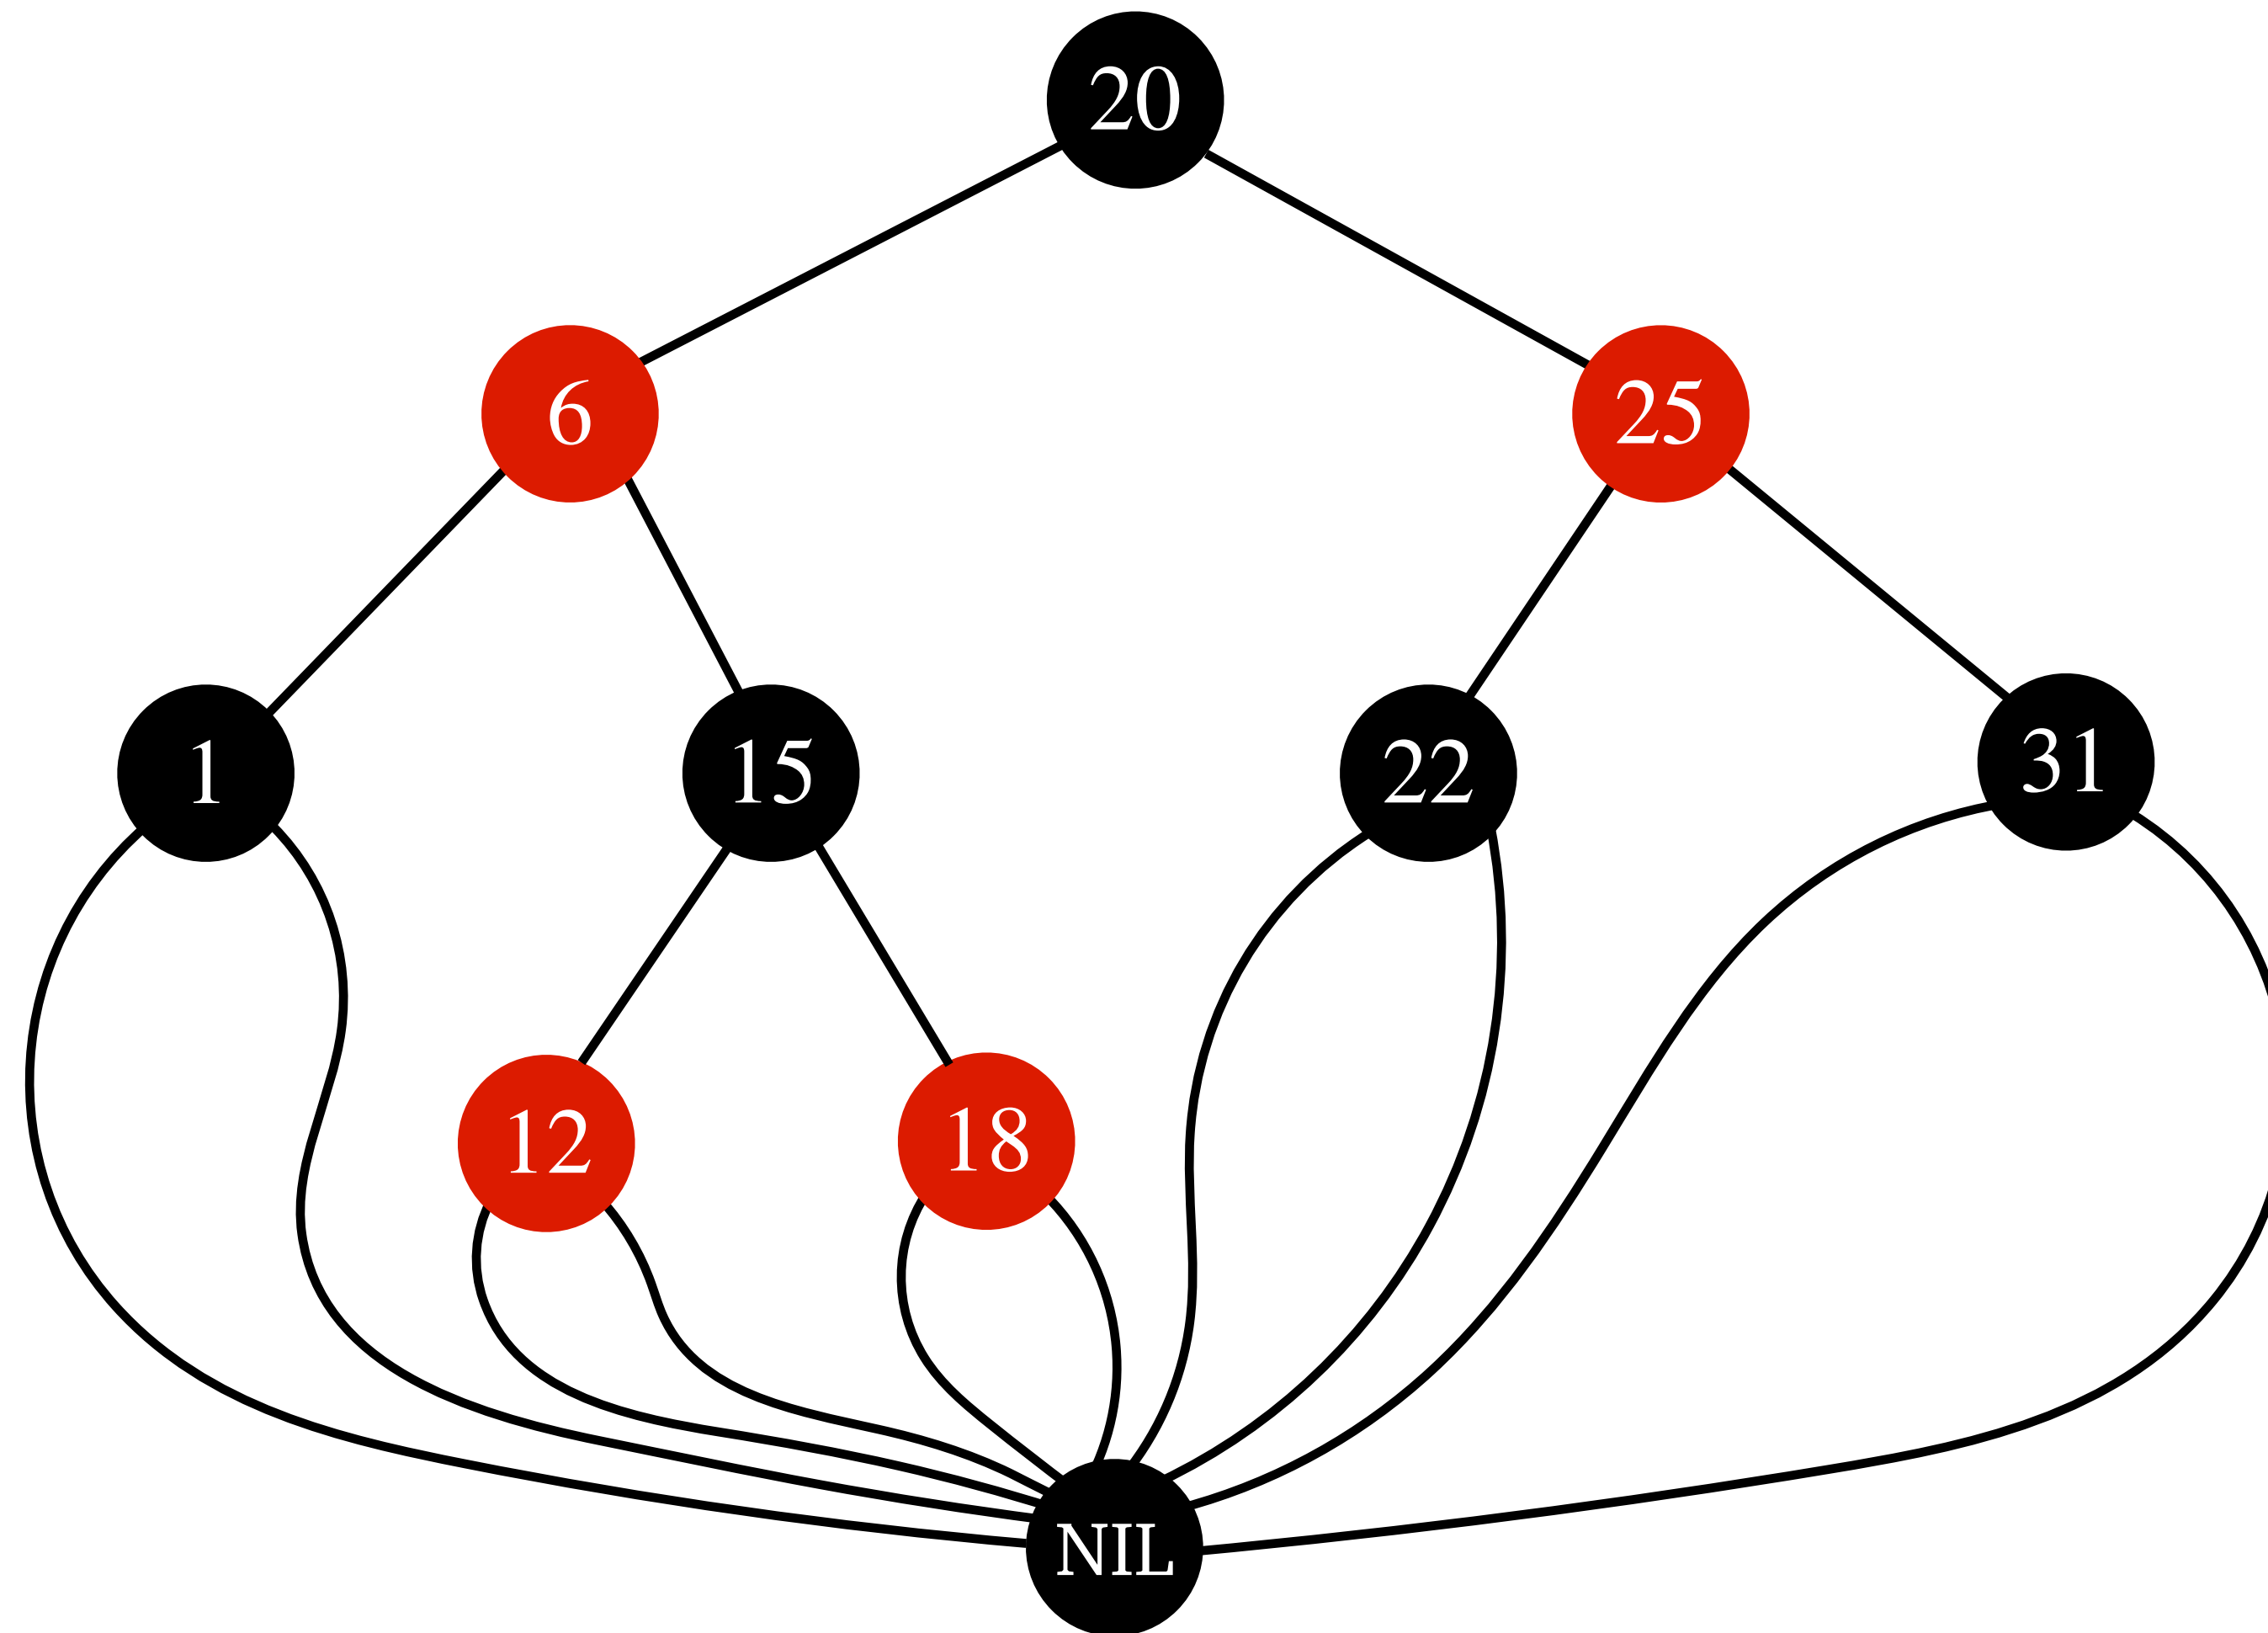
RB-Trees: How do they look like?



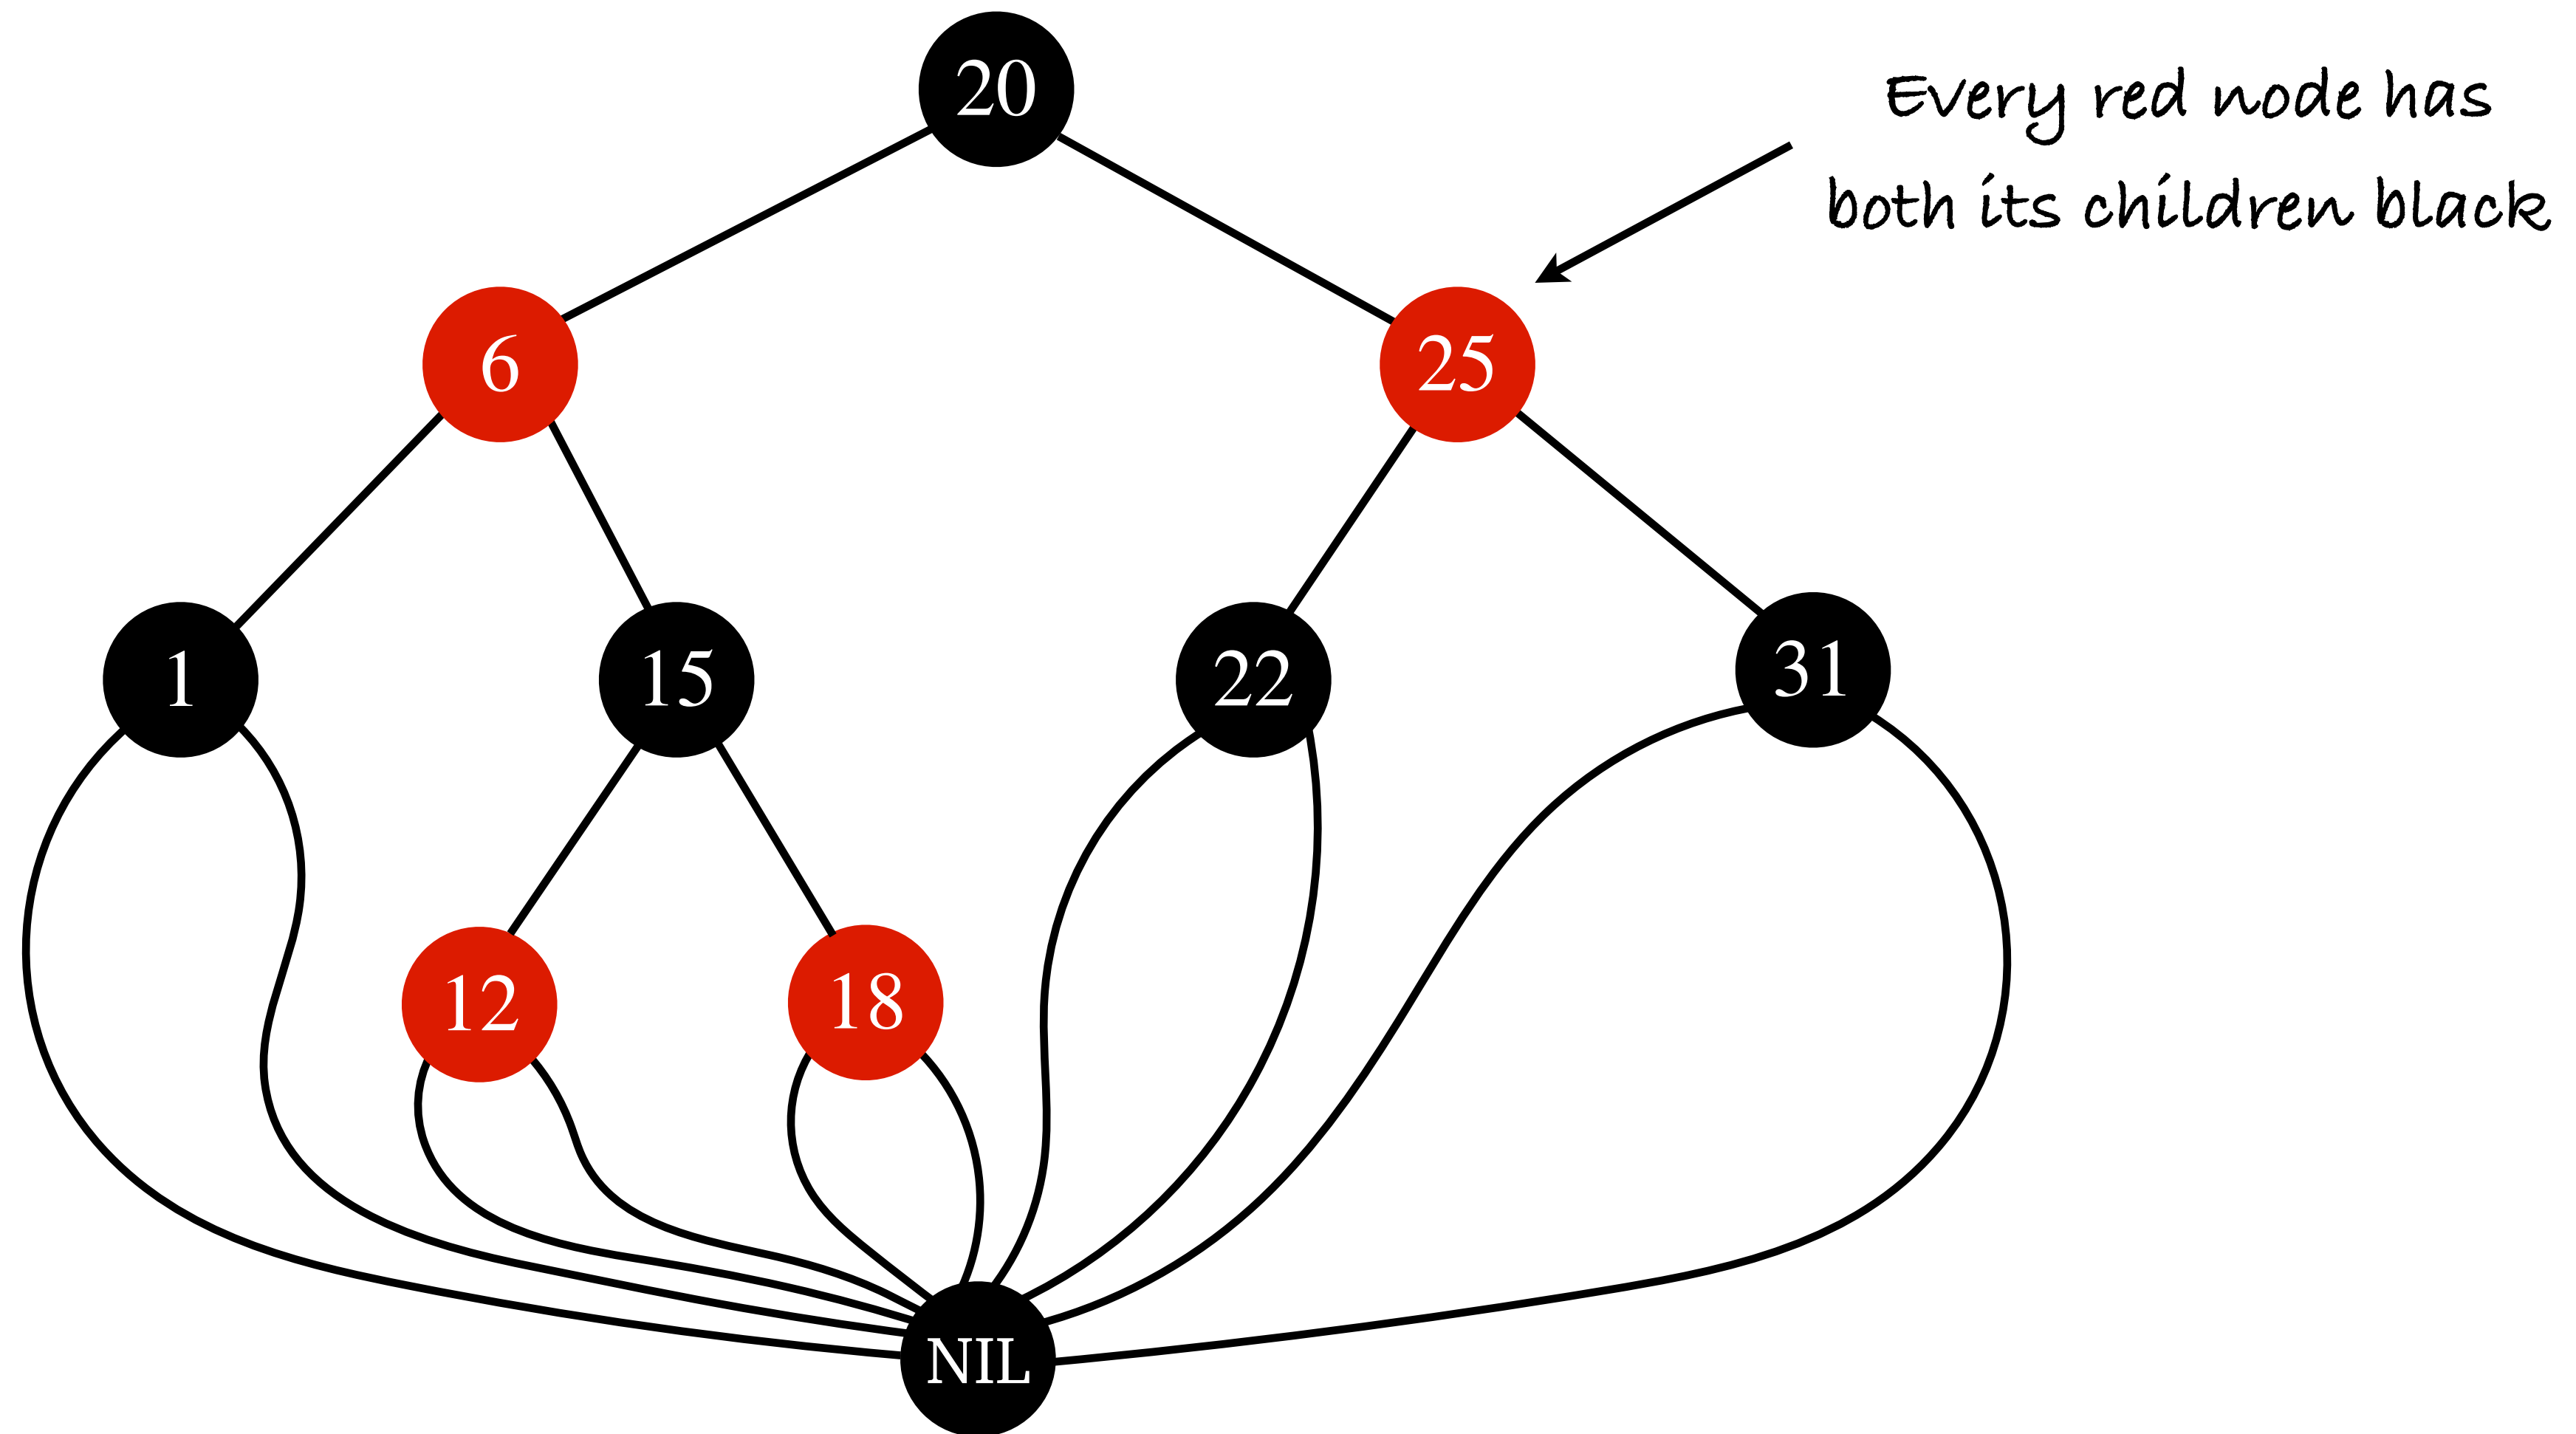
RB-Trees: How do they look like?



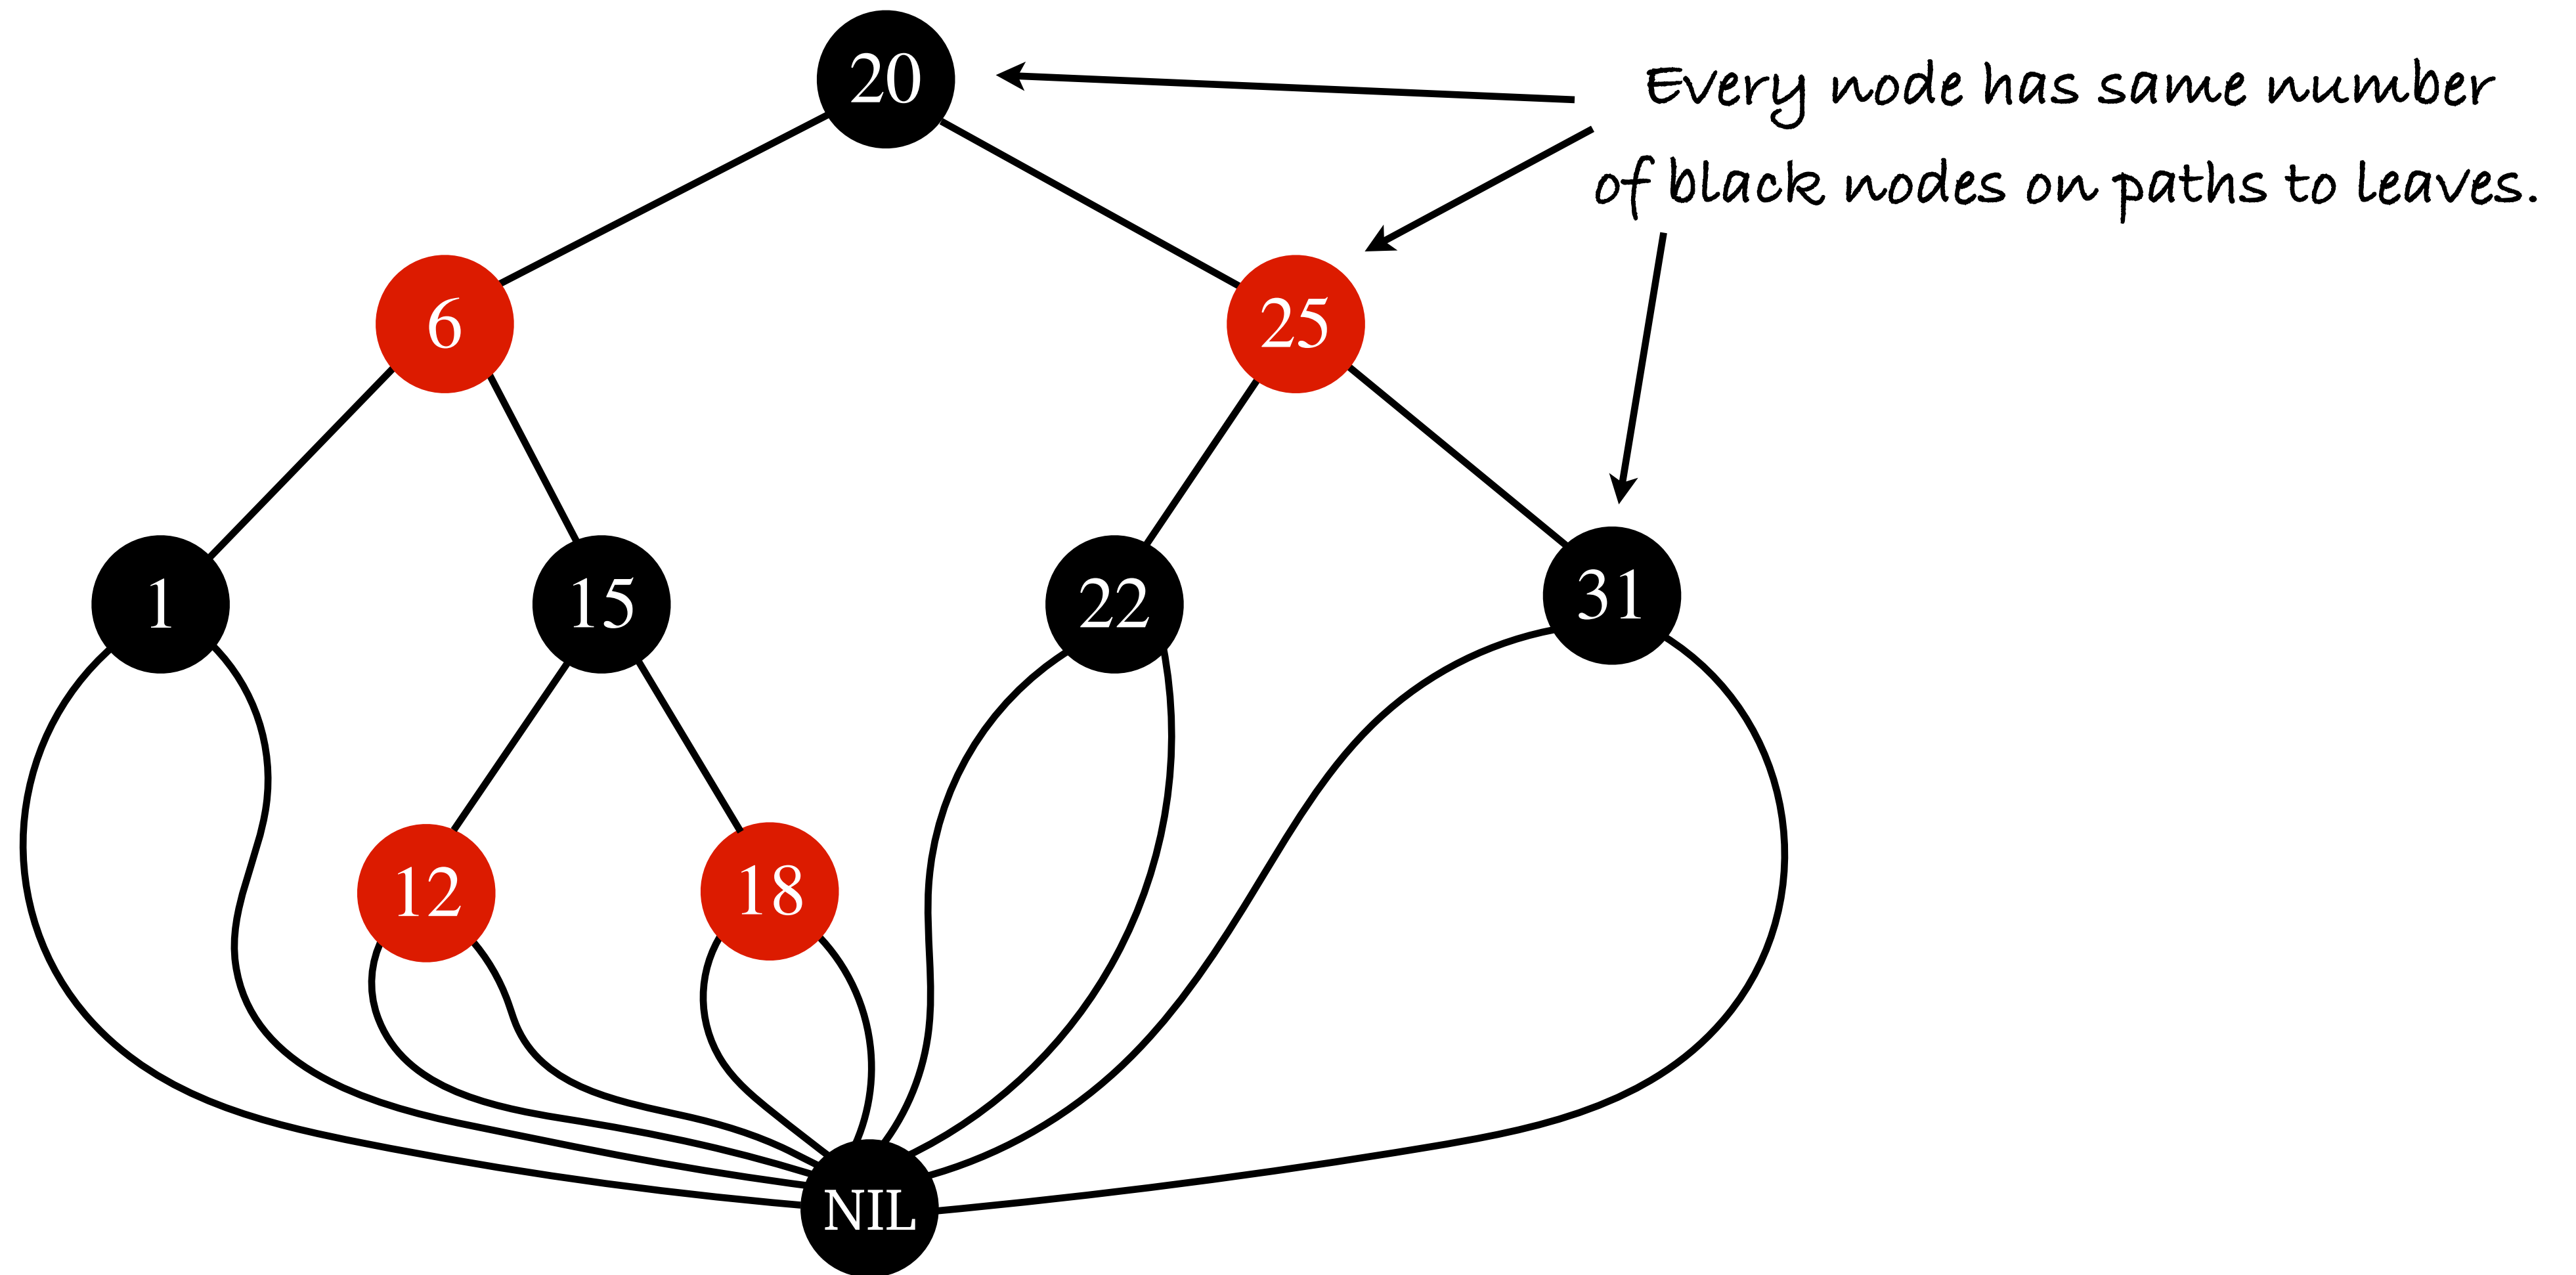
RB-Trees: How do they look like?



RB-Trees: How do they look like?



RB-Trees: How do they look like?



RB-Trees: Formal Description

RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.

RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.

RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.

RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.
- Every internal (non-leaf) node has two children.

RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.
- Every internal (non-leaf) node has two children.
- If a node is **red**, then both its children are **black**.

RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.
- Every internal (non-leaf) node has two children.
- If a node is **red**, then both its children are **black**.
- For every node, all the paths from the node to descendent leaves contain the

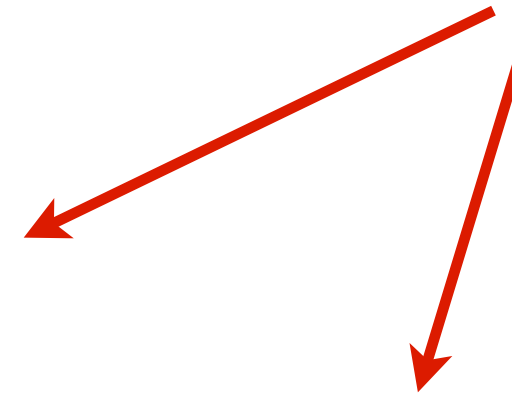
RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.
- Every internal (non-leaf) node has two children.
- If a node is **red**, then both its children are **black**.
- For every node, all the paths from the node to descendent leaves contain the same number of **black** nodes.

RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:


- Every node has a colour either **red** or **black**.
 - Root is **black**.
 - Leaf nodes are NIL nodes which are **black** in colour.
 - Every internal (non-leaf) node has two children.
 - If a node is **red**, then both its children are **black**.
 - For every node, all the paths from the node to descendent leaves contain the same number of **black** nodes.
- 

RB-Trees: Formal Description

Root ●

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.
- Every internal (non-leaf) node has two children.
- If a node is **red**, then both its children are **black**.
- For every node, all the paths from the node to descendent leaves contain the same number of **black** nodes.



RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.
- Every internal (non-leaf) node has two children.
- If a node is **red**, then both its children are **black**.
- For every node, all the paths from the node to descendent leaves contain the same number of **black** nodes.

Root




Leaf 1

Leaf 2

RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:


- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.
- Every internal (non-leaf) node has two children.
- If a node is **red**, then both its children are **black**.
- For every node, all the paths from the node to descendent leaves contain the same number of **black** nodes.



RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.
- Every internal (non-leaf) node has two children.
- If a node is **red**, then both its children are **black**.
- For every node, all the paths from the node to descendent leaves contain the same number of **black** nodes.




RB-Trees: Formal Description

RB-trees are BSTs which satisfy the following properties:

- Every node has a colour either **red** or **black**.
- Root is **black**.
- Leaf nodes are NIL nodes which are **black** in colour.
- Every internal node has two children.
- If a node is **red**, then both its children are **black**.
- For every node, all the paths from the node to descendent leaves contain the same number of **black** nodes.

*Both these properties ensure that
no path from root to a leaf is more than
twice as long as any other.*



RB-Trees: Black Height

RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf,

RB-Trees: Black Height

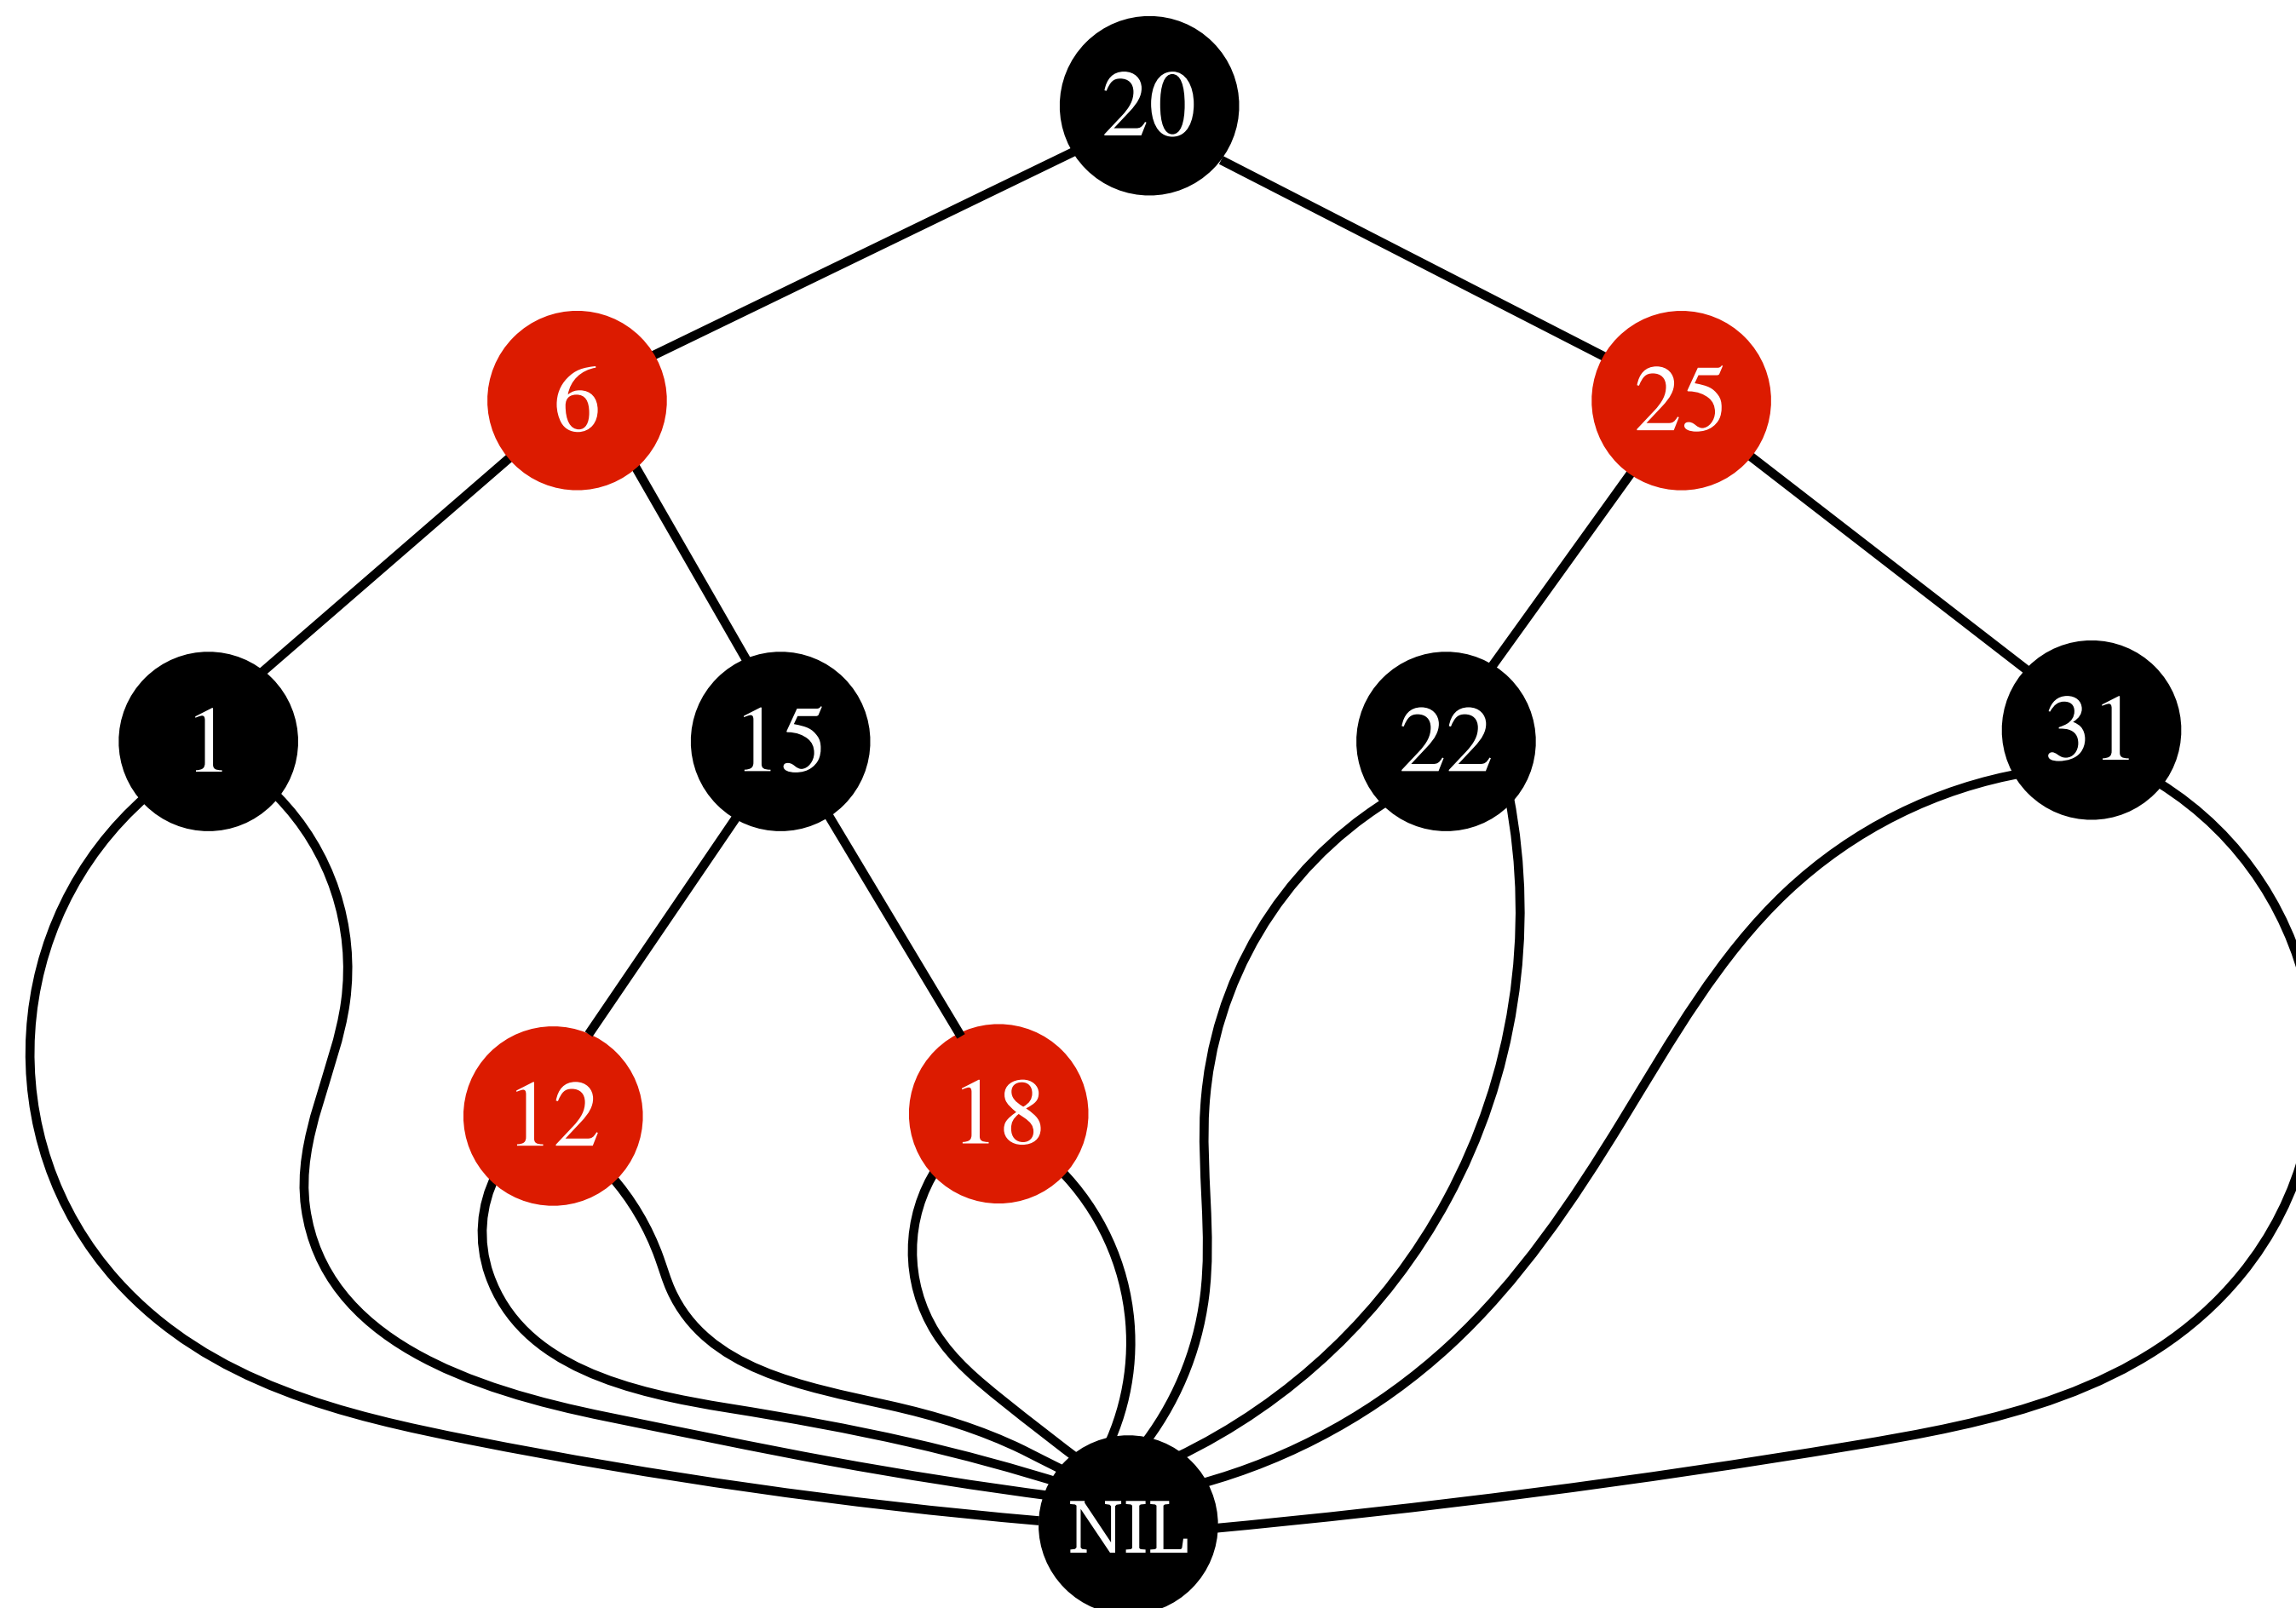
Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x ,

RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.

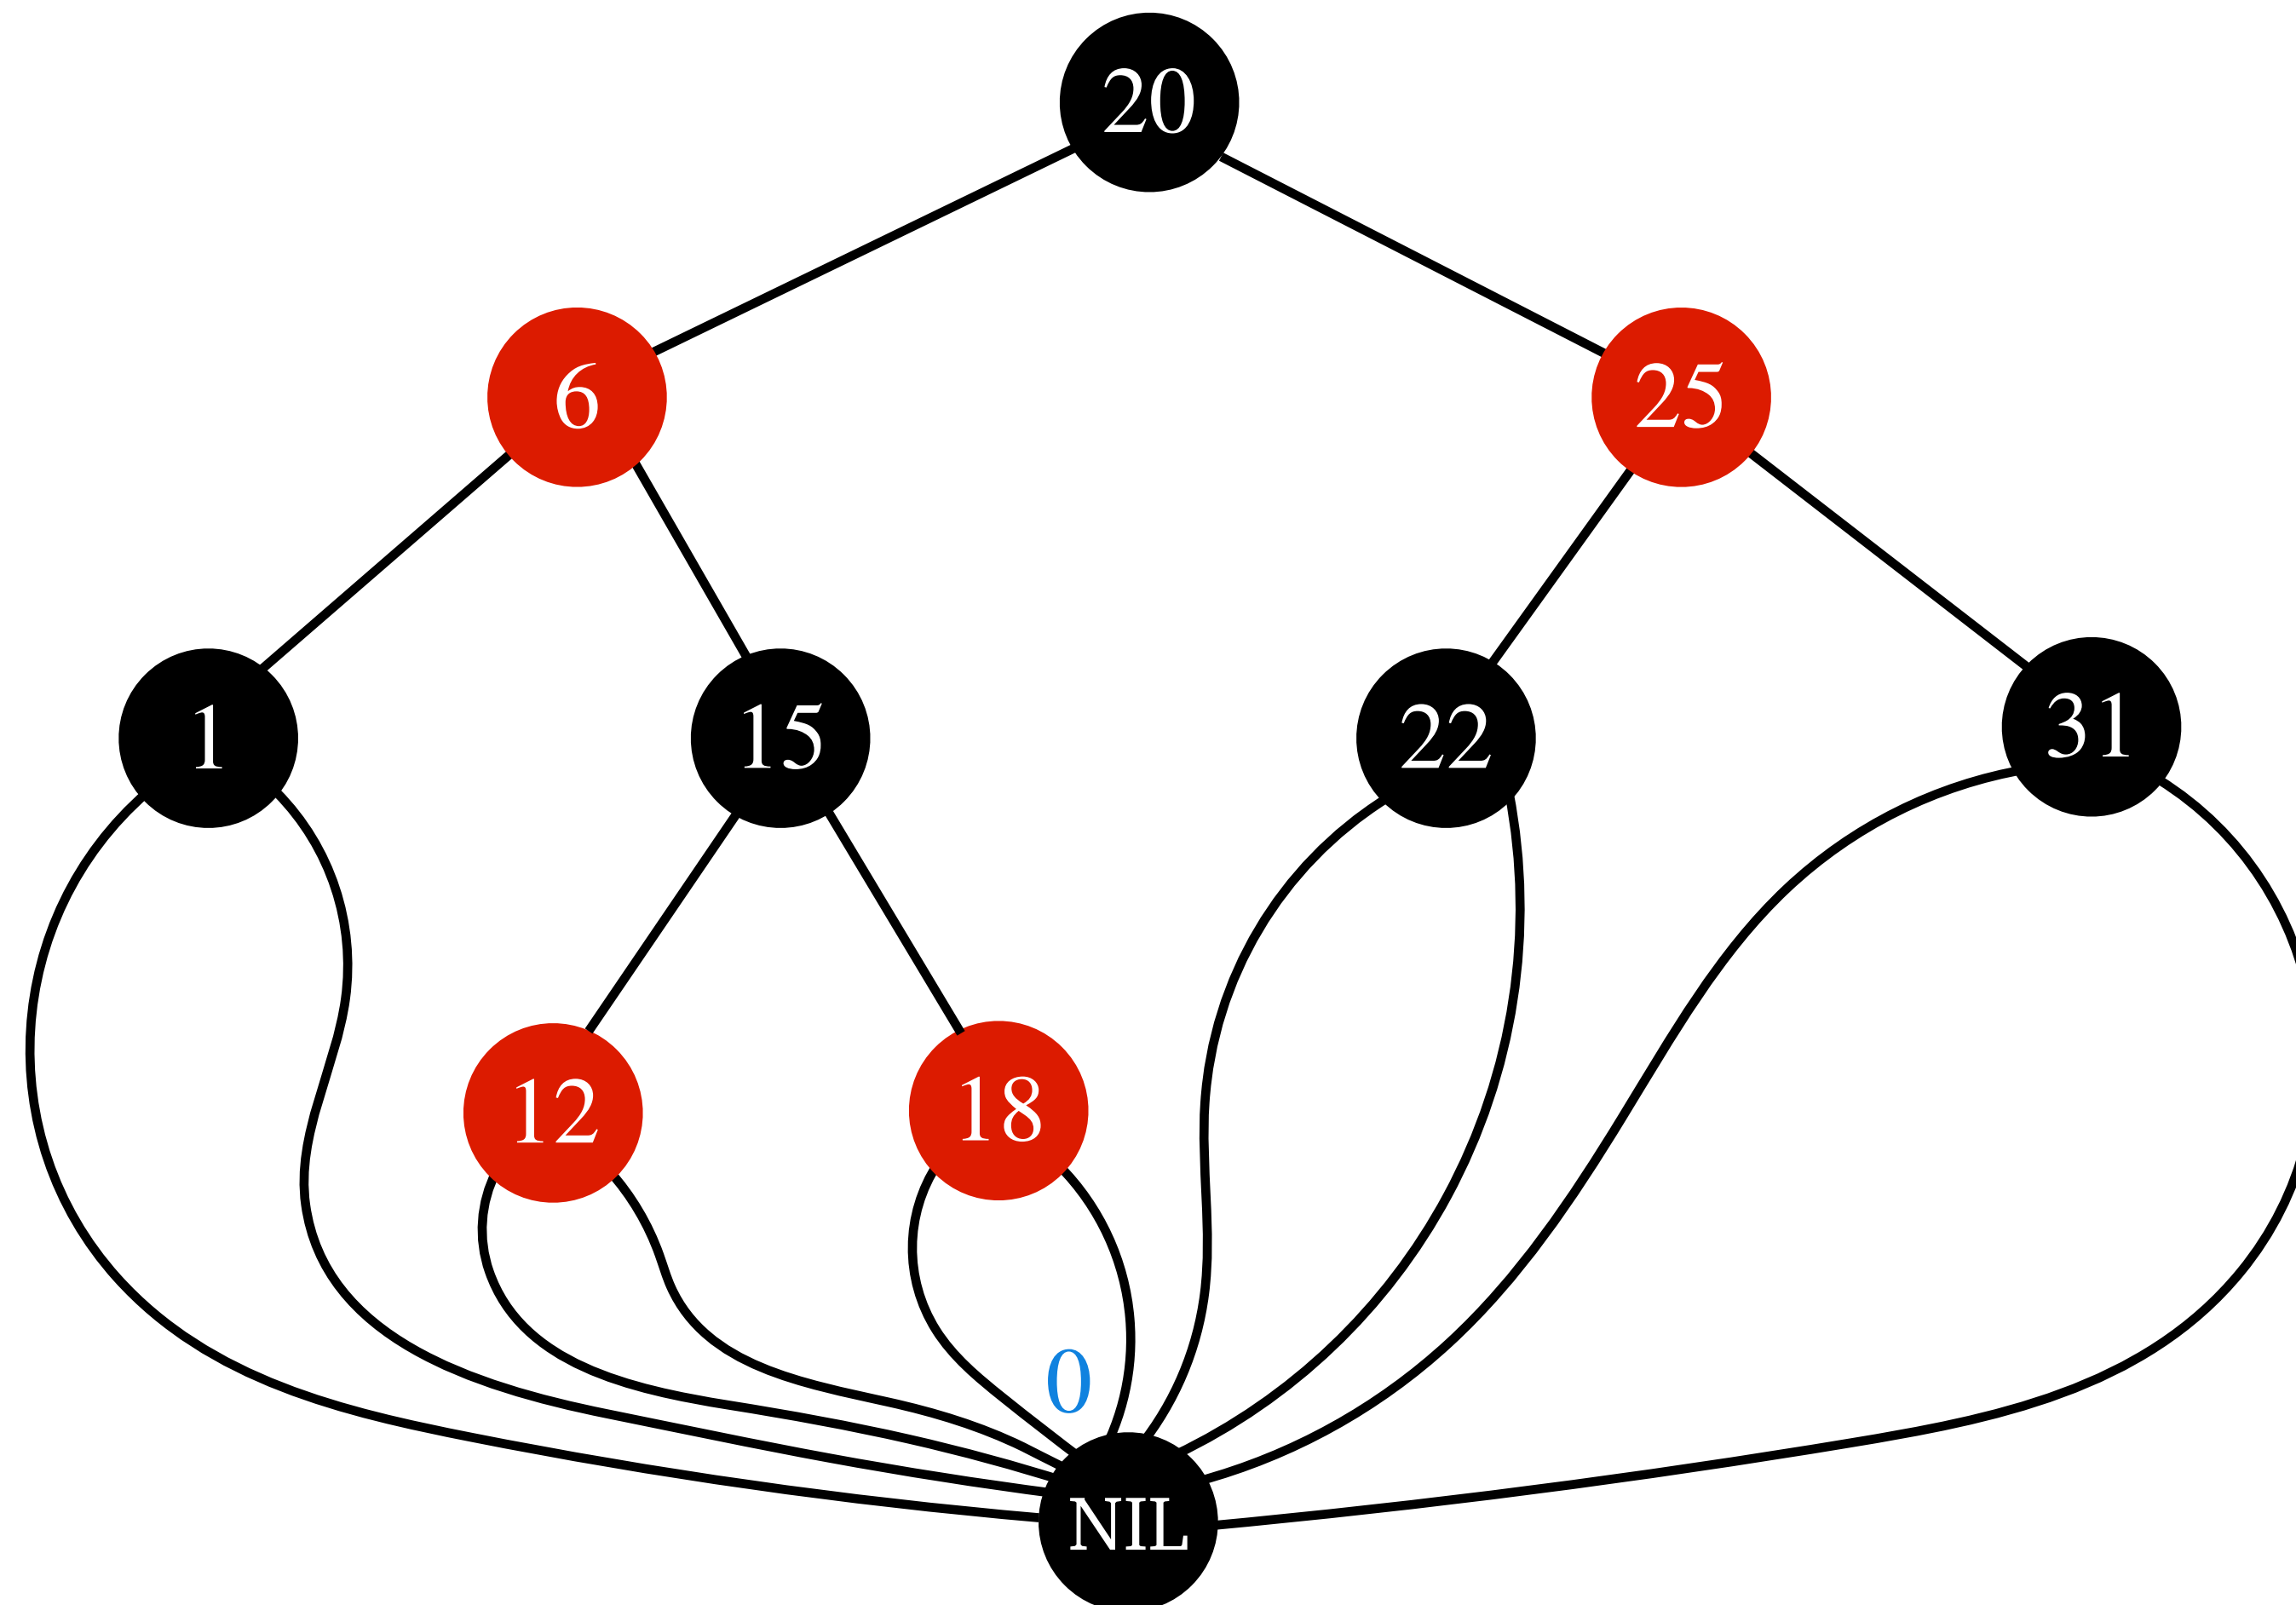
RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



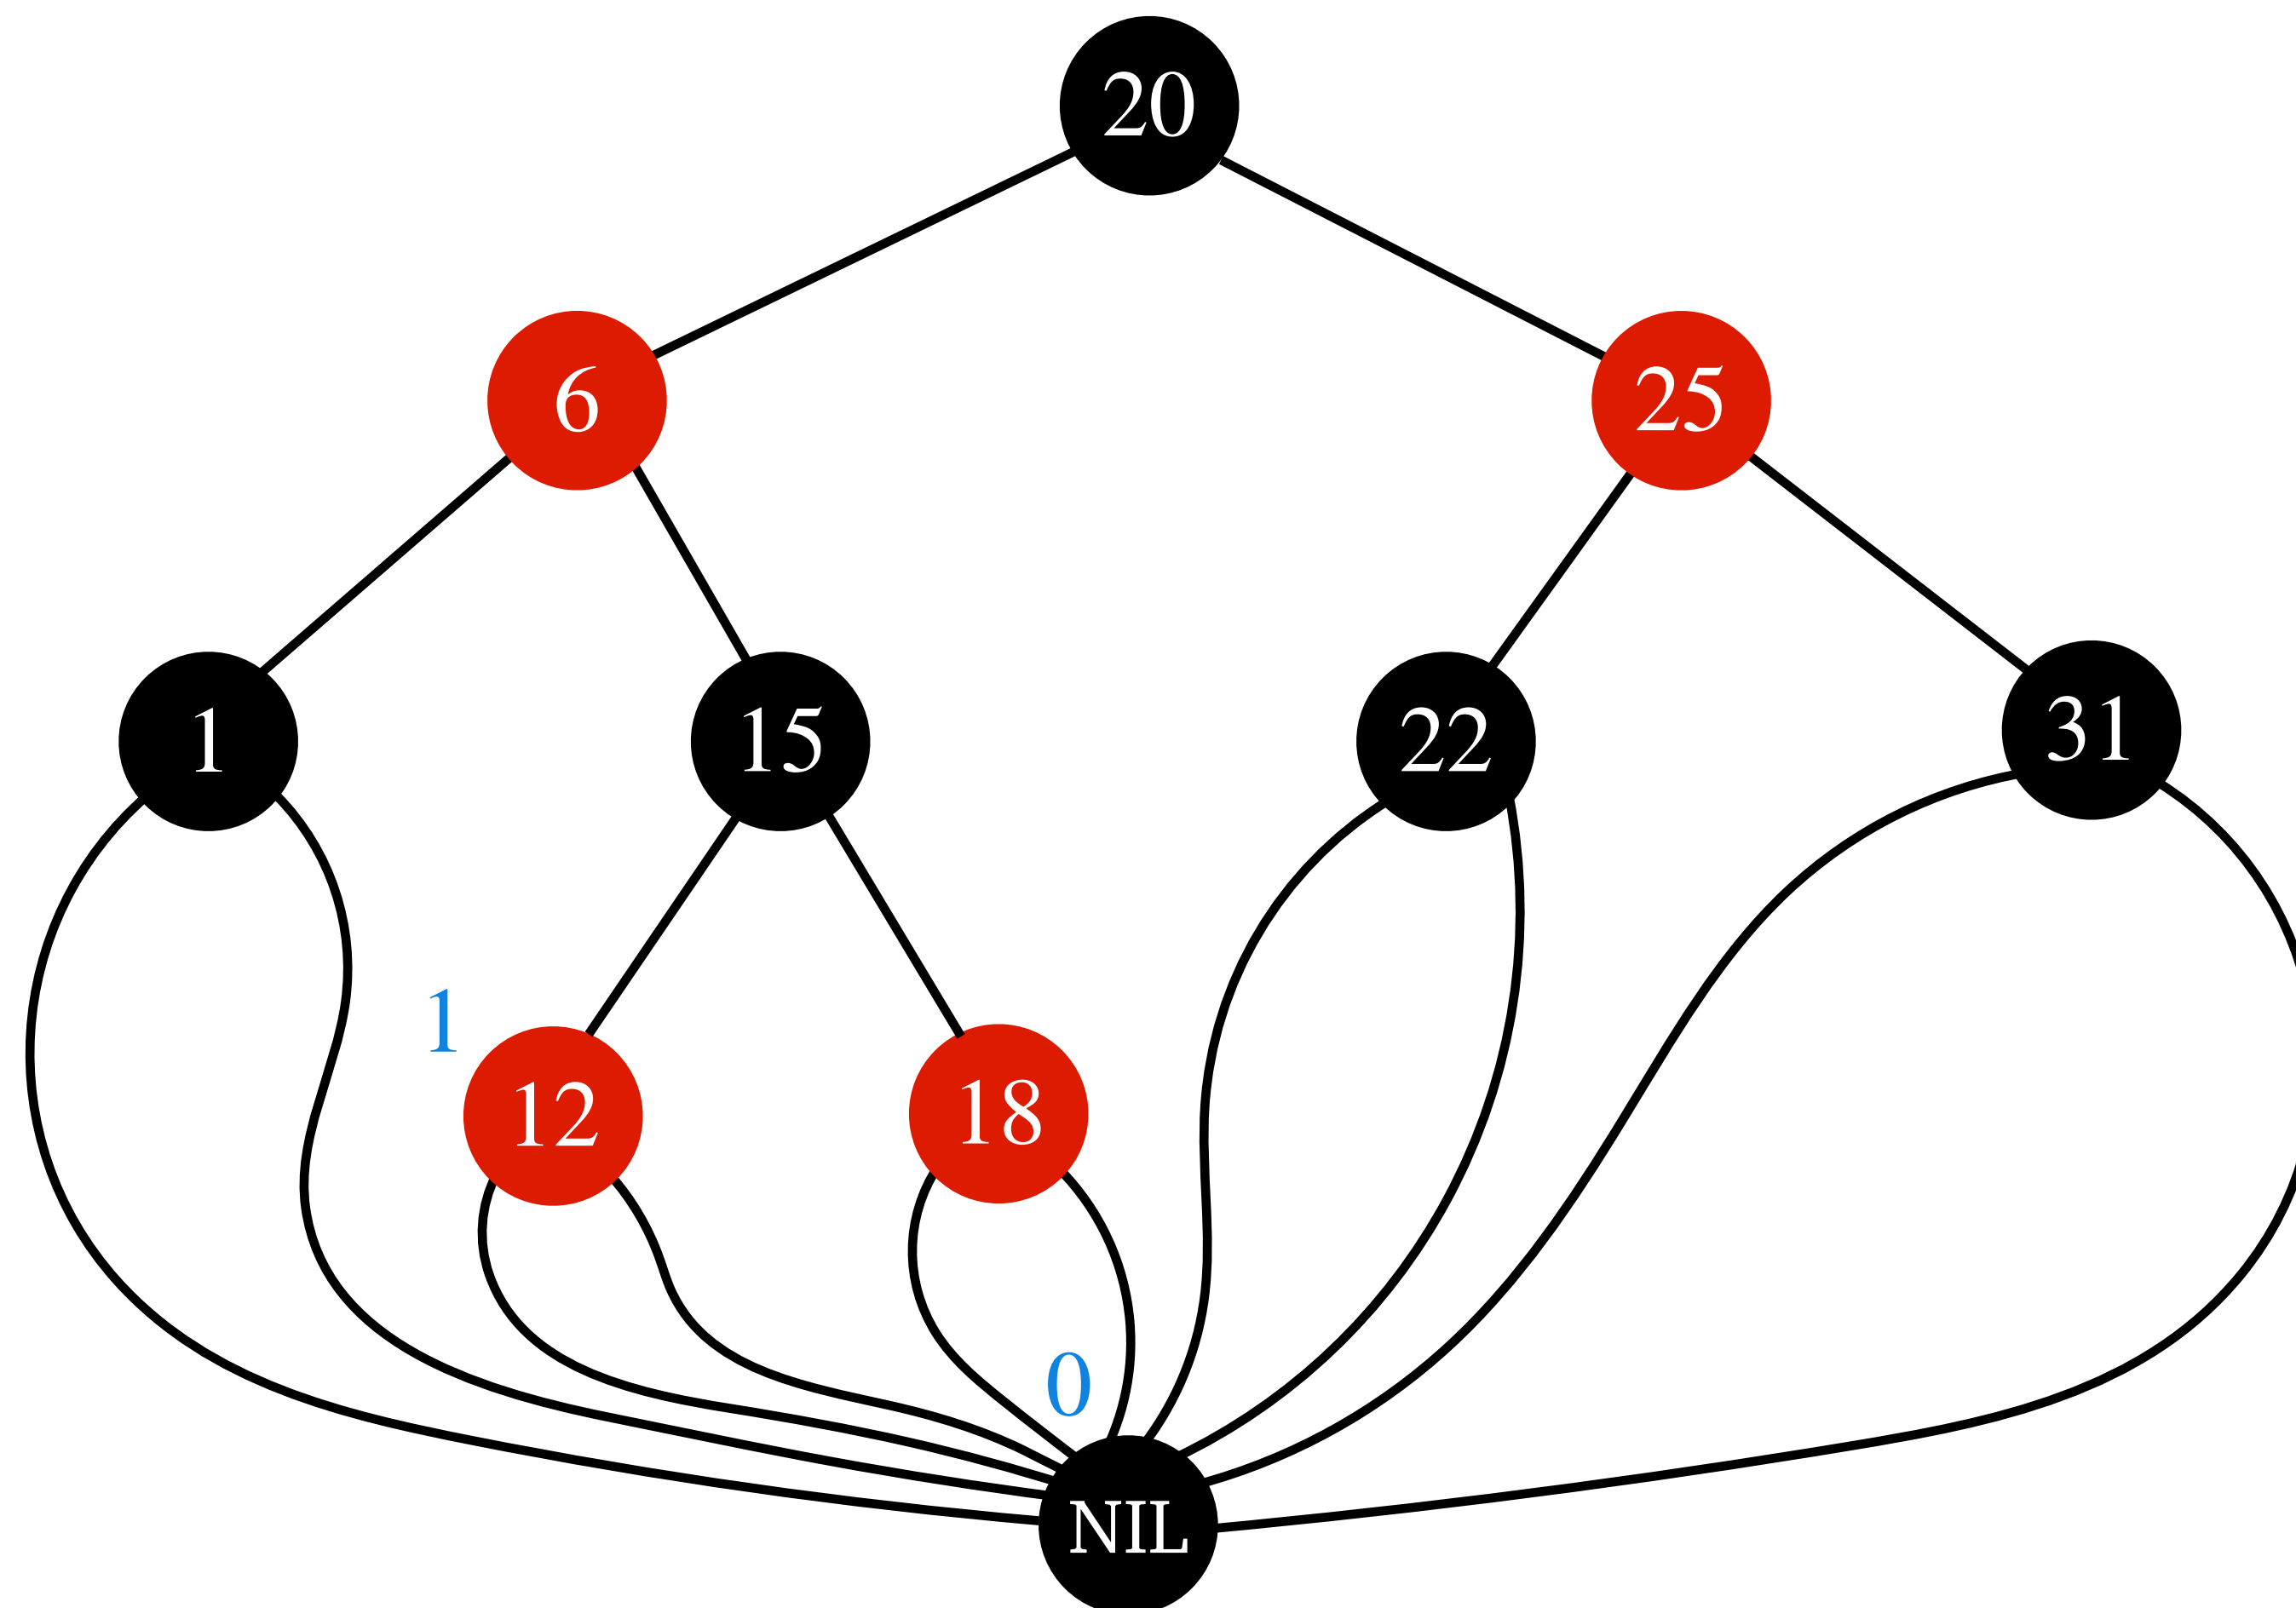
RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



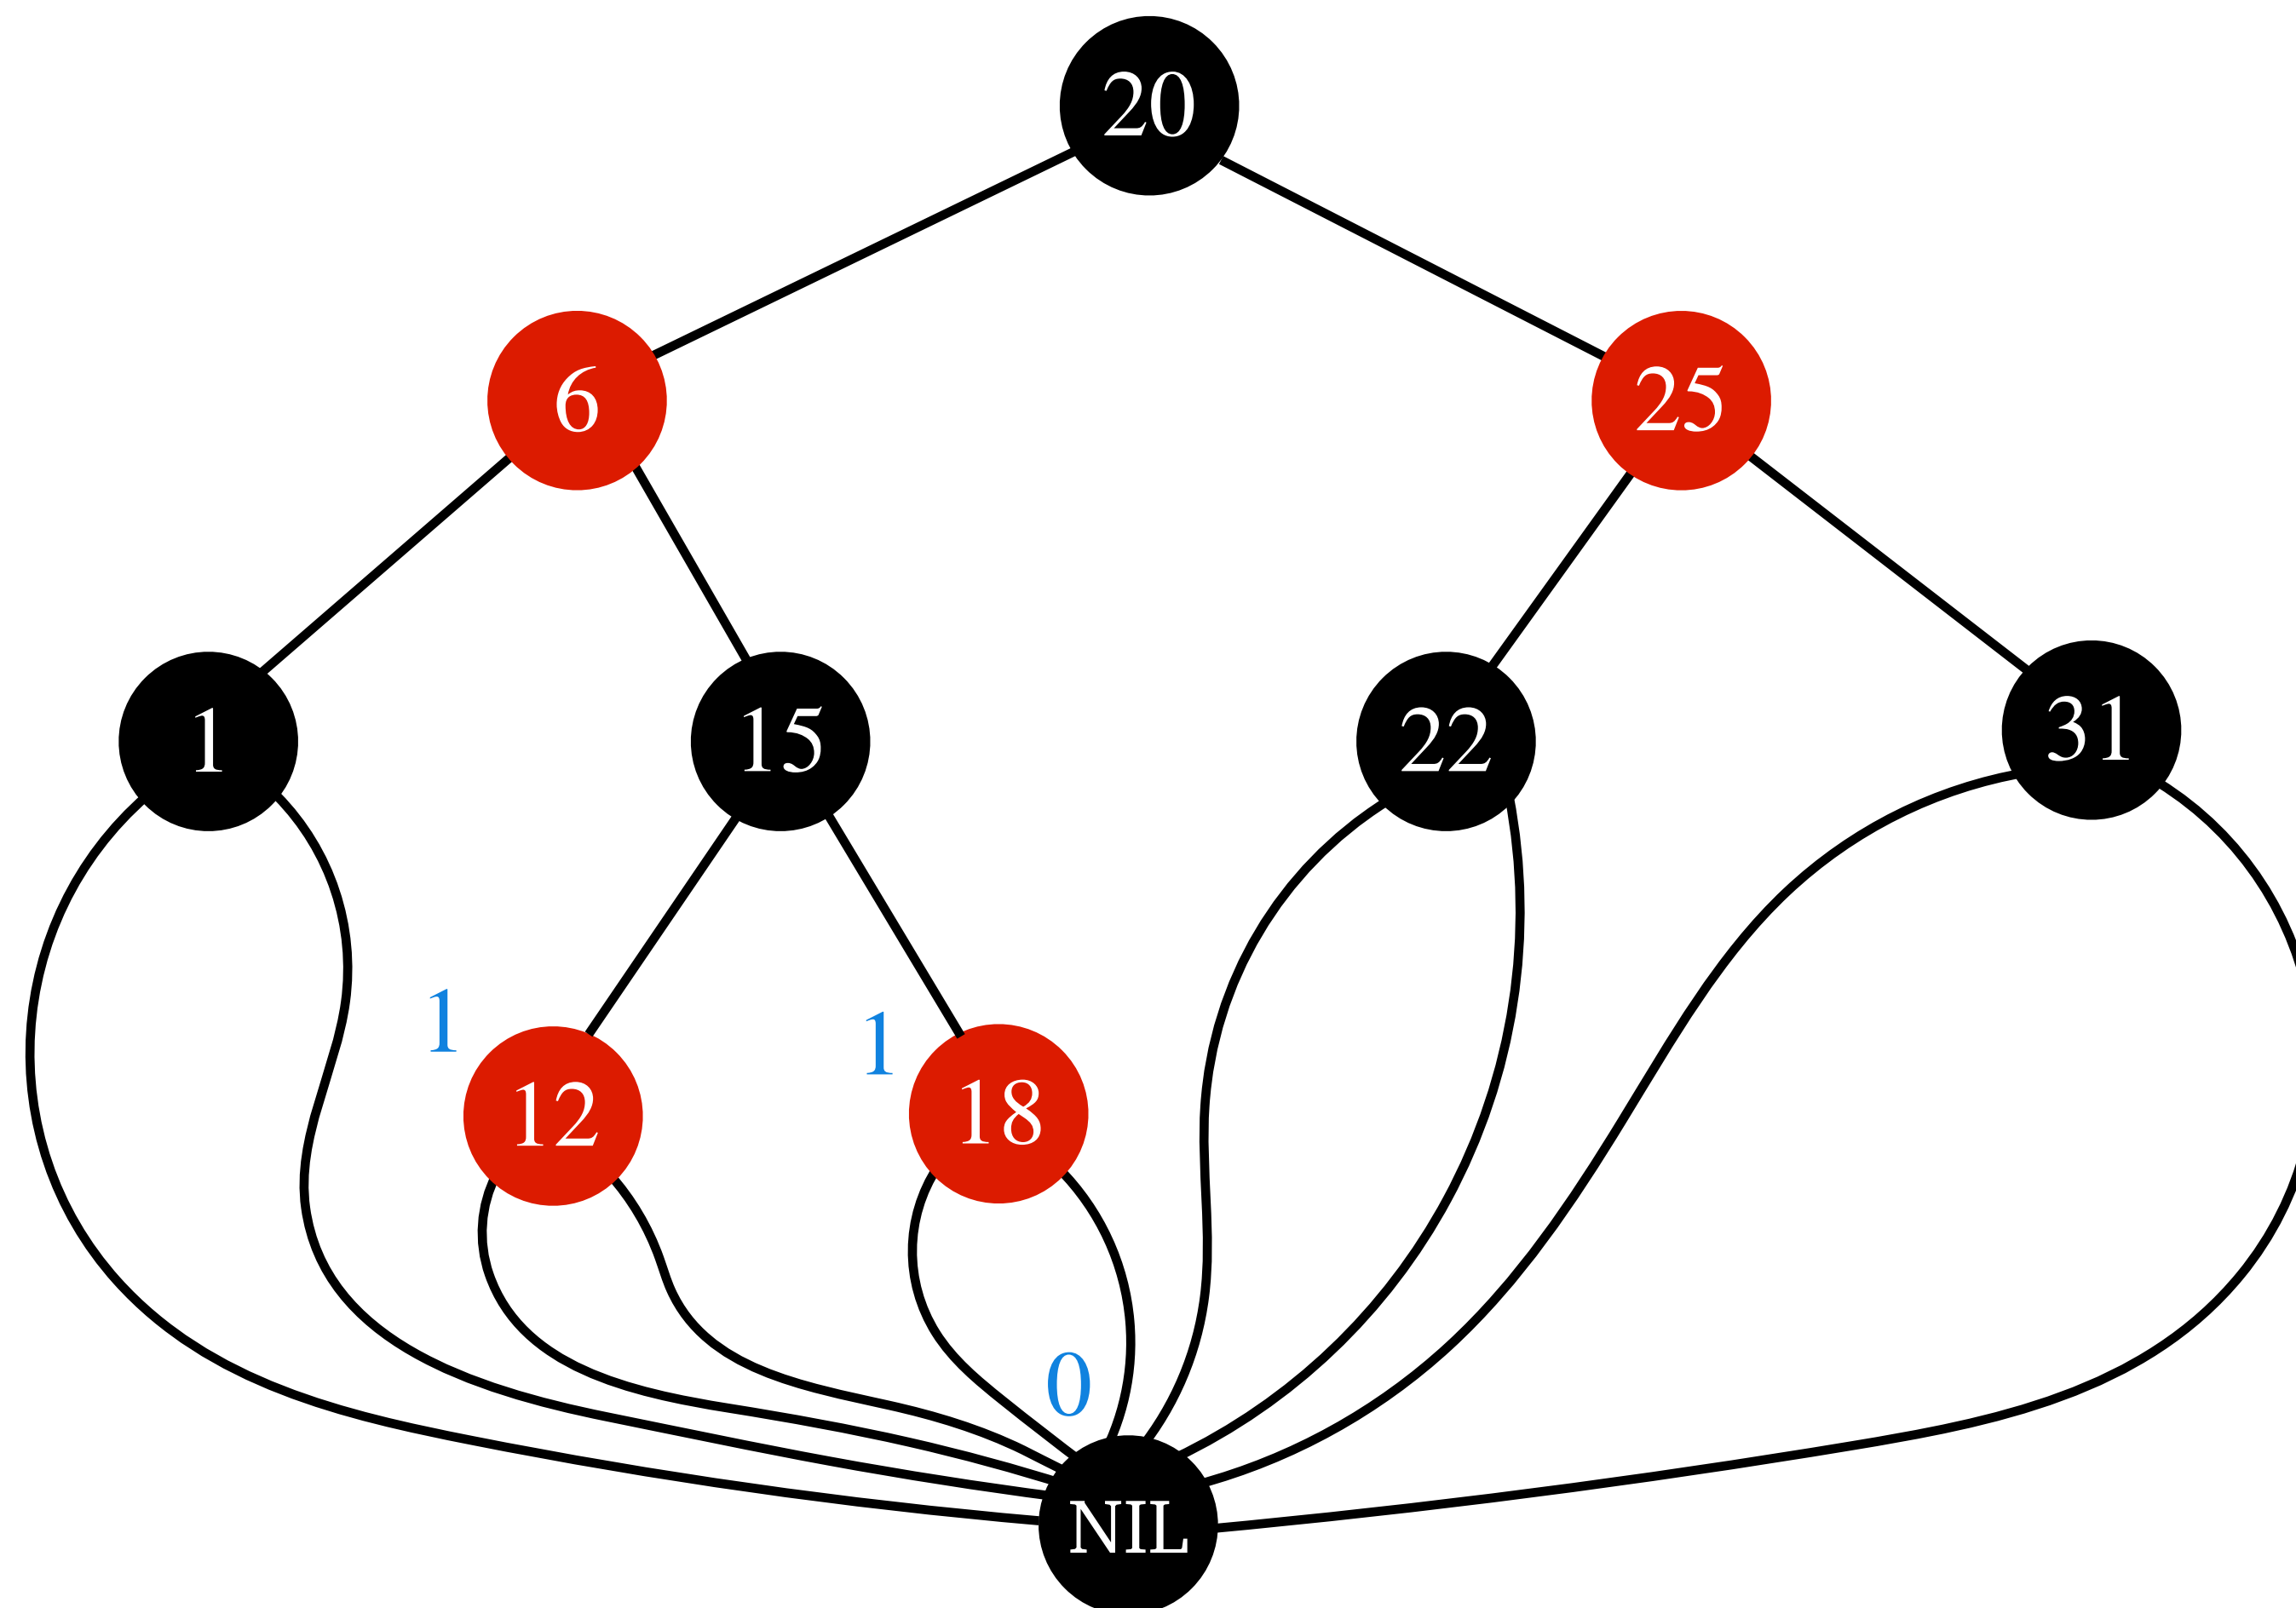
RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



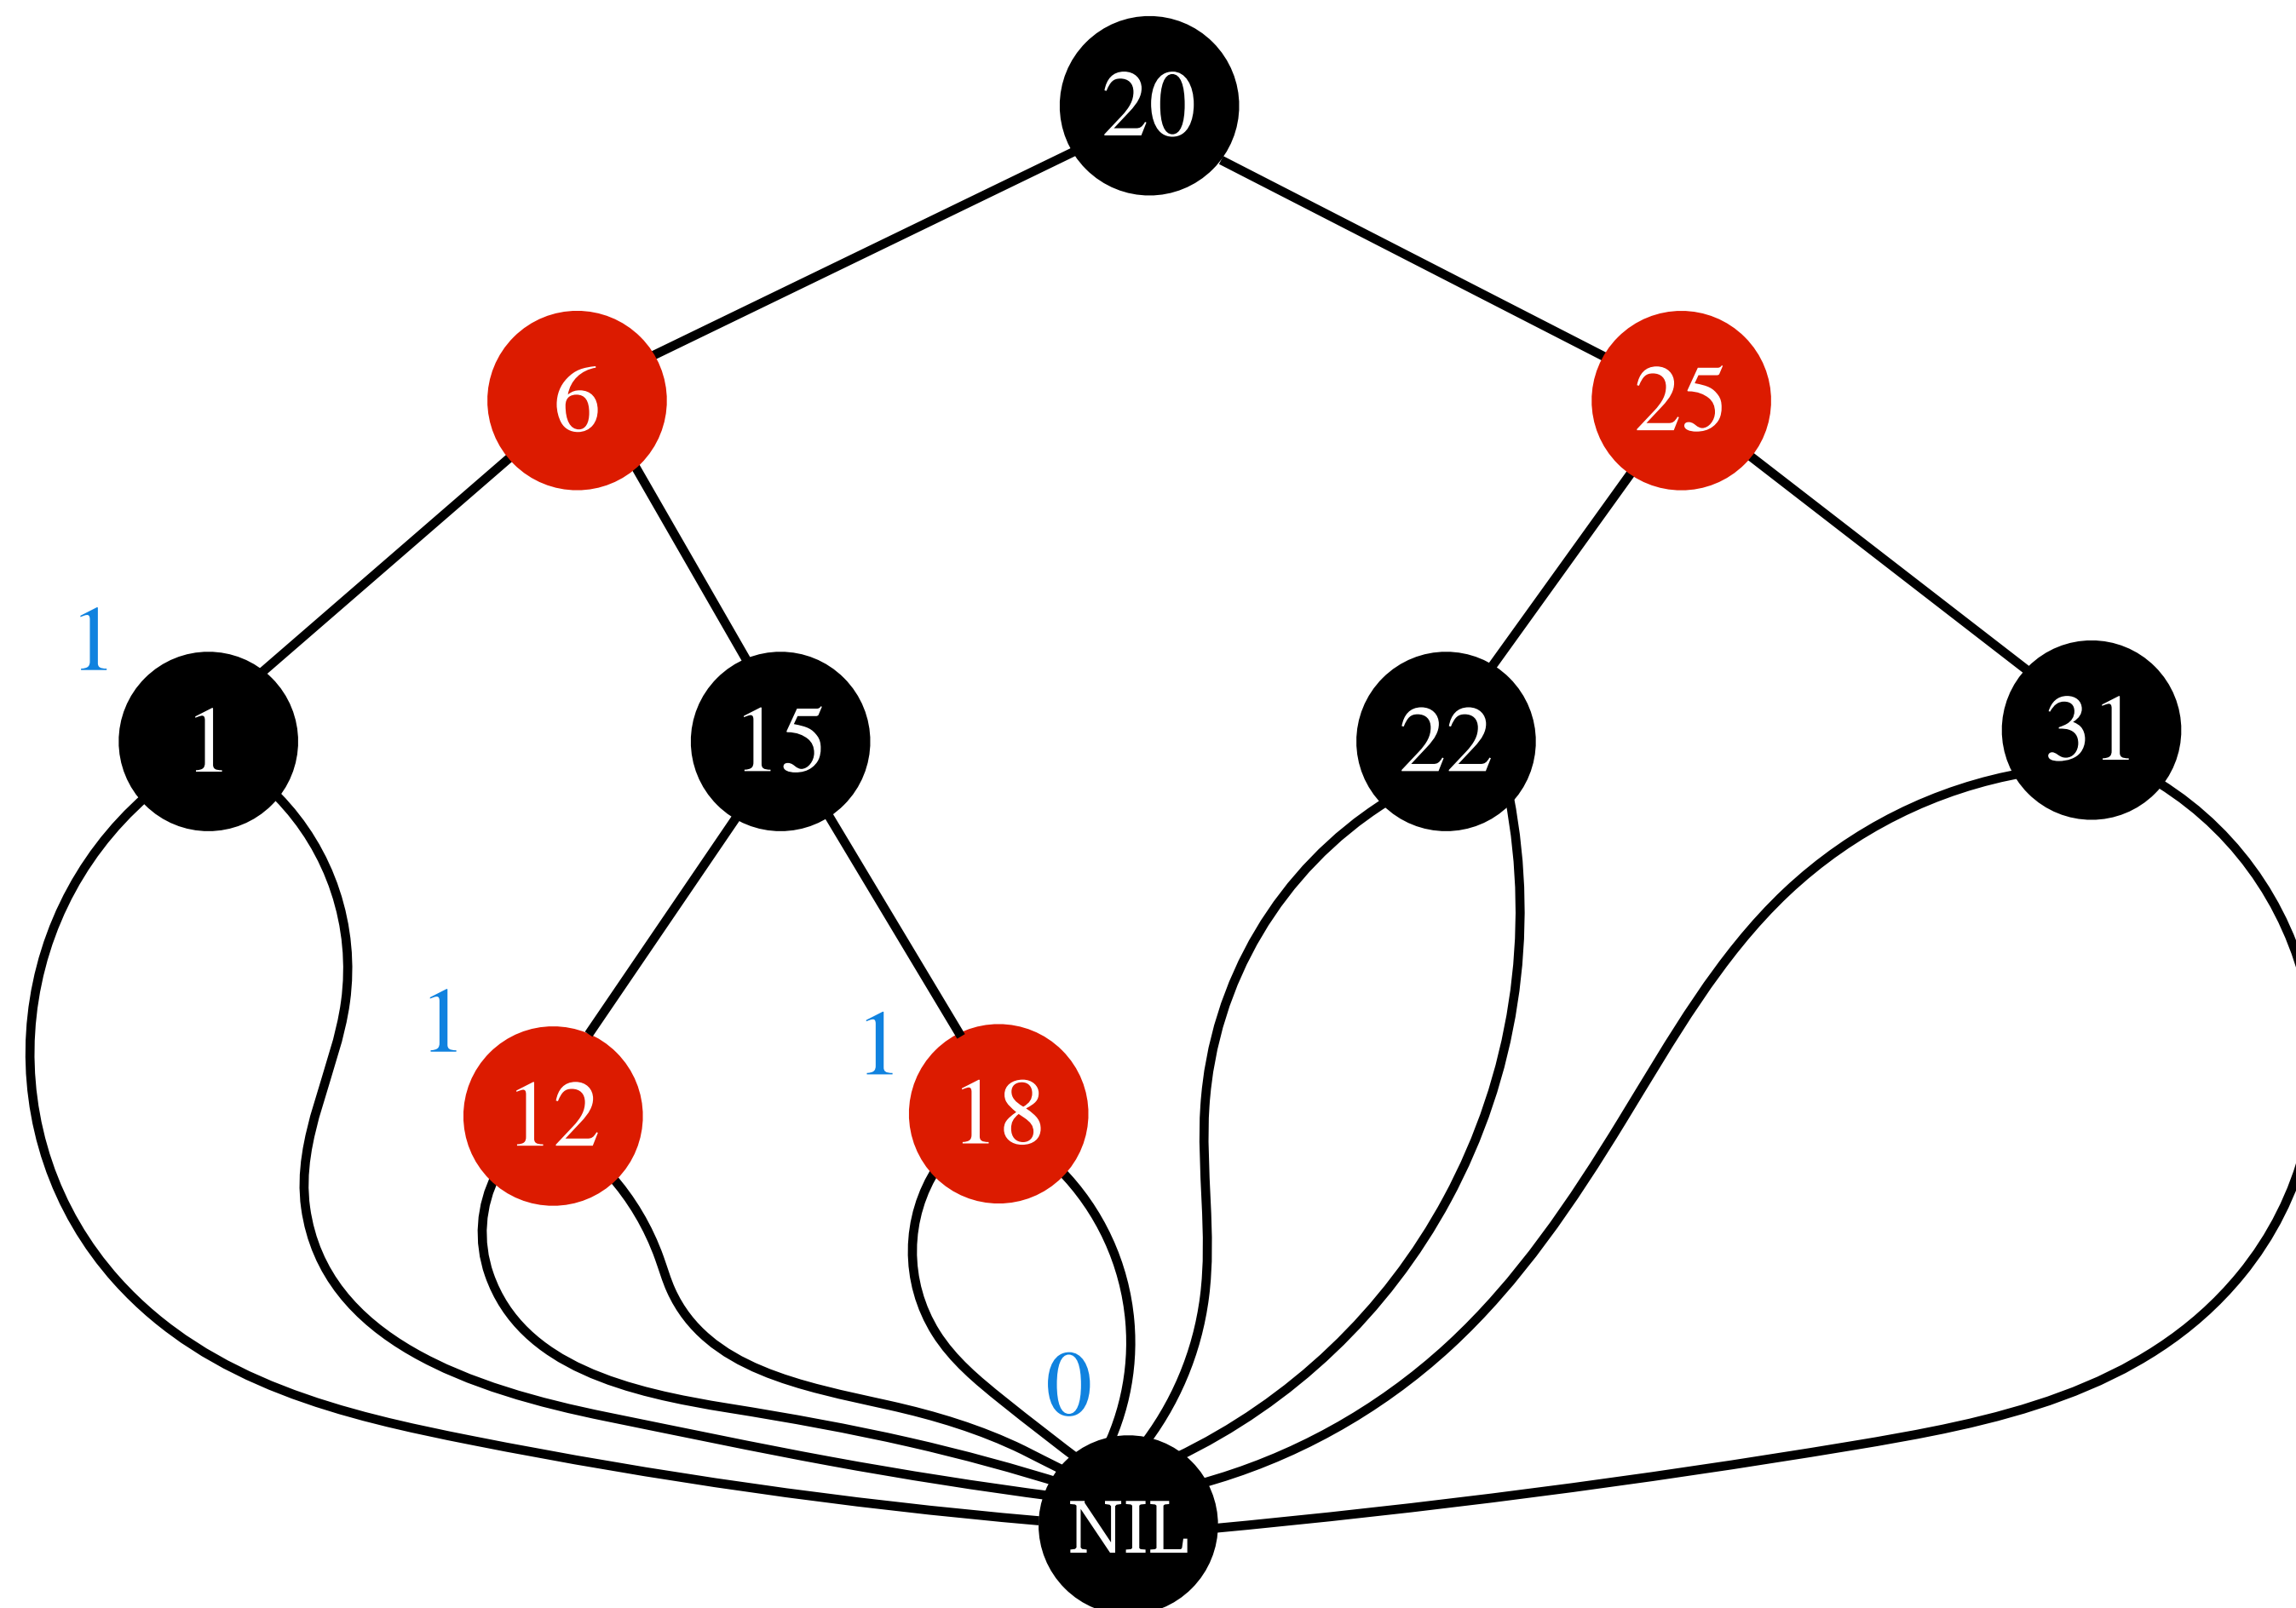
RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



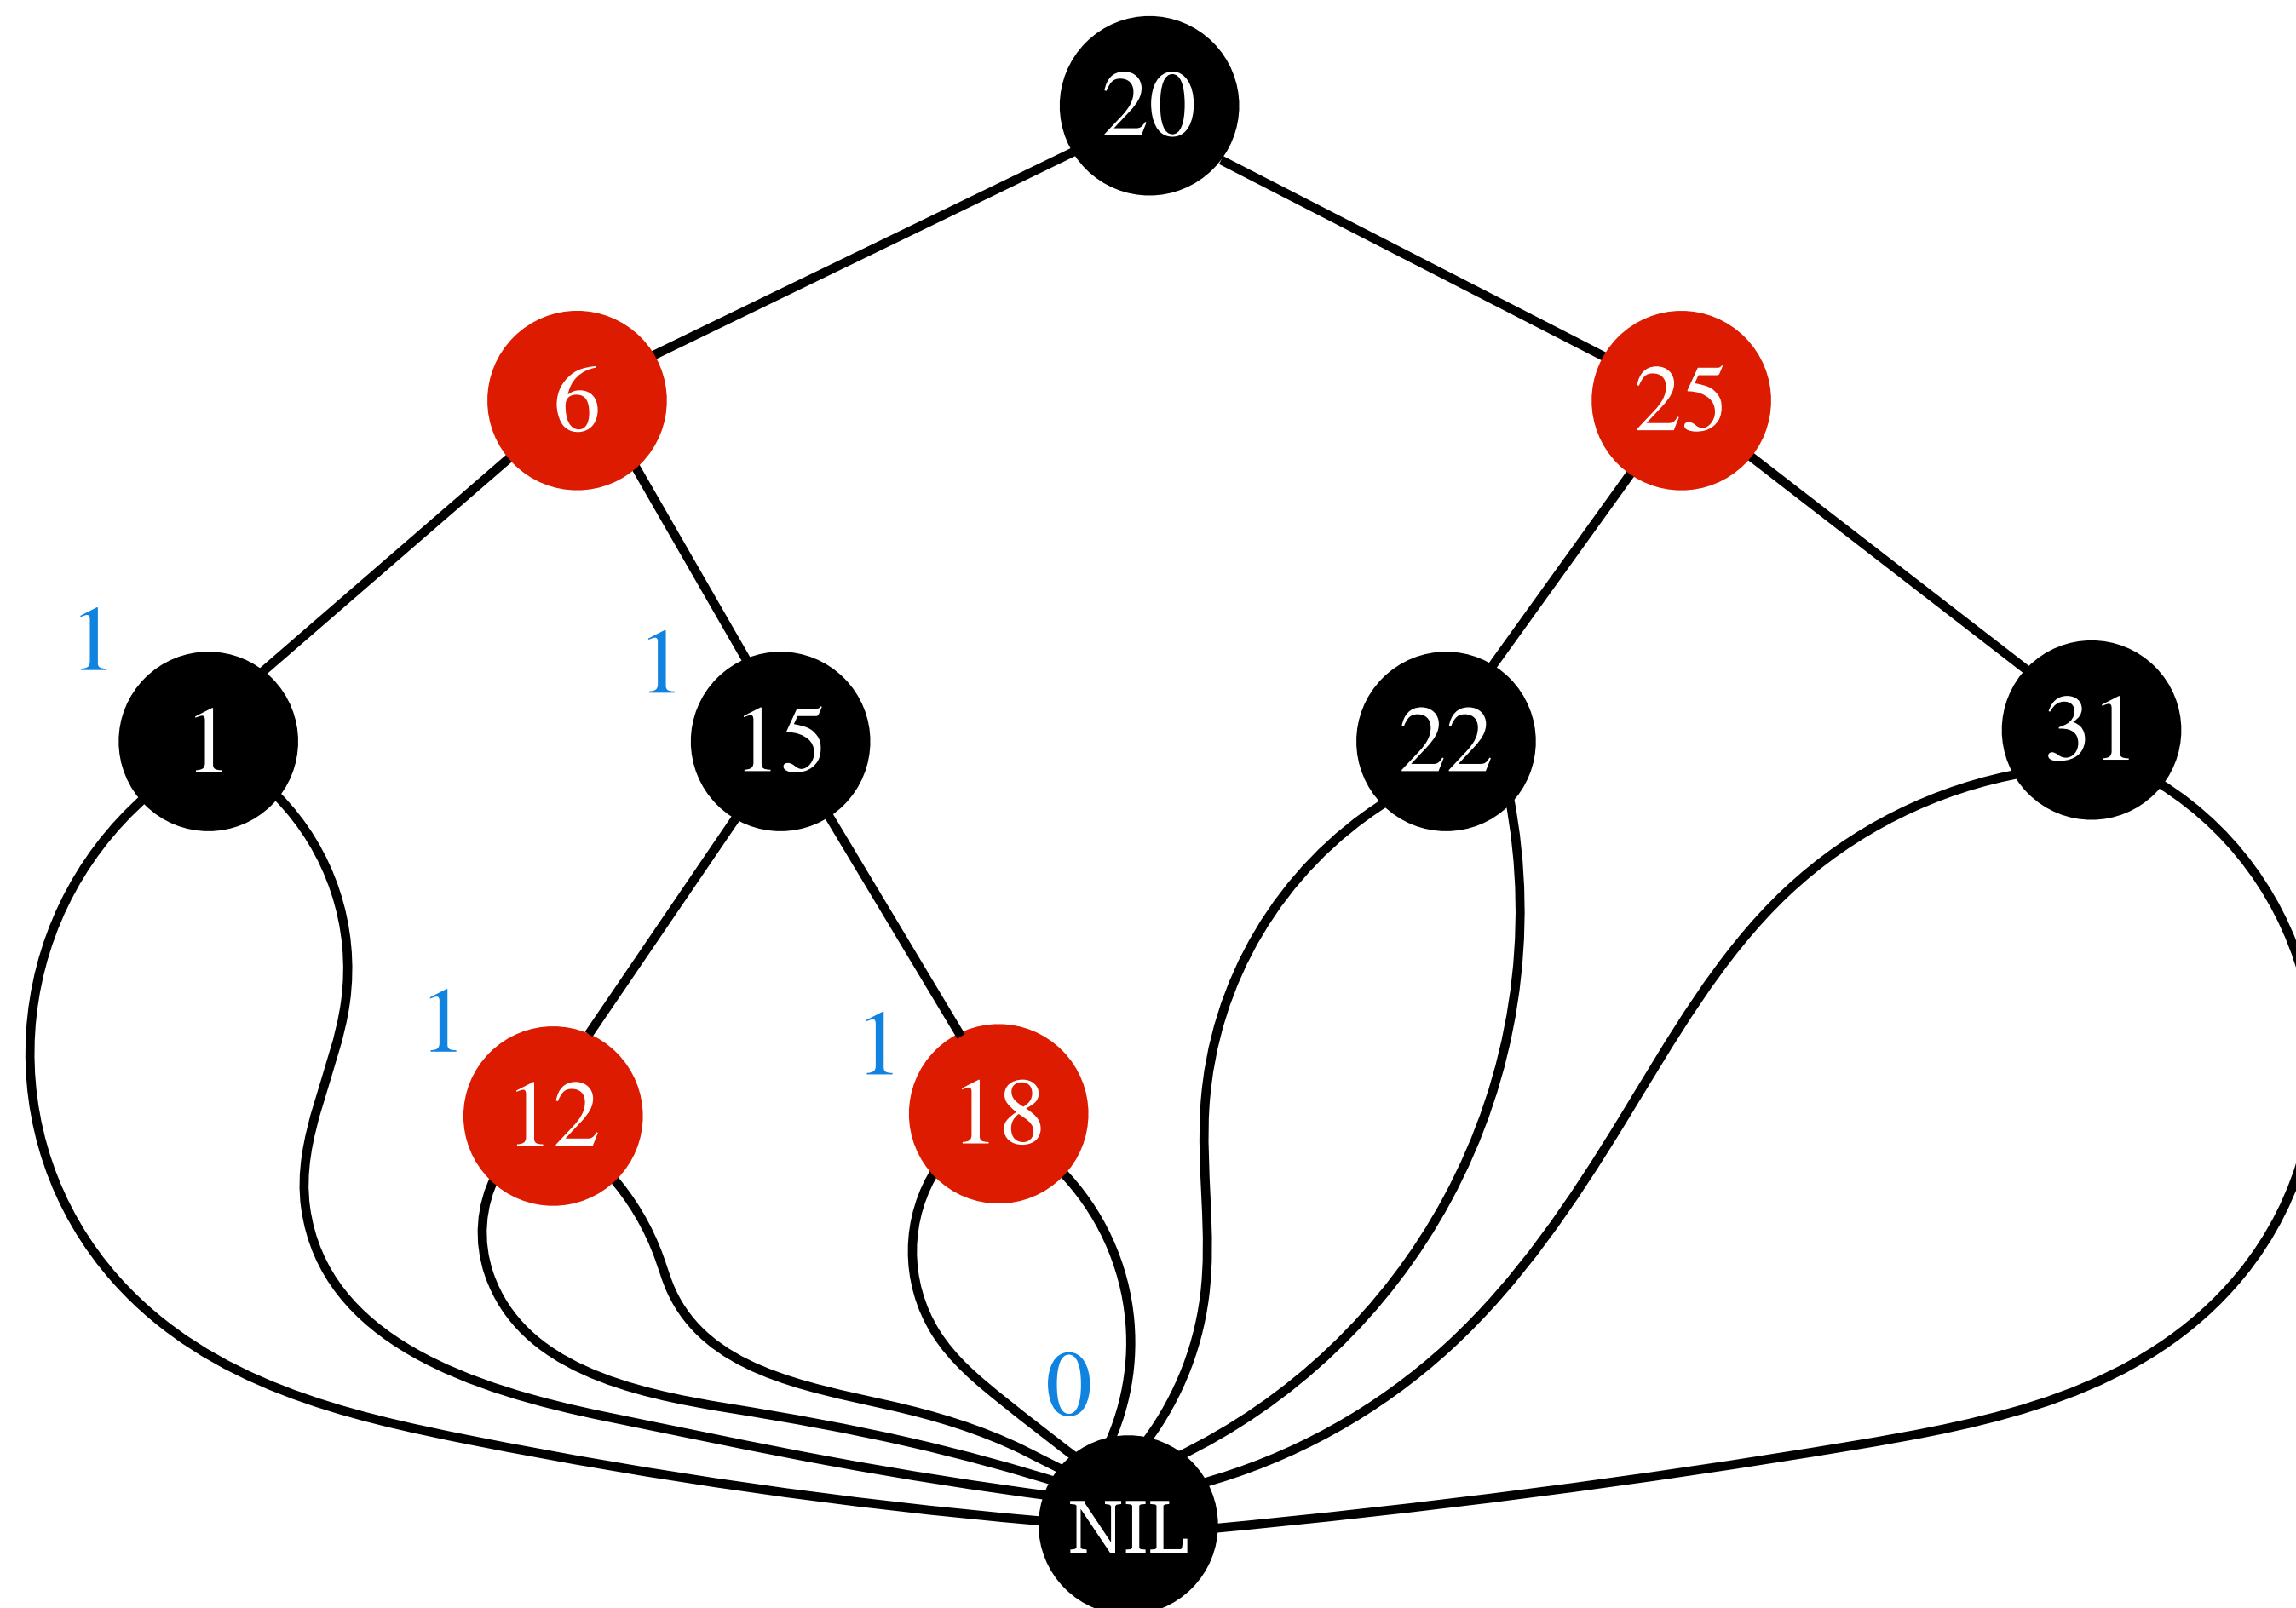
RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



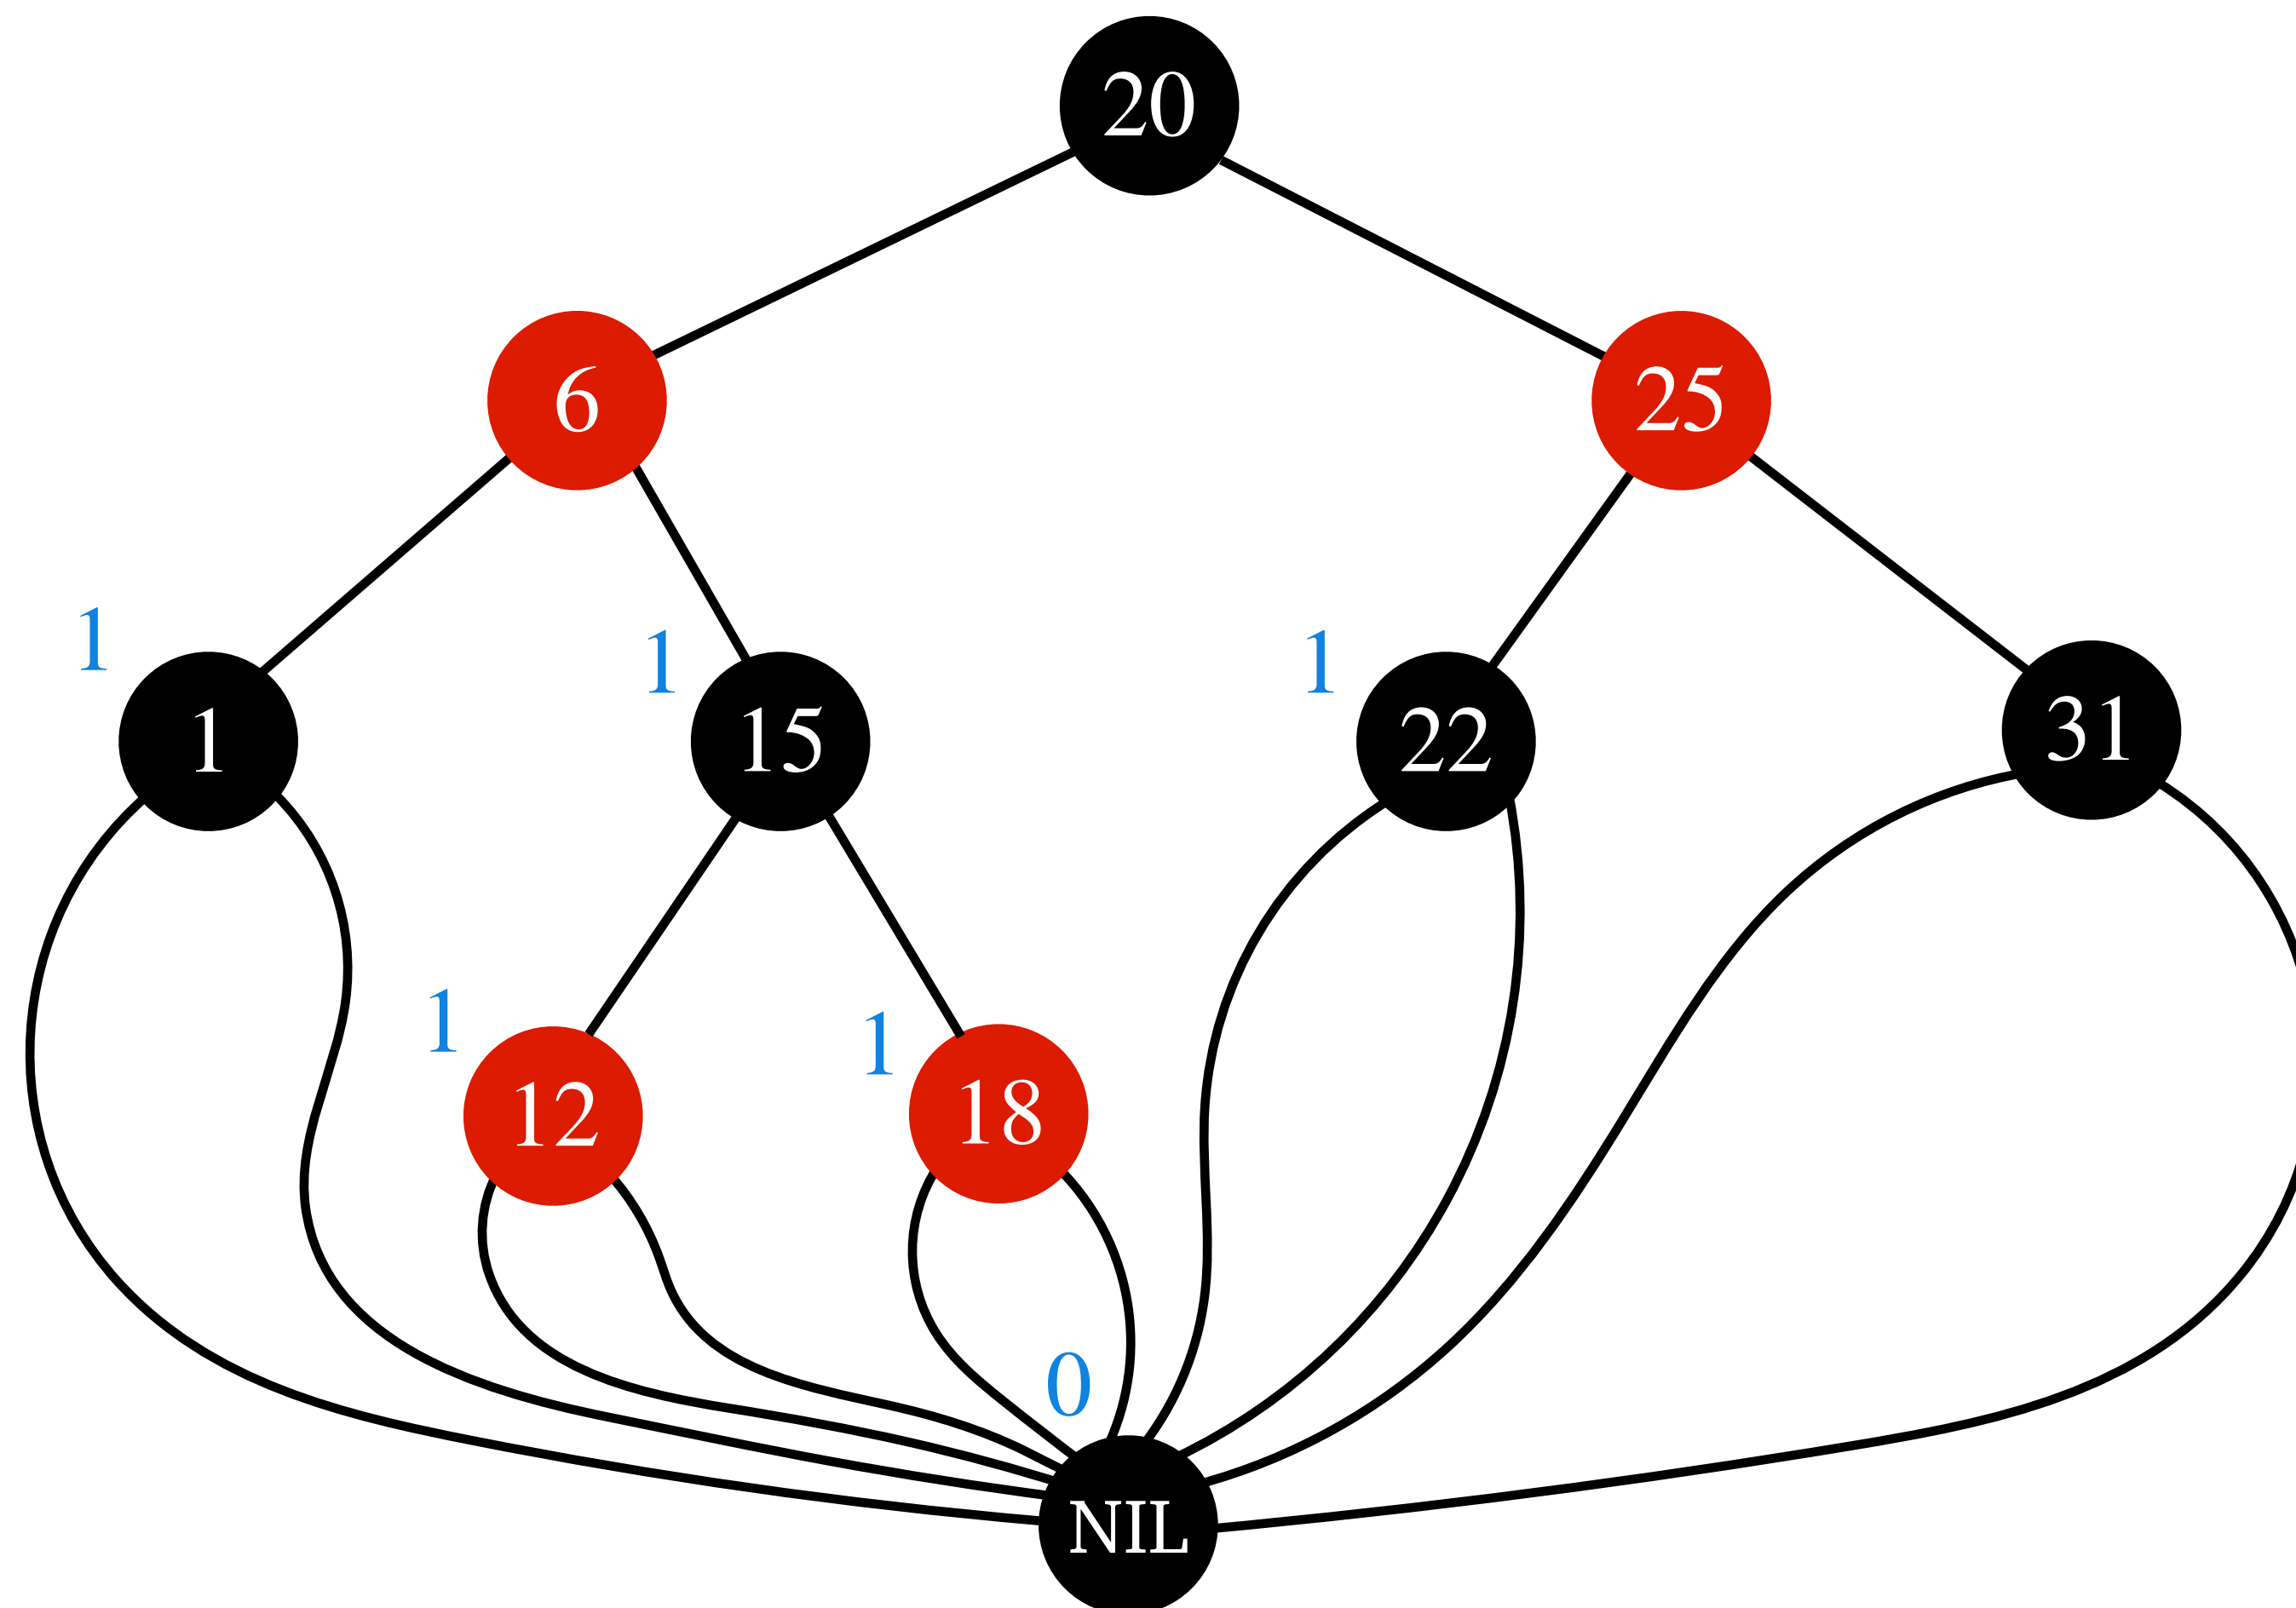
RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



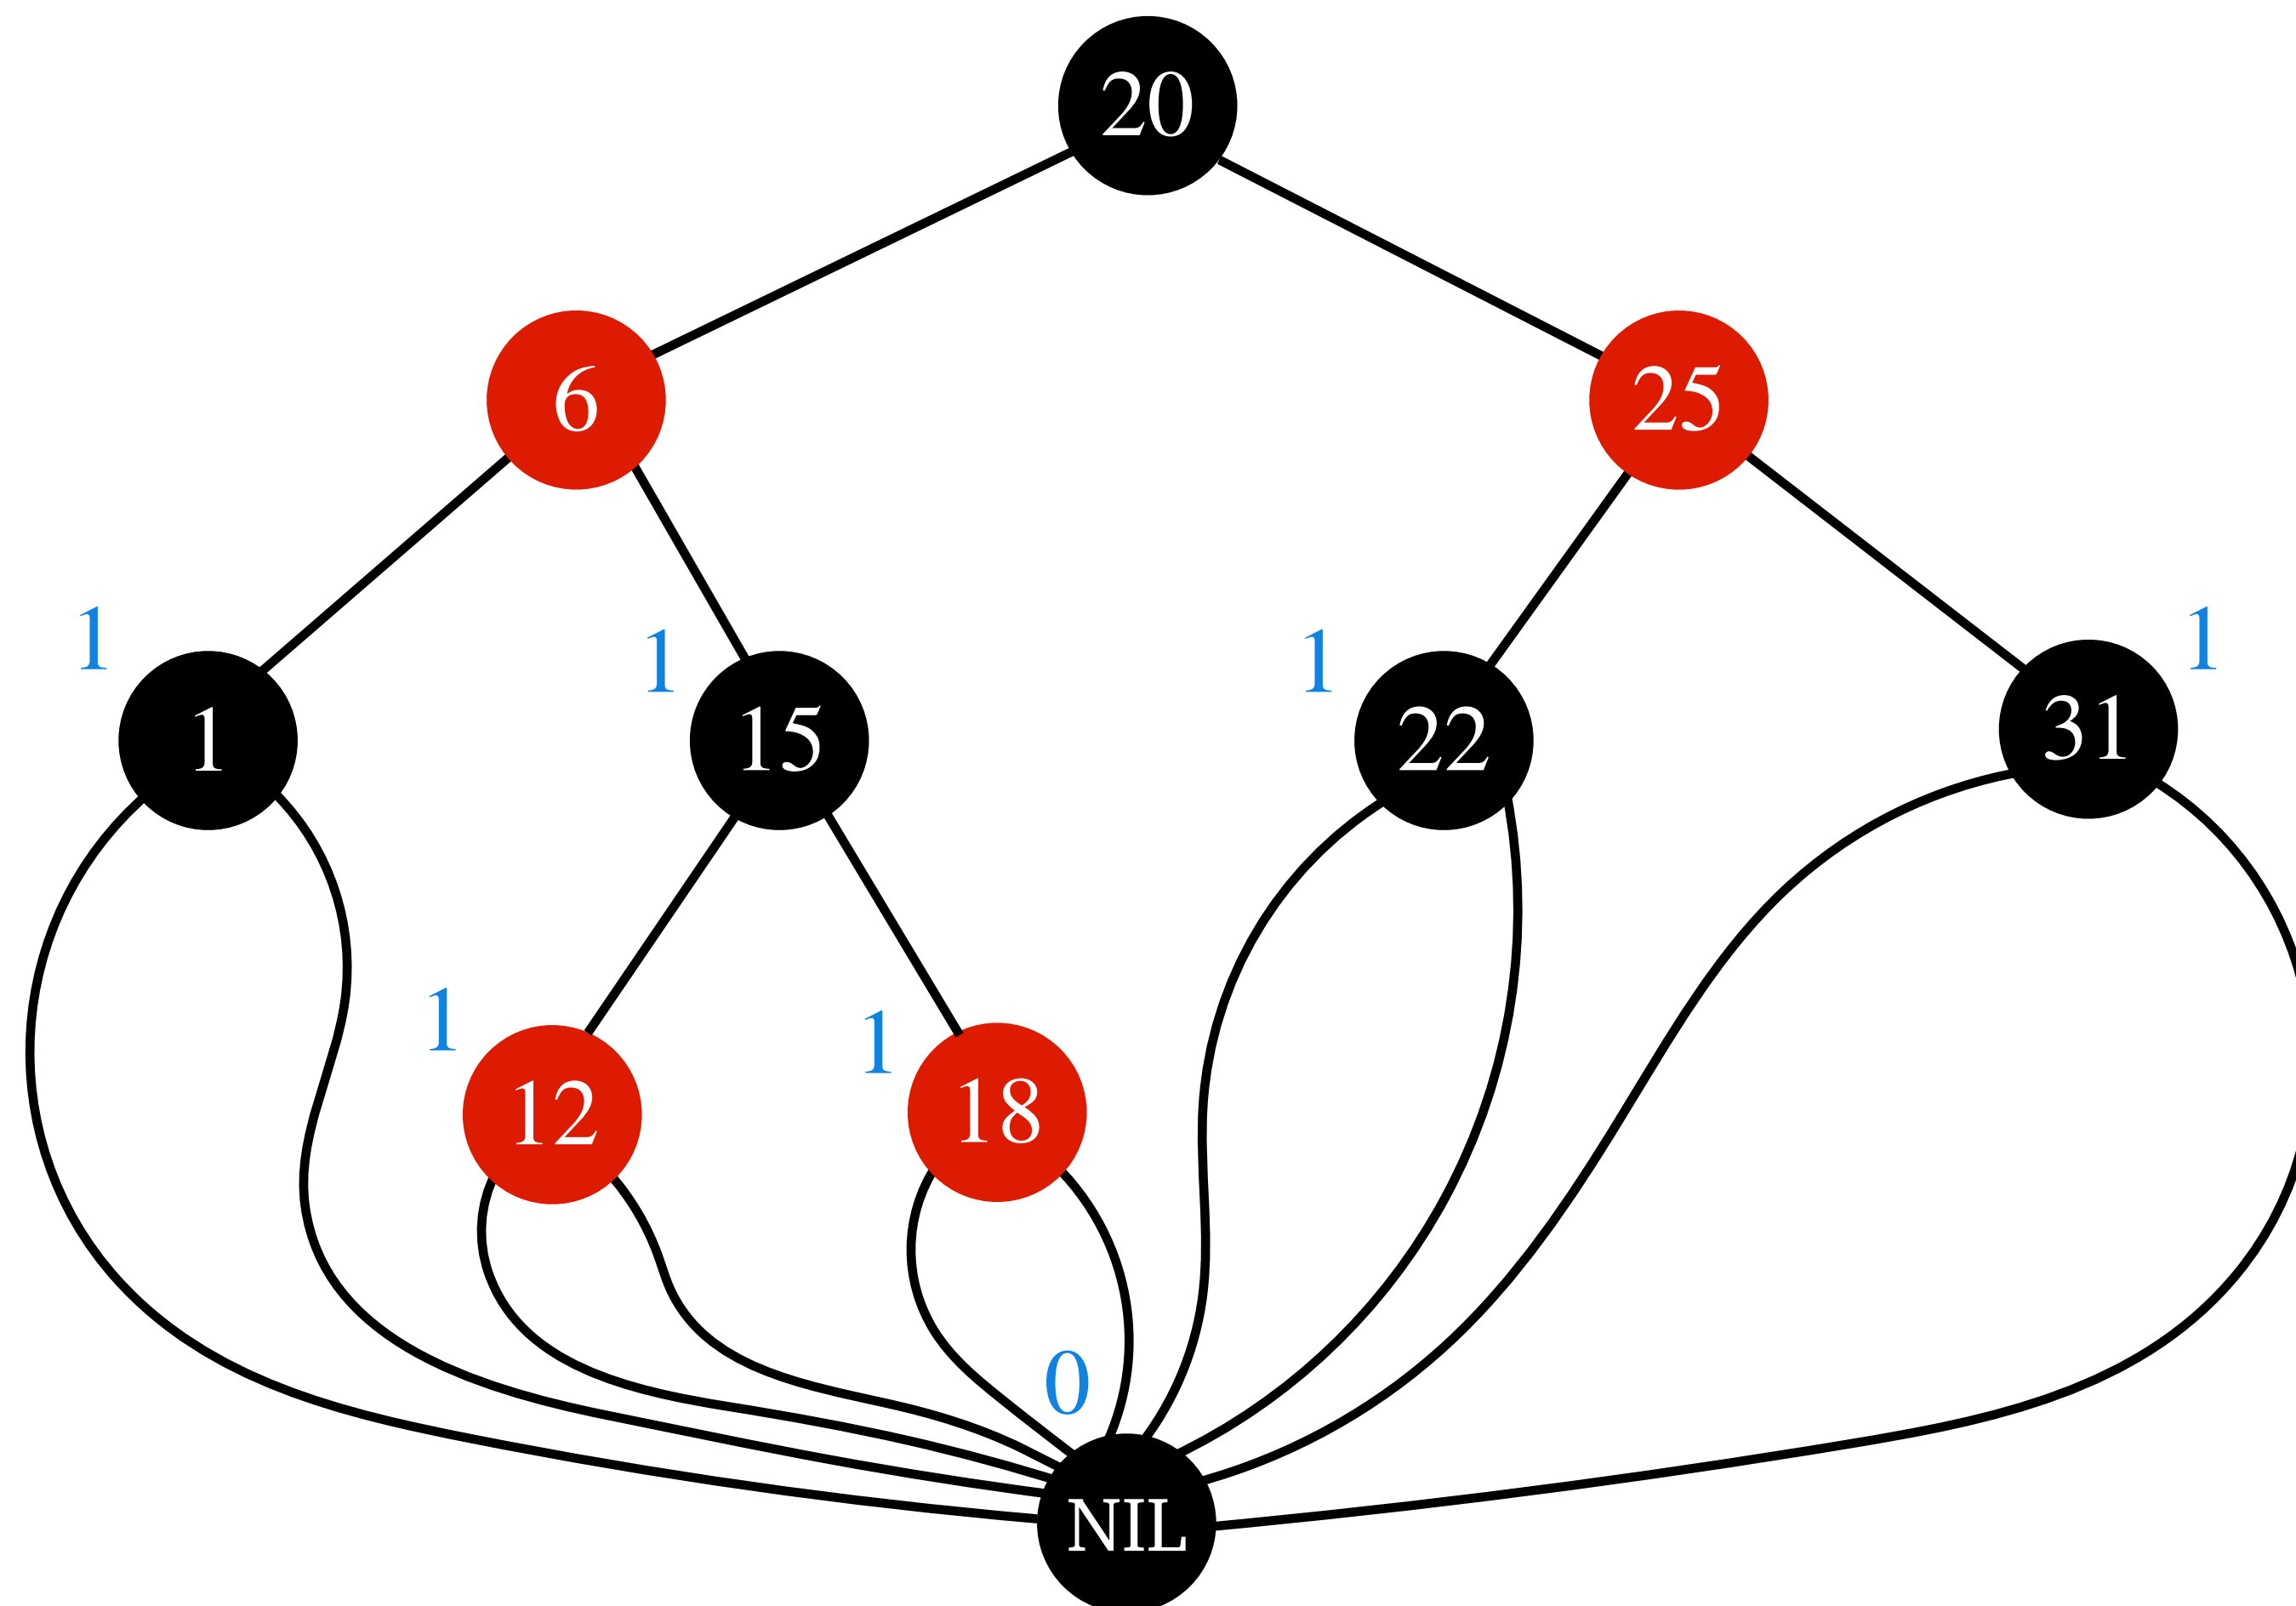
RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



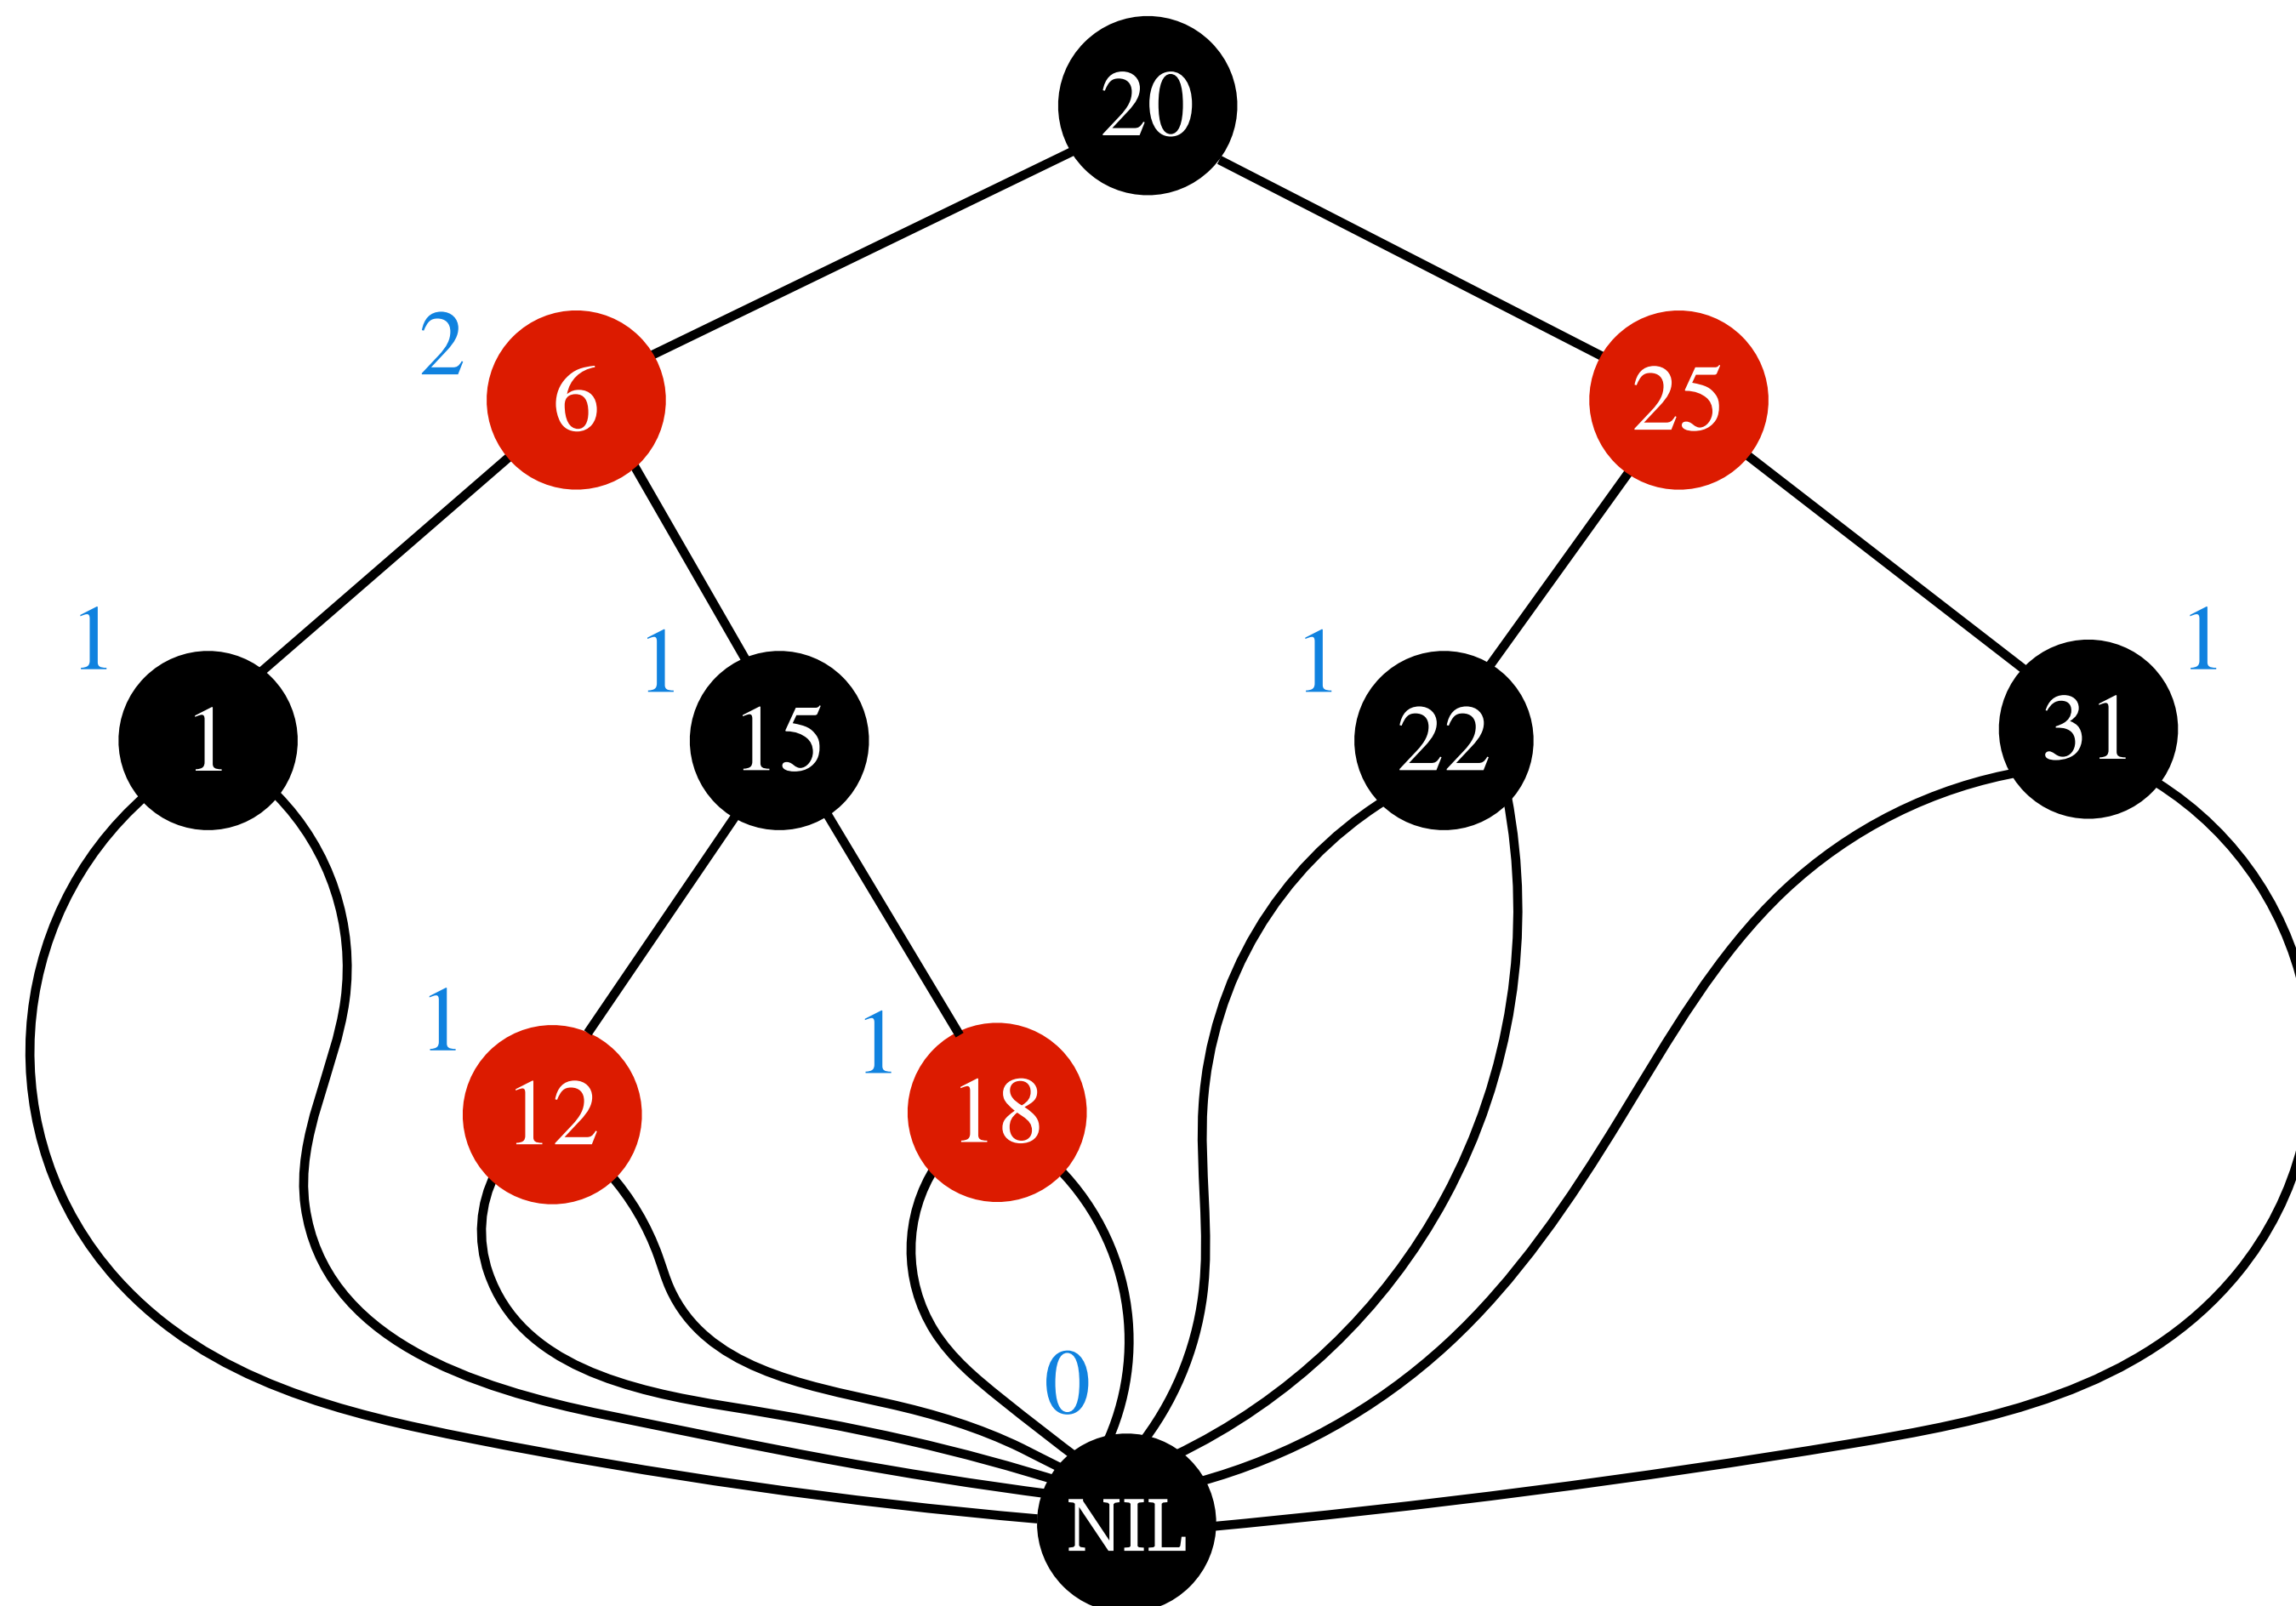
RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



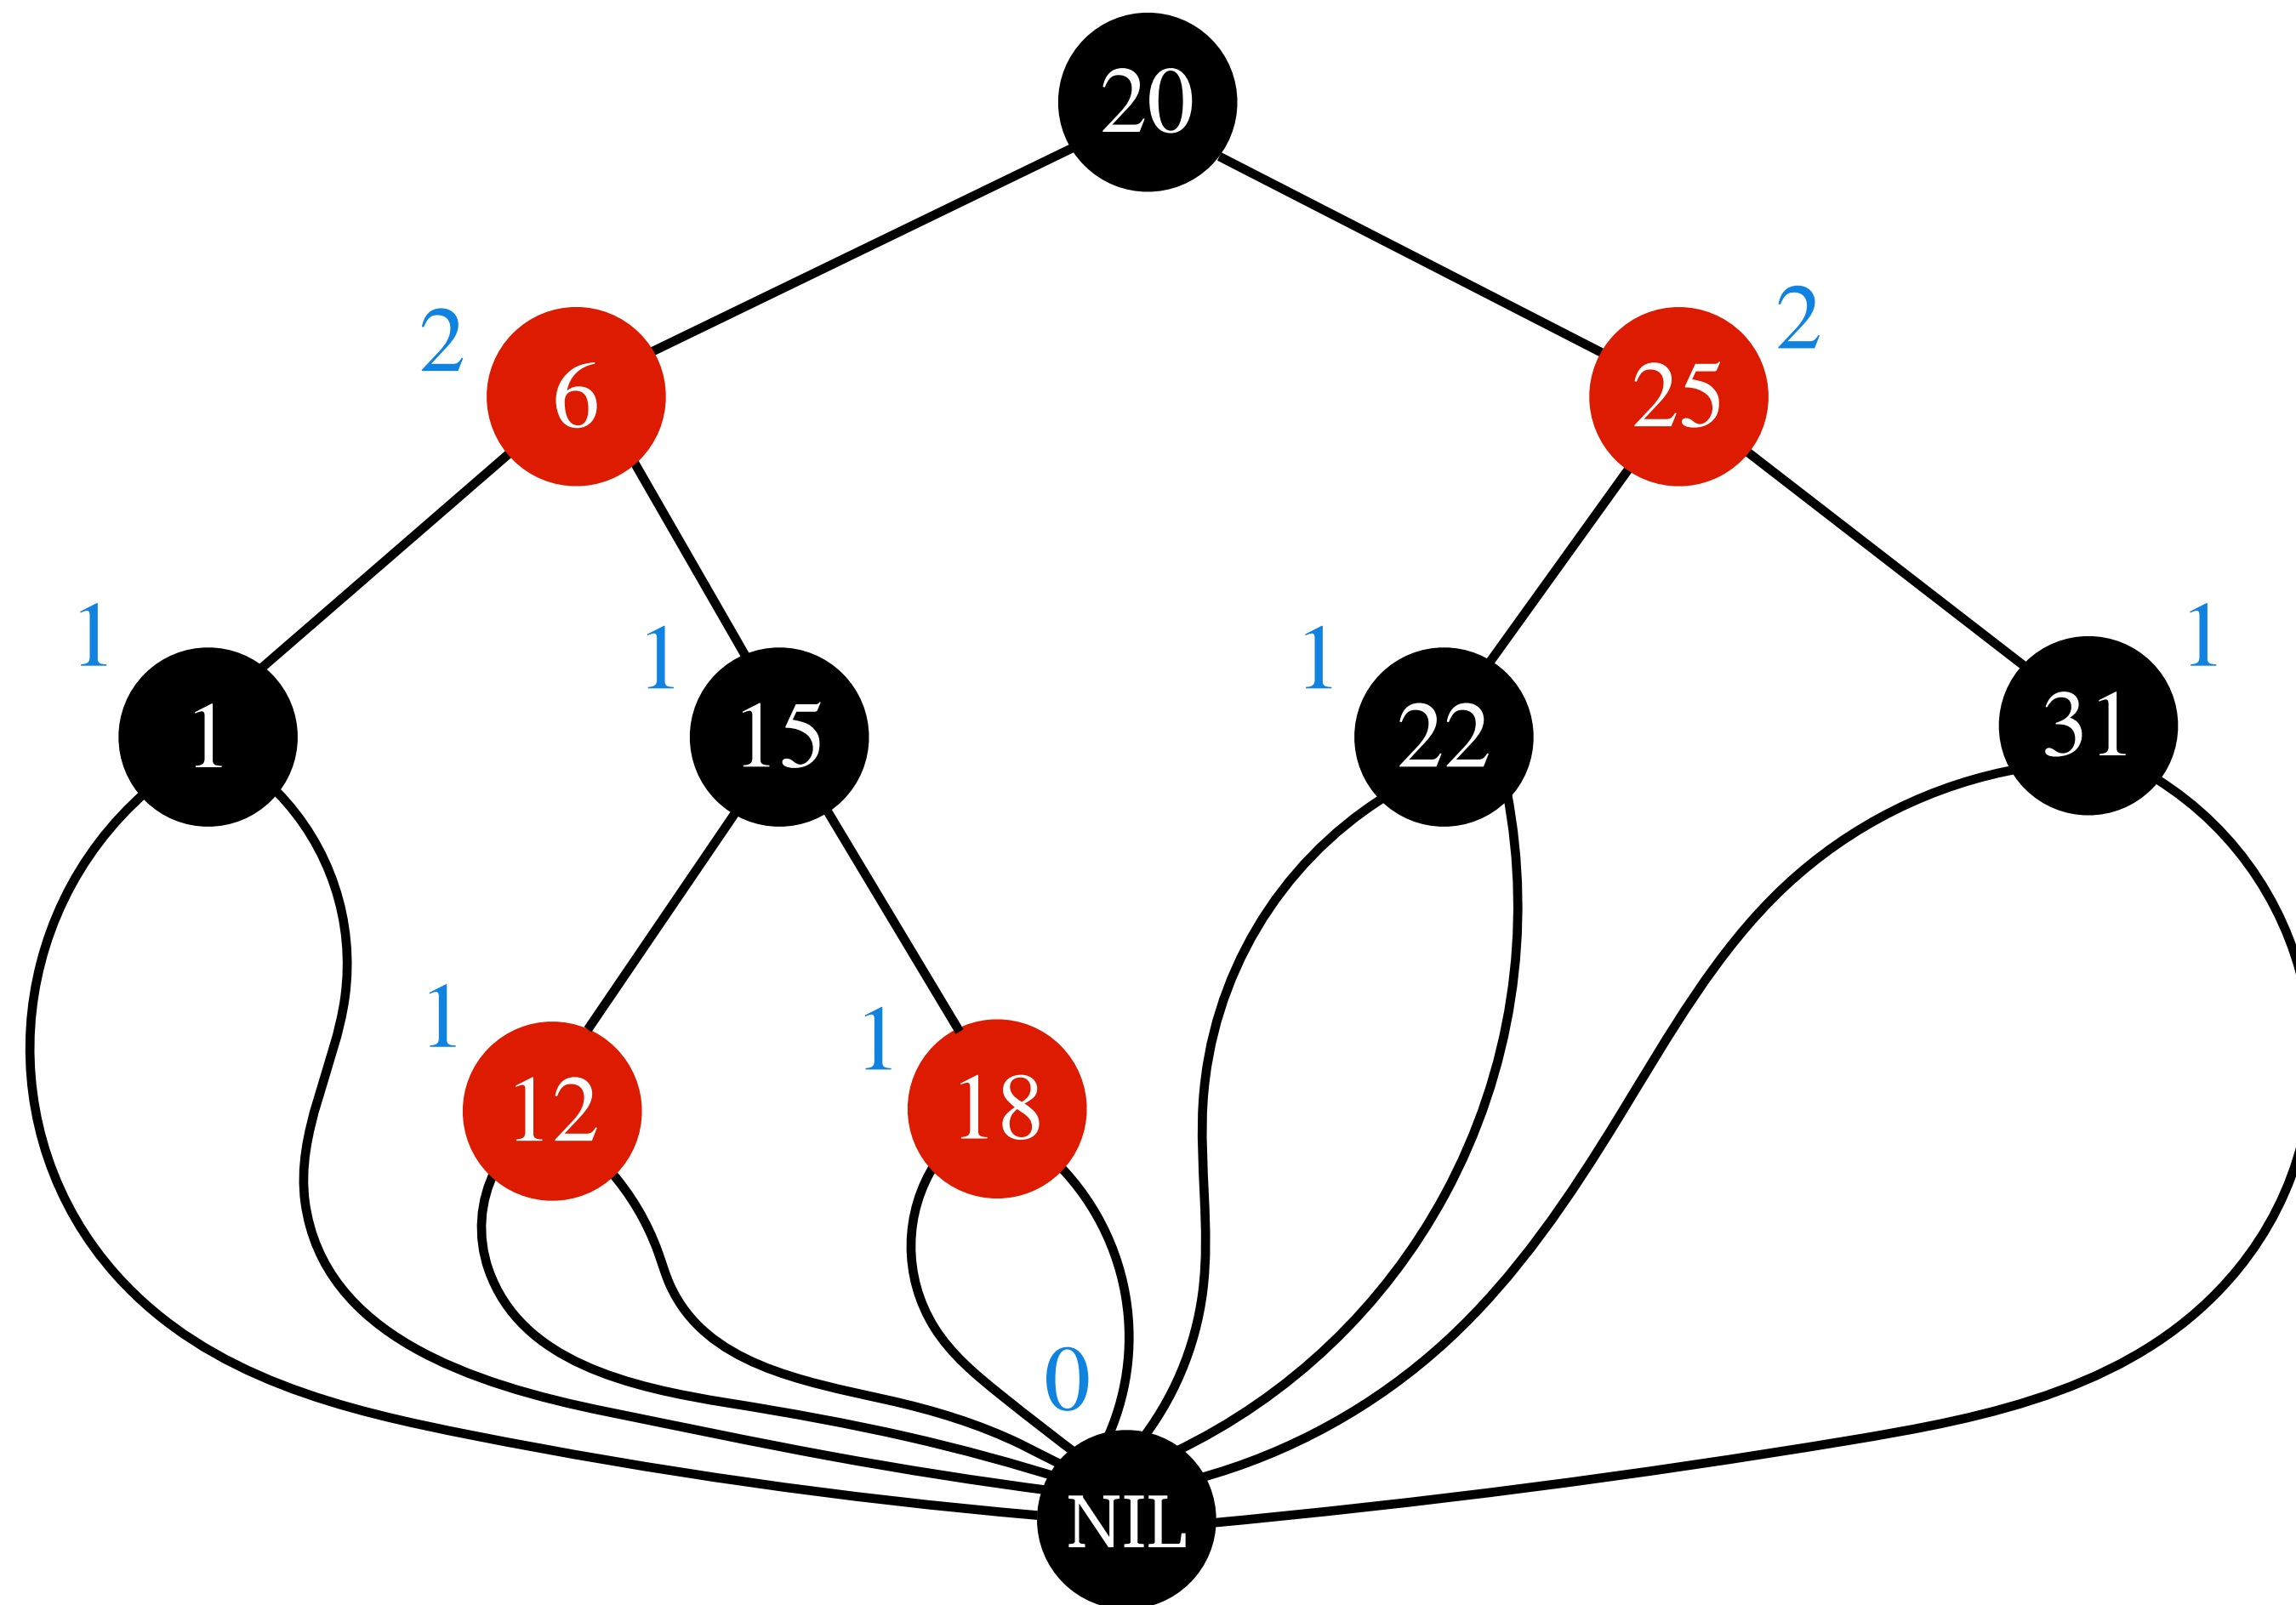
RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.



RB-Trees: Black Height

Def: For any node x , the number of black nodes on any path from x to a leaf, excluding x , is called black height of x , denoted by $bh(x)$.

