


Understanding the Three Types of Machine Learning


Machine Learning (ML) is broadly categorized into three types: **Supervised Learning**, **Unsupervised Learning**, and **Reinforcement Learning**. Each type has its own unique characteristics, applications, and algorithms.



Made with Napkin

Core Learning Paradigms in Machine Learning

1. Supervised Learning (Labeled Data) → Classification & Regression



Made with Napkin

Supervised Learning involves training a model on labeled data, where each input has a corresponding output.


Classification (Predict Discrete Labels)

- **Logistic Regression:** Estimates probabilities using a sigmoid function for binary classification. *Example:* Predicting whether an email is spam or not.
- **Decision Trees:** Splits data recursively based on feature thresholds to classify samples. *Example:* Classifying loan applicants as low or high risk.
- **Random Forest:** An ensemble of decision trees that reduces overfitting and improves accuracy. *Example:* Diagnosing diseases from patient symptoms.

- **Support Vector Machines (SVM):** Finds the optimal hyperplane to separate classes with maximum margin. *Example:* Handwritten digit recognition.

- **Naïve Bayes:** Uses Bayes' Theorem with feature independence assumption for fast classification. *Example:* Sentiment analysis in text reviews.

- **Deep Learning Example** → Convolutional Neural Networks (CNNs): Extracts features from images for complex classification tasks. *Example:* Face recognition in security systems.




Made with Napkin

Regression (Predict Continuous Values)


- **Linear Regression:** Models the relationship between input and output using a straight-line equation. *Example:* Predicting house prices based on area.
- **Polynomial Regression:** Extends linear regression by fitting a polynomial curve to capture non-linearity. *Example:* Predicting temperature variations over time.

- **Deep Learning Example** → Recurrent Neural Networks (RNNs): Captures sequential dependencies for complex regression problems. *Example:* Stock price prediction based on historical trends.



Made with Napkin

2. Unsupervised Learning (Unlabeled Data) → Clustering & Dimensionality Reduction



Made with Napkin

Unsupervised Learning works with unlabeled data to identify hidden patterns and structures.


Clustering

- **K-Means:** Partitions data into K clusters by minimizing intra-cluster variance. *Example:* Customer segmentation in e-commerce.
- **Hierarchical Clustering:** Builds a tree-like structure of nested clusters for flexible grouping. *Example:* Taxonomy of animal species.

- **DBSCAN:** Groups data based on density, identifying clusters of varying shapes. *Example:* Detecting anomalies in network traffic.

- **Deep Learning Example** → Self-Organizing Maps (SOMs): Uses neural networks to discover patterns in high-dimensional data. *Example:* Grouping similar handwritten digits without labels.

Unsupervised Learning: Clustering Techniques and Applications




Made with Napkin

Dimensionality Reduction

- **Principal Component Analysis (PCA):** Reduces data dimensions by transforming features into principal components. *Example:* Image compression.
- **t-SNE:** Projects high-dimensional data into a lower space while preserving local structure. *Example:* Visualizing word embeddings.

- **Autoencoders:** Neural networks that learn efficient data representations through encoding-decoding. *Example:* Noise reduction in images.

- **Deep Learning Example** → Variational Autoencoders (VAEs): Learns compact representations for complex data distributions. *Example:* Generating new human face images from latent space.




Made with Napkin

3. Reinforcement Learning (Agent-Environment Interaction) → Policy Optimization & Value-Based Learning

Reinforcement Learning trains an agent to interact with an environment and learn optimal actions through rewards and penalties.

Reinforcement Learning Cycle




Made with Napkin

Policy Optimization (Directly learn the best action policy)

- **Policy Gradient:** Optimizes a policy function using gradient ascent. *Example:* Controlling robotic arms.
- **Actor-Critic:** Combines policy-based and value-based methods for stable learning. *Example:* Self-driving car decision-making.

- **Deep Learning Example** → Proximal Policy Optimization (PPO): Enhances policy gradient methods for better exploration. *Example:* Training virtual humanoids to walk in simulations.

Which Reinforcement Learning Method to Use?




Made with Napkin

Value-Based Learning (Estimates future rewards to make decisions)

- **Q-Learning:** Uses a Q-table to estimate the best action for each state. *Example:* Robot navigation in a maze.
- **Deep Q-Networks (DQN):** Uses neural networks to approximate Q-values for large state spaces. *Example:* Playing Atari games.

- **Deep Learning Example** → AlphaZero: Uses deep neural networks and reinforcement learning to master complex games. *Example:* Playing chess at superhuman levels.



Made with Napkin