

Lecture 6

Logistic Regression & Classification

TOPIC 5

Classification

Truth vs. Prediction

Truth

Baseline reference











Dog Car No ? D Ph H D D B

D N D D D N N D D N

Truth

Prediction











*5 correct
5 incorrect*

Index	Image	Truth Label	Prediction Label	Correct?
1		Dog	Dog	✓
2		Car	Dog	✗
3		Dog	Dog	✓
4		Dog	No Dog	✗
5		Dog	Dog	✓
6		Ph	No Dog	✓
7		H	Dog	✗
8		Dog	No Dog	✗
9		Dog	Dog	✓
10		B	Dog	✗

Is it a dog?

Class 1: (+ve) class \Rightarrow DOG
Class 2: (-ve) class \Rightarrow No-Dog.

TP = 1
TN = 4
FP = 3
FN = 2











Truth										
Prediction	Dog	Dog	Dog	No Dog	Dog	No Dog	Dog	No Dog	Dog	Dog
	✓	✗	✓		✓		✗		✓	✗

$P_r: \frac{TP}{TP+FP} = \frac{1}{4}$ $Re: \frac{TP}{TP+FN} = \frac{1}{3}$

True Positive = 4 **False Positive = 3**

Out of all dog truth samples how many did you get right? Recall

Out of all dog predictions how many did you get right? Precision

Truth											⚡
Prediction	Dog	Dog	Dog	No Dog	Dog	No Dog	Dog	No Dog	Dog	Dog	
				×	✓			×			

1/3

}

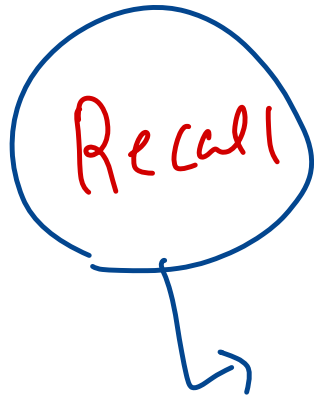
True Negative = 1

False Negative = 2

$$\text{Precision} = \frac{TP}{TP + FP} =$$

D	N-D
$\frac{4}{4+3} = \frac{4}{7}$	$\frac{1}{4}$
$\frac{4}{4+2} = \frac{2}{3}$	$\frac{1}{3}$

Recall = $\frac{TP}{TP + FN}$



$$\frac{\text{No. of correct}}{\text{Total. instances}} = \frac{5}{10} = 50\%$$

$$\frac{y}{i}$$


Truth										
Prediction	Dog	Dog	Dog	No Dog	Dog	No Dog	Dog	No Dog	Dog	Dog
	✓	✗	✓	✗	✓	✓	✗	✗	✓	✗

$$\frac{TP + TN}{TP + FN + FP + TN}$$

How many we got right? → 5

Accuracy → 5/10 → 0.5

Predict:
N.



0	1	2	5	6	7	8	9	1	0
0	1	0	5	6	7	8	9	1	0
N	N		N	N	N	N	N	N	N

6	2	3	4	5	1	1	1	0	4
4	3	4	5	1	1	1	0	4	4
N	N	N	N	N	N	N	N	N	N

Identify all 2's

Not 2.

$$\frac{18}{20} = 90\%$$

Overall

TP=0

2

Precision vs. Recall

- Precision is Out of all **dog predictions** how many you got right (Prediction baseline)
- Precision = $4/7$
- **Precision = $TP / (TP+FP)$**
- Recall (truth is the baseline)
- Out of all **dog truth** samples how many did you get right?
- Total dog truth samples = 6; True positive = 4
- **Recall = $TP / (TP+FN)$**
- Exercise: What is precision and recall for no-dog class ?

1. Precision:

- Precision is the ratio of correctly predicted positive observations to the total predicted positives. It assesses the accuracy of the positive predictions made by the model.

- The precision formula is given by:

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

- In other words, precision is the ability of the classifier not to label as positive a sample that is actually negative.

2. Recall (Sensitivity or True Positive Rate):

- Recall is the ratio of correctly predicted positive observations to the total actual positives. It measures the ability of the model to capture all the relevant positive instances.

- The recall formula is given by:

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

- Recall is concerned with the ability of the classifier to find all the positive instances.

MNIST Data Set

- 70,000 images handwritten by school students
- This “5-detector” will be an example of a *binary classifier*, capable of distinguishing between just two classes, 5 and non-5
- Let’s use stochastic grad. Descent
SGDClassifier -

```
from sklearn.linear_model import SGDClassifier
```

```
sgd_clf = SGDClassifier(random_state=42)  
sgd_clf.fit(X_train, y_train_5)
```



```
from sklearn.datasets import fetch_openml
```

```
mnist = fetch_openml('mnist_784', as_frame=False)
```

- X – data (2D Numpy Array)- 28x28 = 784 pixel image. 70000
- y - target (1D Numpy Array as the target) – 70000
- Training first 60,000; test last 10,000 \Rightarrow
- K-fold cross-validation
- Accuracy 95% across all 3 sets ?
- Dummy classifier for non-5s:

```
>>> from sklearn.model_selection import cross_val_score
>>> cross_val_score(sgd_clf, X_train, y_train_5, cv=3, scoring="accuracy")
array([0.95035, 0.96035, 0.9604 ])
```

```
from sklearn.dummy import DummyClassifier
```

```
dummy_clf = DummyClassifier()
dummy_clf.fit(X_train, y_train_5)
print(any(dummy_clf.predict(X_train))) # prints False: no 5s detected
```

Can you guess this model's accuracy? Let's find out:

```
>>> cross_val_score(dummy_clf, X_train, y_train_5, cv=3, scoring="accuracy")
array([0.90965, 0.90965, 0.90965])
```

A: 95%
P: 80%
P: 65%

Confusion Matrices

Equation 3-1. Precision

$$\text{precision} = \frac{TP}{TP + FP}$$

TP is the number of true positives, and FP is the number of false positives.

Recall, also called *sensitivity* or the *true positive rate* (TPR): this is the ratio of positive instances that are correctly detected by the classifier.

$$\text{recall} = \frac{TP}{TP + FN}$$

```
>>> from sklearn.metrics import precision_score, recall_score
>>> precision_score(y_train_5, y_train_pred) # == 3530 / (687 + 3530)
0.8370879772350012
>>> recall_score(y_train_5, y_train_pred) # == 3530 / (1891 + 3530)
0.6511713705958311
```

```
>>> from sklearn.metrics import confusion_matrix
>>> cm = confusion_matrix(y_train_5, y_train_pred)
>>> cm
array([[53892,   687],
       [ 1891, 3530]])
```

```
>>> y_train_perfect_predictions = y_train_5 # pretend we reached perfection
>>> confusion_matrix(y_train_5, y_train_perfect_predictions)
array([[54579,    0],
       [    0, 5421]])
```

		Predicted		
		Negative	Positive	
Actual	Negative	8 3 9	6	Precision (e.g., 3 out of 4)
	Positive	7 2	5 5 5	
		TN	TP	
		FN	FP	
				Recall (e.g., 3 out of 5)

- When it claims an image represents a 5, it is correct only 83.7% (PRECISION) of the time. Moreover, it only detects 65.1% of the 5s (RECALL).

Evaluation Metrics

Metric	Formula
Average classification accuracy	$(TN + TP) / (TN + TP + FN + FP)$
Type I error (false positive rate)	$FP / (TN + FP)$
Type II error (false negative rate)	$FN / (FN + TP)$
True positive rate	$TP / (TP + FN)$
True negative rate	$TN / (TN + FP)$

- Type I error or false positive rate: The chance of incorrectly classifying a (randomly selected) sample as positive
- Type II error or false negative rate: The chance of incorrectly classification a (randomly selected) sample as negative

Evaluation Metrics

Metric	Formula
Precision	$TP / (TP + FP)$
Recall	$TP / (TP + FN)$

Precision: Fraction of retrieved instances that are relevant

Recall: Fraction of relevant instances that are retrieved

Evaluation Metrics

Metric	Formula
Sensitivity	$TP / (TP + FN)$
Specificity	$TN / (TN + FP)$
Predictive value for a positive result (PV+)	$TP / (TP + FP)$
Predictive value for a negative result (PV-)	$TN / (TN + FN)$

Sensitivity: Proportion of actual positives which are correctly identified

Specificity: Proportion of actual negatives which are correctly identified

Sensitivity: The chance of correctly identifying positive samples. A sensitive test helps rule out disease (when the result is negative)

Specificity: The chance of correctly classifying negative samples. A very specific test rules in disease with a higher degree of confidence.

F1 score

```
>>> from sklearn.metrics import f1_score
>>> f1_score(y_train_5, y_train_pred)
0.7325171197343846
```

- A single metric to combine precision and recall - The harmonic mean
- Unlike regular mean, harmonic mean gives more weight to low values.
- Therefore, the classifier's F1 score is only high if both recall and precision are high.
- The F1 score favors classifiers that have similar precision and recall

$$\frac{2}{\frac{1}{a} + \frac{1}{b}}$$

$$F_1 = \frac{2}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}} = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}} = \frac{TP}{TP + \frac{FN + FP}{2}}$$

$$\left(\frac{2ab}{a+b} \right)$$

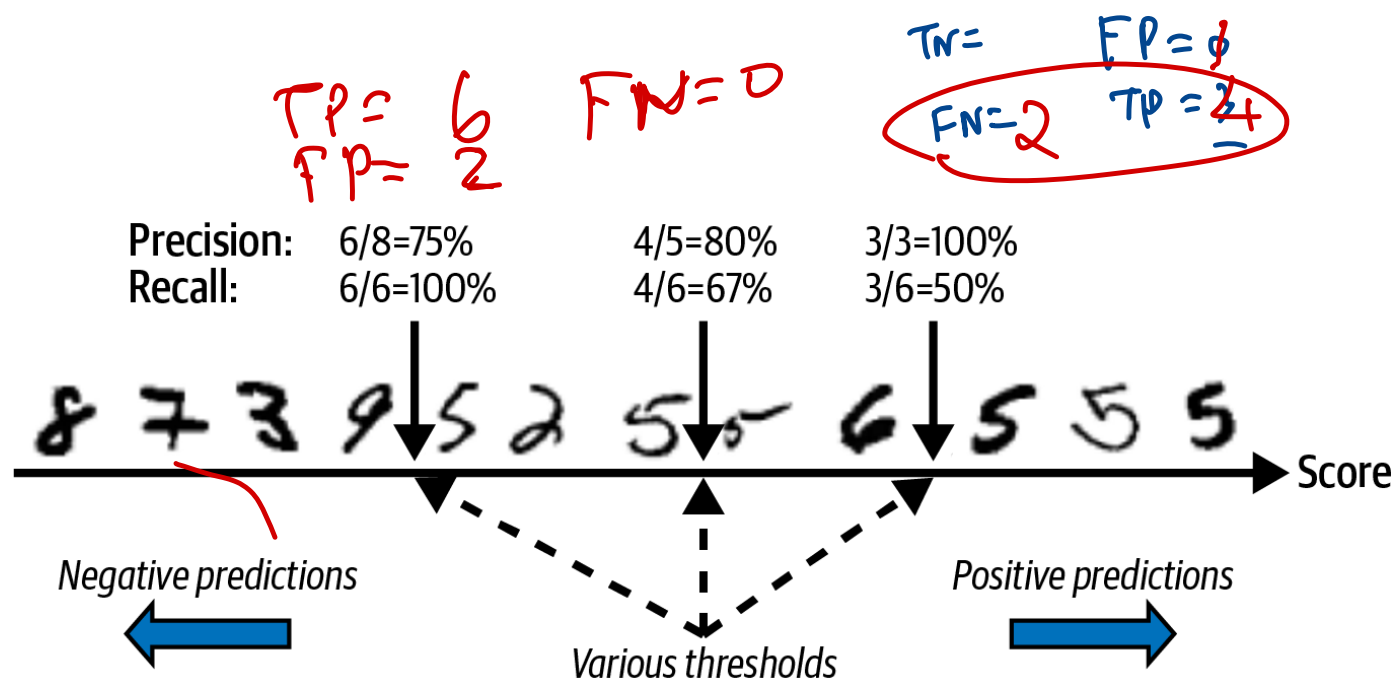

GM
AM

HM

(class 10)

Precision / Recall Tradeoff

- Based on the threshold increasing / decreasing recall / precision.
- Scikit-Learn does not let you set the threshold directly, but it does give you access to the decision scores that it uses to make predictions



Code ...


```
>>> y_scores = sgd_clf.decision_function([some_digit])
>>> y_scores
array([2164.22030239])
>>> threshold = 0
>>> y_some_digit_pred = (y_scores > threshold)
array([ True])
```

The `SGDClassifier` uses a threshold equal to 0, so the preceding code returns the same result as the `predict()` method (i.e., `True`). Let's raise the threshold:

```
>>> threshold = 3000
>>> y_some_digit_pred = (y_scores > threshold)
>>> y_some_digit_pred
array([False])
```


Precision vs. Recall

- A high precision classifier is not that useful if the recall is low. 48% recall is not great.
- Something around 60% Recall in this case is a decent solution.



ROC Curve

- Receiver Operating Characteristic (ROC)
- Plots TPR (Recall) vs. FPR (Fall out)
- $FPR = 1 - TNR$
- Sensitivity vs. 1-specificity
- There is a trade-off: the higher the recall (TPR), the more false positives (FPR) the classifier produces.
- The dotted line represents the ROC curve of a purely random classifier; a good classifier stays as far away from that line as possible (toward the top-left corner).
- One way to compare classifiers is to measure the *area under the curve* (AUC). A perfect classifier will have a ROC AUC equal to 1, whereas a purely random classifier will have a ROC AUC equal to 0.5.
- Scikit-Learn provides a function to estimate the ROC AUC:

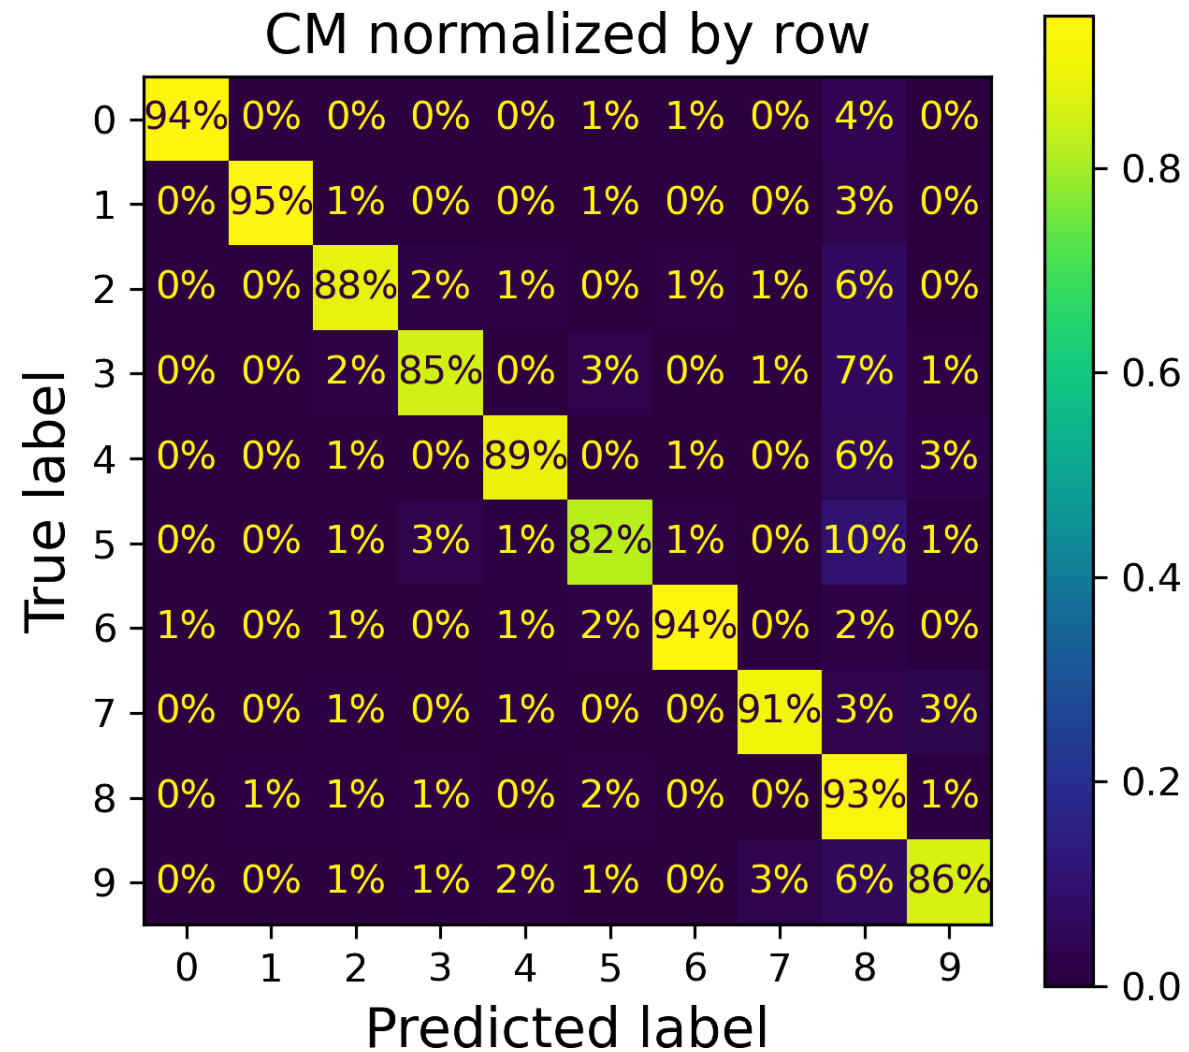



```
>>> from sklearn.metrics import roc_auc_score
>>> roc_auc_score(y_train_5, y_scores)
0.9604938554008616
```

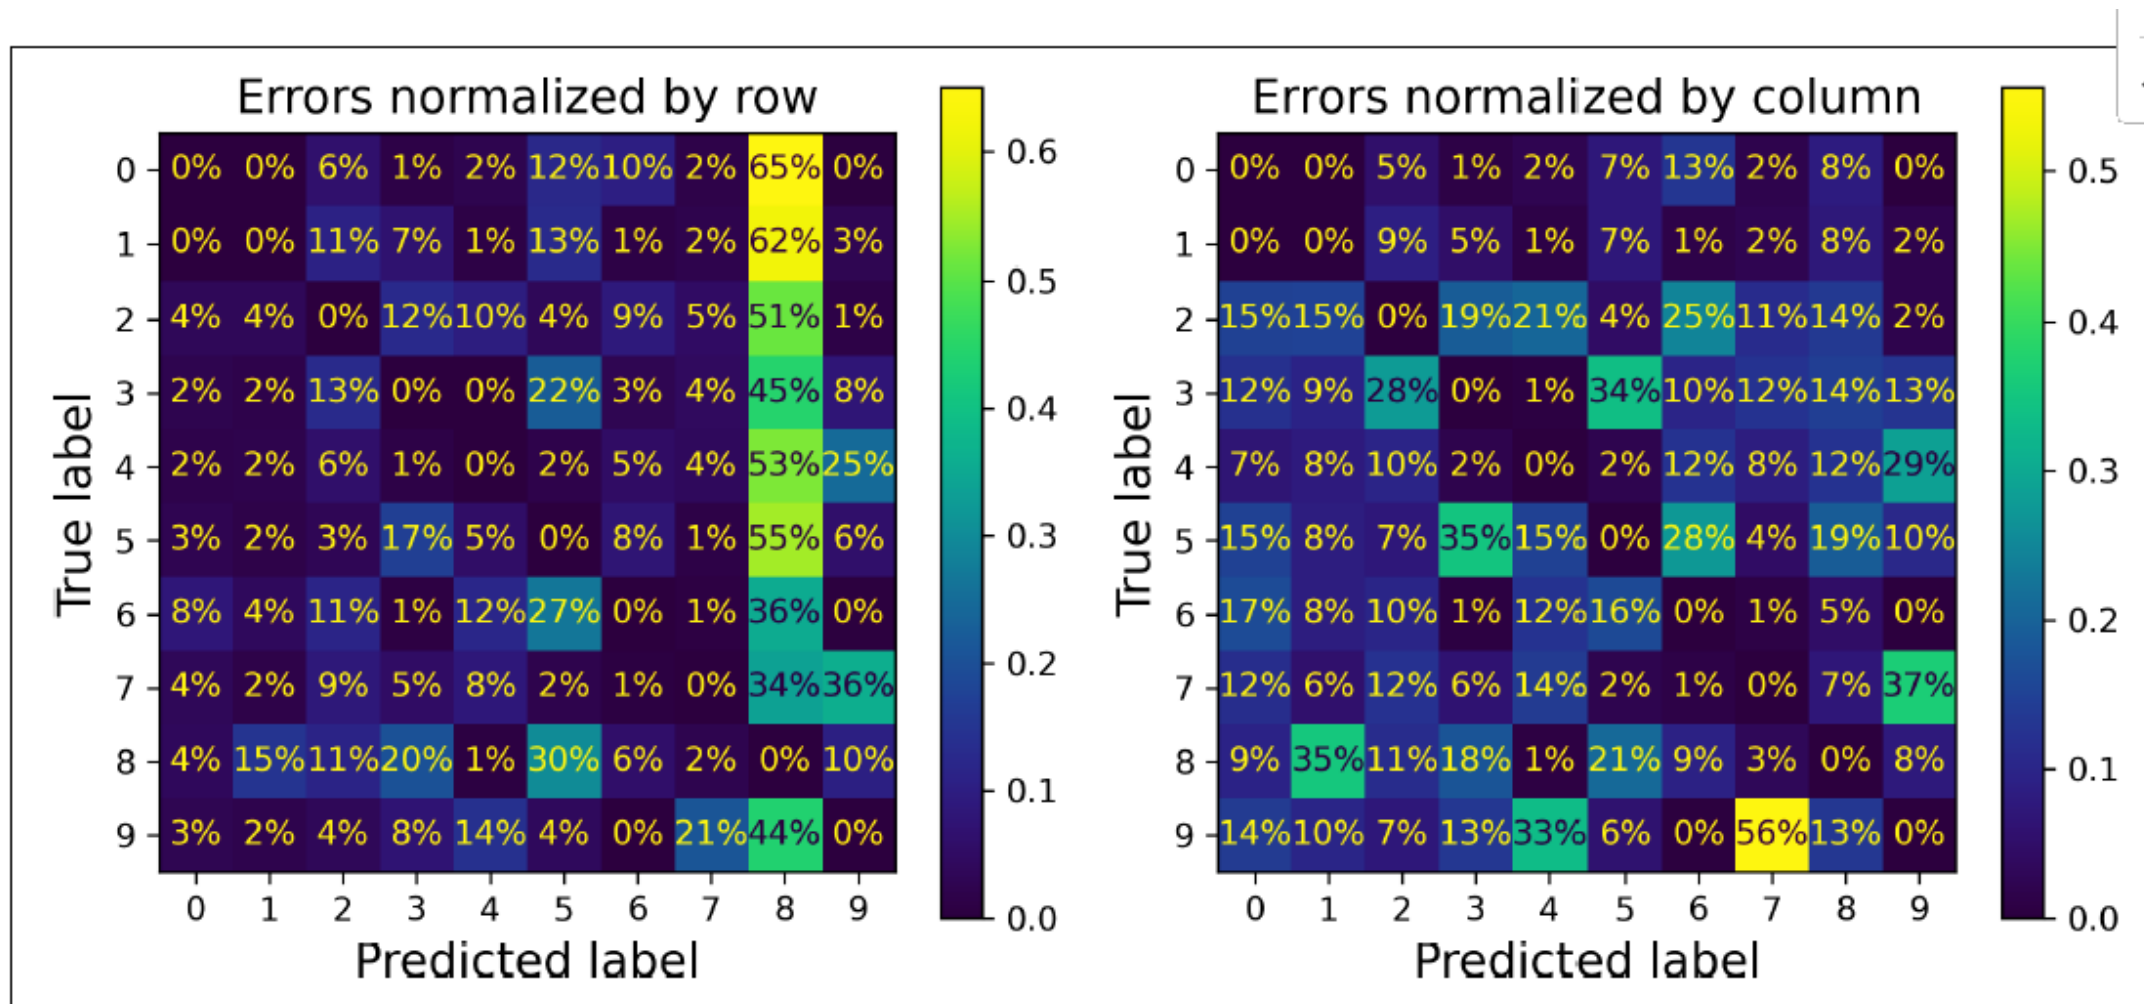
Multiclass Classifier

- Can distinguish between 2 or more classes.
- LogisticRegression, RandomForestClassifier, GaussianNB – multiclass
- SGDClassifier, SVC – are binary classifiers.
- One-versus-one (OvO) strategy
- One-versus-all (OvA) strategy

Error Analysis: Confusion matrix (left) and the same CM normalized by row (right)



- Confusion matrix with errors only, normalized by row (left) and columns (right)



Multilabel Classification

- Each instance has always been assigned to just one class. But in some cases you may want your classifier to output multiple classes for each instance.
- Consider a face-recognition classifier: what should it do if it recognizes several people in the same picture? It should attach one tag per person it recognizes.
- Say the classifier has been trained to recognize three faces: Alice, Bob, and Charlie.
- Then when the classifier is shown a picture of Alice and Charlie, it should output [True, False, True] (meaning “Alice yes, Bob no, Charlie yes”).
- **Self Reading** task: Multioutput Classification

Cross Validation (Self Reading)

- It is used for
 - Performance evaluation: Evaluate the performance of a classifier using the given data
 - Model Selection: Compare the performance of two or more algorithms (DT classifier and neural network) to determine the best algorithm for the given data
 - Tuning model parameters: Compare the performance of two variants of a parametric model

- **TYPES of cross Validation**
 - Resubstitution Validation
 - Hold-Out Validation
 - K-Fold Cross-Validation
 - Leave-One-Out Cross-Validation
 - Repeated K-Fold Cross-Validation

86

- **Repeated K-Fold Cross-Validation**
 - Repeat k-fold cross validation multiple times

- **Hold-Out Validation**
 - The database is partitioned into two non-overlapping parts, one for training and other for testing
 - The results depend a lot on the partition, may be skewed if the test set is too easy or too difficult

- **K-Fold Cross-Validation**
 - Data is partitioned into k equal folds (partitions). k-1 folds are used for training and 1-fold for testing
 - The procedure is repeated k times

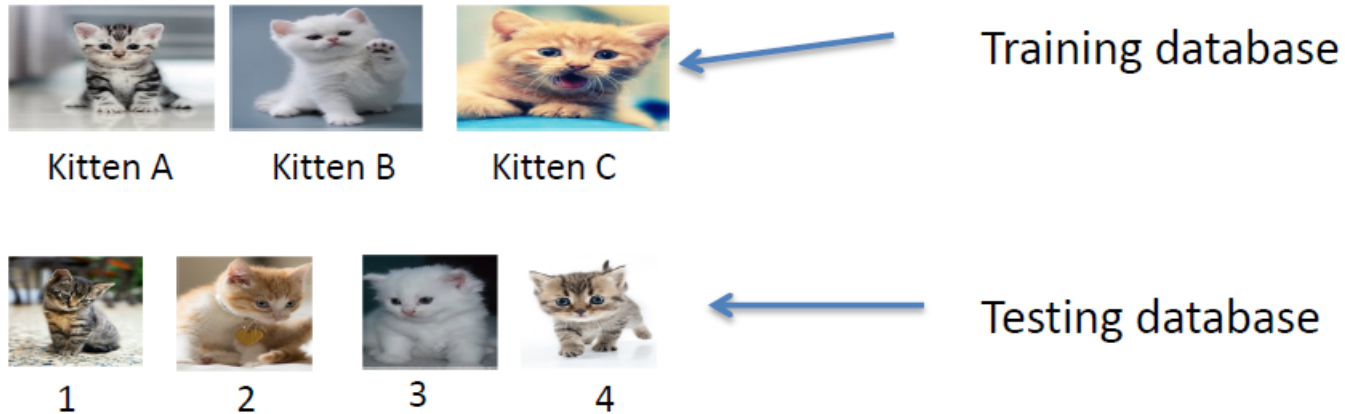
Test	Train	Train	Train
Train	Test	Train	Train
Train	Train	Test	Train
Train	Train	Train	Test

- **Leave-One-Out Cross-Validation**
 - Special case of k-fold cross validation where k=number of instances in the data
 - Testing is performed on a single instance and the remaining are used for training

- **Resubstitution Validation**
 - All the available data is used for training and the same data is used for testing
 - Does not provide any information about generalizability

Cross Validation

- “Cross-Validation is a statistical method of evaluating and comparing learning algorithms.”
- The data is divided into two parts:
 - Training: to learn or train a model
 - Testing: to validate the model



Validation Method	Advantages	Disadvantages
Resubstitution	Simple	Overfitting
Hold-out validation	Independent training and testing sets	Reduced data for training and testing
K-fold cross validation	Accurate performance estimation	Small sample for performance estimation, underestimated performance variance or overestimated degree of freedom for comparison
Leave-one-out cross validation	Unbiased performance estimation	Very large variance
Repeated k-fold cross-validation	Large number of performance estimates	Overlapped training and test data between each round, underestimated performance variance or overestimated degree of freedom for comparison